

xv6 is a re-implementation of Dennis Ritchie's and Ken Thompson's Unix Version 6 (v6). xv6 loosely follows the structure and style of v6, but is implemented for a modern x86-based multiprocessor using ANSI C.

#### ACKNOWLEDGMENTS

xv6 is inspired by John Lions's Commentary on UNIX 6th Edition (Peer to Peer Communications; ISBN: 1-57398-013-7; 1st edition (June 14, 2000)). See also <http://pdos.csail.mit.edu/6.828/2014/xv6.html>, which provides pointers to on-line resources for v6.

xv6 borrows code from the following sources:  
 JOS (asm.h, elf.h, mmu.h, bootasm.S, ide.c, console.c, and others)  
 Plan 9 (entryother.S, mp.h, mp.c, lapic.c)  
 FreeBSD (ioapic.c)  
 NetBSD (console.c)

The following people have made contributions:

Russ Cox (context switching, locking)  
 Cliff Frey (MP)  
 Xiao Yu (MP)  
 Nickolai Zeldovich  
 Austin Clements

In addition, we are grateful for the bug reports and patches contributed by Silas Boyd-Wickizer, Peter Froehlich, Shivam Handa, Anders Kaseorg, Eddie Kohler, Yandong Mao, Hitoshi Mitake, Carmi Merimovich, Joel Nider, Greg Price, Eldar Sehayek, Yongming Shen, Stephen Tu, and Zouchangwei.

The code in the files that constitute xv6 is  
 Copyright 2006-2014 Frans Kaashoek, Robert Morris, and Russ Cox.

#### ERROR REPORTS

If you spot errors or have suggestions for improvement, please send email to Frans Kaashoek and Robert Morris ([kaashoek,rtm@csail.mit.edu](mailto:kaashoek,rtm@csail.mit.edu)).

#### BUILDING AND RUNNING XV6

To build xv6 on an x86 ELF machine (like Linux or FreeBSD), run "make". On non-x86 or non-ELF machines (like OS X, even on x86), you will need to install a cross-compiler gcc suite capable of producing x86 ELF binaries. See <http://pdos.csail.mit.edu/6.828/2014/tools.html>. Then run "make TOOLPREFIX=i386-jos-elf-".

To run xv6, install the QEMU PC simulators. To run in QEMU, run "make qemu".

To create a typeset version of the code, run "make xv6.pdf". This requires the "mpage" utility. See <http://www.mesa.nl/pub/mpage/>.

The numbers to the left of the file names in the table are sheet numbers. The source code has been printed in a double column format with fifty lines per column, giving one hundred lines per sheet (or page). Thus there is a convenient relationship between line numbers and sheet numbers.

# basic headers	34 trapasm.S	73 mp.c
01 types.h	35 trap.c	75 lapic.c
01 param.h	36 syscall.h	78 ioapic.c
02 memlayout.h	37 syscall.c	79 picirq.c
02 defs.h	39 sysproc.c	80 kbd.h
04 x86.h		81 kbd.c
06 asm.h	# file system	82 console.c
07 mmu.h	41 buf.h	85 timer.c
09 elf.h	42 fcntl.h	86 uart.c
	42 stat.h	
# entering xv6	43 fs.h	# user-level
10 entry.S	44 file.h	87 initcode.S
11 entryother.S	45 ide.c	87 usys.S
12 main.c	47 bio.c	88 init.c
	49 log.c	88 sh.c
# locks	51 fs.c	
15 spinlock.h	60 file.c	# bootloader
15 spinlock.c	62 sysfile.c	95 bootasm.S
	67 exec.c	96 bootmain.c
# processes		
17 vm.c	# pipes	# add student files her
23 proc.h	68 pipe.c	97 sfdate.c
24 proc.c		98 time.c
31 swtch.S	# string operations	99 user.h
32 kalloc.c	70 string.c	99 halt.c
		100 ps.h
# system calls	# low-level hardware	100 sfps.c
33 traps.h	71 mp.h	101 test.c
34 vectors.pl		

The source listing is preceded by a cross-reference that lists every defined constant, struct, global variable, and function in xv6. Each entry gives, on the same line as the name, the line number (or, in a few cases, numbers) where the name is defined. Successive lines in an entry list the line numbers where the name is used. For example, this entry:

```
swtch 2658
      0374 2428 2466 2657 2658
```

indicates that swtch is defined on line 2658 and is mentioned on five lines on sheets 03, 24, and 26.

```

acquire 1574
  0383 1574 1578 2460 2620
  2655 2683 2767 2816 2868
  2883 2918 2931 2956 2971
  3034 3073 3107 3276 3293
  3566 4058 4078 4607 4665
  4770 4831 5030 5057 5074
  5131 5408 5441 5461 5490
  5510 5520 6029 6054 6068
  6913 6934 6955 8260 8431
  8477 8513
allocproc 2455
  2455 2507 2580
allocuvm 1953
  0428 1953 1967 2559 6746
  6758
alltraps 3454
  3409 3417 3430 3435 3453
  3454
ALT 8010
  8010 8038 8040
argfd 6219
  6219 6256 6271 6283 6294
  6306
argint 3745
  0401 3745 3758 3774 3984
  4016 4026 4038 4056 4121
  6224 6271 6283 6508 6576
  6577 6631
argptr 3754
  0402 3754 4106 4125 6271
  6283 6306 6657
argstr 3771
  0403 3771 6318 6408 6508
  6557 6575 6607 6631
__attribute__ 1310
  0272 0365 1209 1310 9906
BACK 8862
  8862 8977 9270 9539
backcmd 8900 9264
  8900 8914 8978 9264 9266
  9392 9505 9540
BACKSPACE 8350
  8350 8367 8409 8441 8447
ballocc 5204
  5204 5224 5567 5575 5579
BLOCK 4360
  4360 5211 5235
B_BUSY 4159
  4159 4658 4776 4777 4790
  4793 4817 4828 4840
B_DIRTY 4161
  4161 4593 4616 4621 4660
  4678 4790 4819 5139
begin_op 5028
  0336 2650 5028 6083 6174
  6321 6411 6511 6556 6574
  6606 6720
bfree 5229
  5229 5614 5624 5627
bget 4766
  4766 4798 4806
binit 4739
  0263 1231 4739
bmap 5560
  5322 5560 5586 5669 5719
bootmain 9667
  9618 9667
BPB 4357
  4357 4360 5210 5212 5236
bread 4802
  0264 4802 4977 4978 4990
  5006 5088 5089 5182 5193
  5211 5235 5360 5381 5468
  5576 5620 5669 5719
brelse 4826
  0265 4826 4829 4981 4982
  4997 5014 5092 5093 5184
  5196 5217 5222 5242 5366
  5369 5390 5476 5582 5626
  5672 5723
BSIZE 4305
  4157 4305 4323 4351 4357
  4581 4595 4617 4958 4979
  5090 5194 5669 5670 5671
  5715 5719 5720 5721
buf 4150
  0250 0264 0265 0266 0308
  0335 2120 2123 2132 2134
  4150 4154 4155 4156 4512
  4528 4531 4575 4604 4654
  4656 4659 4727 4731 4735
  4741 4753 4765 4768 4801
  4804 4815 4826 4905 4977
  4978 4990 4991 4997 5006
  5007 5013 5014 5088 5089
  5122 5169 5180 5191 5207
  5231 5356 5378 5455 5563
  5609 5655 5705 8229 8240
  8244 8247 8418 8439 8453

```

```

  8487 8508 8515 8987 8990
  8991 8992 9105 9117 9119
  9122 9123 9124 9127 9128
  9132
B_VALID 4160
  4160 4620 4660 4678 4807
bwrite 4815
  0266 4815 4818 4980 5013
  5091
bzero 5189
  5189 5218
C 8031 8424
  8031 8079 8104 8105 8106
  8107 8108 8110 8424 8434
  8437 8444 8455 8488
CAPSLOCK 8012
  8012 8045 8186
cgaputc 8355
  8355 8413
clearpteu 2029
  0437 2029 2035 6760
cli 0557
  0557 0559 1126 1660 8310
  8404 9562
cmd 8866
  8866 8878 8887 8888 8893
  8894 8902 8907 8911 8920
  8923 8928 8936 8942 8946
  8954 8978 8980 9069 9081
  9085 9086 9202 9205 9207
  9208 9209 9210 9213 9214
  9216 9218 9219 9220 9221
  9222 9223 9224 9225 9226
  9229 9230 9232 9234 9235
  9236 9237 9238 9239 9250
  9251 9253 9255 9256 9257
  9258 9259 9260 9263 9264
  9266 9268 9269 9270 9271
  9272 9362 9363 9364 9365
  9367 9371 9374 9380 9381
  9384 9387 9389 9392 9396
  9398 9400 9403 9405 9408
  9410 9413 9414 9425 9428
  9431 9435 9450 9453 9458
  9462 9463 9466 9471 9472
  9478 9487 9488 9494 9495
  9501 9502 9511 9514 9516
  9522 9523 9528 9534 9540
  9541 9544
CMOS_PORT 7685
  7685 7699 7700 7738
CMOS_RETURN 7686
  7686 7741
CMOS_STATA 7725
  7725 7773
CMOS_STATB 7726
  7726 7766
CMOS_UIP 7727
  7727 7773
COM1 8613
  8613 8623 8626 8627 8628
  8629 8630 8631 8634 8640
  8641 8657 8659 8667 8669
commit 5101
  4953 5073 5101
CONSOLE 4437
  4437 8527 8528
consoleinit 8523
  0269 1227 8523
consoleintr 8427
  0271 8198 8427 8675
consoleread 8470
  8470 8528
consolewrite 8508
  8508 8527
conspugetc 8401
  8216 8247 8268 8286 8289
  8293 8294 8401 8441 8447
  8454 8515
context 2360
  0251 0380 2308 2360 2382
  2488 2489 2490 2491 2780
  2808 3020 3093
CONV 7782
  7782 7783 7784 7785 7786
  7787 7788 7789
copyout 2118
  0436 2118 6768 6779
copyuvm 2053
  0433 2053 2064 2066 2584
cprintf 8252
  0270 1224 1264 1967 3018
  3022 3024 3590 3603 3608
  3901 4086 5322 7419 7439
  7661 7862 8252 8312 8313
  8314 8317
cpu 2306
  0311 1224 1264 1266 1278
  1506 1566 1587 1608 1646
  1661 1662 1670 1672 1718

```

```

1731 1737 1876 1877 1878
1879 2306 2316 2320 2331
2780 2801 2807 2808 2809
3093 3565 3590 3591 3603
3604 3608 3610 7313 7314
7661 8312
cpunum 7651
0326 1288 1724 7651 7873
7882
CR0_PE 0727
0727 1135 1171 9593
CR0_PG 0737
0737 1050 1171
CR0_WP 0733
0733 1050 1171
CR4_PSE 0739
0739 1043 1164
create 6457
6457 6477 6490 6494 6514
6557 6578
CRTPORT 8351
8351 8360 8361 8362 8363
8381 8382 8383 8384
CTL 8009
8009 8035 8039 8185
DAY 7732
7732 7755
deallocuvn 1982
0429 1968 1982 2016 2562
DEVSPACE 0204
0204 1832 1845
devsw 4430
4430 4435 5658 5660 5708
5710 6011 8527 8528
dinode 4327
4327 4351 5357 5361 5379
5382 5456 5469
dirent 4365
4365 5764 5805 6366 6404
dirlink 5802
0288 5771 5802 5817 5825
6341 6489 6493 6494
dirlookup 5761
0289 5761 5767 5809 5925
6423 6467
DIRSIZ 4363
4363 4367 5755 5822 5878
5879 5942 6315 6405 6461
dobuiltin 9081
9081 9128
DPL_USER 0779
0779 1727 1728 2514 2515
3523 3618 3627
EOESC 8016
8016 8170 8174 8175 8177
8180
elfhdr 0955
0955 6715 9669 9674
ELF_MAGIC 0952
0952 6731 9680
ELF_PROG_LOAD 0986
0986 6742
end_op 5053
0337 2652 5053 6085 6179
6323 6330 6348 6357 6413
6447 6452 6516 6521 6527
6536 6540 6558 6562 6579
6583 6608 6614 6619 6722
6752 6805
entry 1040
0961 1036 1039 1040 3402
3403 6792 7171 9671 9695
9696
EOI 7515
7515 7634 7675
ERROR 7536
7536 7627
ESR 7518
7518 7630 7631
exec 6710
0275 3858 6647 6710 8768
8829 8830 8931 8932 9844
9845 9913
EXEC 8858
8858 8927 9209 9515
execcmd 8870 9203
8870 8915 8928 9203 9205
9471 9477 9478 9506 9516
exit 2633 9858
0359 2633 2672 3555 3559
3619 3628 3853 3969 8716
8719 8761 8826 8831 8921
8930 8940 8983 9135 9142
9760 9764 9816 9828 9841
9847 9851 9858 9906 10067
10076 10136
EXTMEM 0202
0202 0208 1829
fdalloc 6238
6238 6258 6532 6662

```

```

fetchint 3717
0404 3717 3747 6638
fetchstr 3729
0405 3729 3776 6644
file 4400
0252 0278 0279 0280 0282
0283 0284 0351 2385 4400
5170 6008 6014 6024 6027
6030 6051 6052 6064 6066
6102 6115 6152 6213 6219
6222 6238 6253 6267 6279
6292 6303 6505 6654 6856
6871 8210 8608 8879 8938
8939 9214 9222 9422
filealloc 6025
0278 6025 6532 6877
fileclose 6064
0279 2644 6064 6070 6297
6534 6665 6666 6904 6906
filedup 6052
0280 2602 6052 6056 6260
fileinit 6018
0281 1232 6018
fileread 6115
0282 6115 6130 6273
filestat 6102
0283 6102 6308
filewrite 6152
0284 6152 6184 6189 6285
FL_IF 0710
0710 1662 1668 2518 2805
7658
fork 2574
0360 2574 3852 3963 8760
8823 8825 9155 9157 9821
9850 9905
forkl 9151
8905 8947 8957 8964 8979
9131 9151
forkret 2826
2435 2491 2826
freerange 3251
3211 3234 3240 3251
freevm 2010
0430 2010 2015 2078 2696
6795 6802
FSSIZE 0162
0162 4579
gatedesc 0901
0523 0526 0901 3511
getbuiltin 9051
9051 9076
getcallerpcs 1626
0384 1588 1626 3020 8315
getcmd 8987
8987 9117
getgid 2985
0376 2985 3876 4003 8784
9060 9929 10123 10128
getprocs 3029
0372 3029 3880 4128 8788
9933 10062 10066
gettoken 9306
9306 9391 9395 9407 9420
9421 9457 9461 9483
getuid 2979
0375 2979 3875 3998 8783
9056 9928 10115 10120
GID_DEFAULT 2303
2303 2527
growproc 2553
0361 2553 4041
havedisk1 4530
4530 4564 4662
holding 1644
0385 1577 1604 1644 2799
HOURS 7731
7731 7754
ialloc 5353
0290 5353 5371 6476 6477
IBLOCK 4354
4354 5360 5381 5468
I_BUSY 4425
4425 5462 5464 5487 5491
5513 5515
ICRHI 7529
7529 7637 7707 7719
ICRLO 7519
7519 7638 7639 7708 7710
7720
ID 7512
7512 7548 7666
IDE_BSY 4515
4515 4539
IDE_CMD_READ 4520
4520 4597
IDE_CMD_WRITE 4521
4521 4594
IDE_DF 4517
4517 4541

```

```

IDE_DRDY 4516          0299 0300 0301 0302 0303
    4516 4539          0432 1918 2386 4406 4412
IDE_ERR 4518          4431 4432 5173 5314 5326
    4518 4541          5352 5376 5403 5406 5412
ideinit 4551          5438 5439 5453 5485 5508
    0306 1233 4551      5530 5560 5606 5637 5652
ideintr 4602          5702 5760 5761 5802 5806
    0307 3574 4602      5904 5907 5939 5950 6316
idelock 4527          6363 6403 6456 6460 6506
    4527 4555 4607 4609 4628  6554 6569 6604 6716 8470
    4665 4679 4682      8508
iderw 4654          INPUT_BUF 8416
    0308 4654 4659 4661 4663  8416 8418 8439 8451 8453
    4808 4820          8455 8487
idestart 4575        insl 0462
    4531 4575 4578 4584 4626  0462 0464 4617 9723
    4675          install_trans 4972
idewait 4535          4972 5021 5106
    4535 4558 4586 4616      INT_DISABLED 7819
idtinit 3529          7819 7867
    0412 1265 3529      ioapic 7827
idup 5439            7407 7429 7430 7824 7827
    0291 2603 5439 5912      7836 7837 7843 7844 7858
iget 5404            IOAPIC 7808
    5326 5367 5404 5424 5779  7808 7858
    5910          ioapicenable 7873
iinit 5318            0311 4557 7873 8532 8643
    0292 2837 5318      ioapicid 7317
ilock 5453            0312 7317 7430 7447 7861
    0293 5453 5459 5479 5915  7862
    6105 6124 6175 6327 6340  ioapicinit 7851
    6353 6417 6425 6465 6469  0313 1226 7851 7862
    6479 6524 6611 6725 8482  ioapicread 7834
    8502 8517          7834 7859 7860
inb 0453            ioapicwrite 7841
    0453 4539 4563 7454 7741  7841 7867 7868 7881 7882
    8164 8167 8361 8363 8634  IO_PIC1 7907
    8640 8641 8657 8667 8669  7907 7920 7935 7944 7947
    9573 9581 9704          7952 7962 7976 7977
initlock 1562        IO_PIC2 7908
    0386 1562 2443 3232 3525  7908 7921 7936 7965 7966
    4555 4743 4962 5320 6020  7967 7970 7979 7980
    6885 8525          IO_TIMER1 8559
initlog 4956          8559 8568 8578 8579
    0334 2838 4956 4959      IPB 4351
inituvm 1903          4351 4354 5361 5382 5469
    0431 1903 1908 2511      iput 5508
inode 4412            0294 2651 5508 5514 5533
    0253 0288 0289 0290 0291  5810 5933 6084 6346 6618
    0293 0294 0295 0296 0297  IRQ_COM1 3383

```

```

    3383 3584 8642 8643      0207 0208 0212 0213 0217
IRQ_ERROR 3385          0218 0220 0221 1315 1633
    3385 7627          1829 1958 2016
IRQ_IDE 3384          KERNLINK 0208
    3384 3573 3577 4556 4557  0208 1830
IRQ_KBD 3382          KEY_DEL 8028
    3382 3580 8531 8532      8028 8069 8091 8115
IRQ_SLAVE 7910          KEY_DN 8022
    7910 7914 7952 7967      8022 8065 8087 8111
IRQ_SPURIOUS 3386      KEY_END 8020
    3386 3589 7607          8020 8068 8090 8114
IRQ_TIMER 3381          KEY_HOME 8019
    3381 3564 3623 7614 8580  8019 8068 8090 8114
isdirempty 6363        KEY_INS 8027
    6363 6370 6429          8027 8069 8091 8115
ismp 7315            KEY_LF 8023
    0340 1234 7315 7412 7420  8023 8067 8089 8113
    7440 7443 7855 7875      KEY_PGDN 8026
itrunc 5606            8026 8066 8088 8112
    5173 5517 5606          KEY_PGUP 8025
iunlock 5485            8025 8066 8088 8112
    0295 5485 5488 5532 5922  KEY_RT 8024
    6107 6127 6178 6336 6539  8024 8067 8089 8113
    6617 8475 8512          KEY_UP 8021
iunlockput 5530        8021 8065 8087 8111
    0296 5530 5917 5926 5929  kfree 3265
    6329 6342 6345 6356 6430  0317 1998 2000 2020 2023
    6441 6445 6451 6468 6472  2585 2694 3256 3265 3270
    6496 6526 6535 6561 6582  6902 6923
    6613 6751 6804          kill 2927
iupdate 5376          0362 2927 3609 3857 3986
    0297 5376 5519 5632 5728  8767 9912
    6335 6355 6439 6444 6483  kinit1 3230
    6487          0318 1219 3230
I_INVALID 4426        kinit2 3238
    4426 5467 5477 5511      0319 1237 3238
kalloc 3288          KSTACKSIZE 0151
    0316 1294 1763 1842 1909  0151 1054 1063 1295 1879
    1965 2069 2473 3288 6879  2477
KBDATAP 8004          kvmalloc 1857
    8004 8167          0424 1220 1857
kbdgetc 8156          lapiceoi 7672
    8156 8198          0328 3571 3575 3582 3586
kbdtintr 8196        3592 7672
    0322 3581 8196          lapicinit 7601
KBS_DIB 8003          0329 1222 1256 7601
    8003 8165          lapicstartap 7691
KBSTATP 8002          0330 1299 7691
    8002 8164          lapicw 7545
KERNBASE 0207        7545 7607 7613 7614 7615

```

```

7618 7619 7624 7627 7630
7631 7634 7637 7638 7643
7675 7707 7708 7710 7719
7720
lcr3 0590
    0590 1868 1883
lgdt 0512
    0512 0520 1133 1733 9591
lidt 0526
    0526 0534 3531
LINT0 7534
    7534 7618
LINT1 7535
    7535 7619
LIST 8861
    8861 8945 9257 9533
listcmd 8891 9251
    8891 8916 8946 9251 9253
    9396 9507 9534
loadgs 0551
    0551 1734
loadvm 1918
    0432 1918 1924 1927 6748
log 4937 4950
    4937 4950 4962 4964 4965
    4966 4976 4977 4978 4990
    4993 4994 4995 5006 5009
    5010 5011 5022 5030 5032
    5033 5034 5036 5038 5039
    5057 5058 5059 5060 5061
    5063 5066 5068 5074 5075
    5076 5077 5087 5088 5089
    5103 5107 5126 5128 5131
    5132 5133 5136 5137 5138
    5140
logheader 4932
    4932 4944 4958 4959 4991
    5007
LOGSIZE 0160
    0160 4934 5034 5126 6167
log_write 5122
    0335 5122 5129 5195 5216
    5241 5365 5389 5580 5722
ltr 0538
    0538 0540 1880
makeint 9013
    9013 9034 9040
mappages 1779
    1779 1848 1911 1972 2072
MAXARG 0158

```

```

0158 6627 6714 6765
MAXARGS 8864
    8864 8872 8873 9490
MAXFILE 4324
    4324 5715
MAXOPBLOCKS 0159
    0159 0160 0161 5034
MAX_PROC 10054
    10054 10060 10062
memcmp 7015
    0392 7015 7345 7388 7776
memmove 7031
    0393 1285 1912 2071 2132
    4979 5090 5183 5388 5475
    5671 5721 5879 5881 7031
    7054 8376 9937
memset 7004
    0394 1766 1844 1910 1971
    2490 2513 3273 5194 5363
    6434 6634 7004 8378 8990
    9208 9219 9235 9256 9269
    9943
microdelay 7681
    0331 7681 7709 7711 7721
    7739 8658
min 5172
    5172 5670 5720 9806 9831
    9836 9839
MINS 7730
    7730 7753
MONTH 7733
    7733 7756
mp 7152
    7152 7308 7337 7344 7345
    7346 7355 7360 7364 7365
    7368 7369 7380 7383 7385
    7387 7394 7404 7410 7450
mpbcpu 7320
    0341 7320
MPBUS 7202
    7202 7433
mpconf 7163
    7163 7379 7382 7387 7405
mpconfig 7380
    7380 7410
mpenter 1252
    1252 1296
mpinit 7401
    0342 1221 7401 7419 7439
mpioapic 7189

```

```

7189 7407 7429 7431
MPIOAPIC 7203
    7203 7428
MPIOINTR 7204
    7204 7434
MPLINTR 7205
    7205 7435
mpmain 1262
    1209 1240 1257 1262
mpproc 7178
    7178 7406 7417 7426
MPPROC 7201
    7201 7416
mpsearch 7356
    7356 7385
mpsearch1 7338
    7338 7364 7368 7371
multiboot_header 1025
    1024 1025
namecmp 5753
    0298 5753 5774 6420
namei 5940
    0299 2523 5940 6322 6520
    6607 6721
nameiparent 5951
    0300 5905 5920 5932 5951
    6338 6412 6463
namex 5905
    5905 5943 5953
NBUF 0161
    0161 4731 4753
ncpu 7316
    1224 1287 2321 4557 7316
    7418 7419 7423 7424 7425
    7445
NCPU 0152
    0152 2320 7313
NDEV 0156
    0156 5658 5708 6011
NDIRECT 4322
    4322 4324 4333 4423 5565
    5570 5574 5575 5612 5619
    5620 5627 5628
NELEM 0440
    0440 1847 3014 3892 6636
nextpid 2434
    2434 2469
NFILE 0154
    0154 6014 6030
NINDIRECT 4323

```

```

4323 4324 5572 5622
NINODE 0155
    0155 5314 5412
NO 8006
    8006 8052 8055 8057 8058
    8059 8060 8062 8074 8077
    8079 8080 8081 8082 8084
    8102 8103 8105 8106 8107
    8108
NOFILE 0153
    0153 2385 2600 2642 6226
    6242
NPENTRIES 0821
    0821 1311 2017
NPROC 0150
    0150 2426 2461 2661 2687
    2768 2907 2932 3011 3035
NPTENTRIES 0822
    0822 1994
NSEGS 2301
    1711 2301 2310
NULL 2410
    2410 2532 3077 3080 3086
    3090 3136 3140
nulterminate 9502
    9365 9380 9502 9523 9529
    9530 9535 9536 9541
NUMLOCK 8013
    8013 8046
O_CREATE 4203
    4203 6513 9428 9431
O_RDONLY 4200
    4200 6525 9425
O_RDWR 4202
    4202 6546 8814 8816 9109
outb 0471
    0471 4561 4570 4587 4588
    4589 4590 4591 4592 4594
    4597 7453 7454 7699 7700
    7738 7920 7921 7935 7936
    7944 7947 7952 7962 7965
    7966 7967 7970 7976 7977
    7979 7980 8360 8362 8381
    8382 8383 8384 8577 8578
    8579 8623 8626 8627 8628
    8629 8630 8631 8659 9578
    9586 9714 9715 9716 9717
    9718 9719
outsl 0483
    0483 0485 4595

```

outw 0477  
     0477 1181 1183 4087 9624  
     9626  
 O\_WRONLY 4201  
     4201 6545 6546 9428 9431  
 P2V 0218  
     0218 1219 1237 7362 7701  
     8352  
 panic 8305 9139  
     0272 1578 1605 1669 1671  
     1790 1846 1882 1908 1924  
     1927 1998 2015 2035 2064  
     2066 2510 2639 2672 2800  
     2802 2804 2806 2856 2859  
     3270 3605 4578 4580 4584  
     4659 4661 4663 4798 4818  
     4829 4959 5060 5127 5129  
     5224 5239 5371 5424 5459  
     5479 5488 5514 5586 5767  
     5771 5817 5825 6056 6070  
     6130 6184 6189 6370 6428  
     6436 6477 6490 6494 8263  
     8305 8312 8373 8906 8925  
     8956 9139 9157 9378 9422  
     9456 9460 9486 9491  
 panicked 8218  
     8218 8318 8403  
 parseblock 9451  
     9451 9456 9475  
 parsecmd 9368  
     8907 9132 9368  
 parseexec 9467  
     9364 9405 9467  
 parseline 9385  
     9362 9374 9385 9396 9458  
 parsepipe 9401  
     9363 9389 9401 9408  
 parseredirs 9414  
     9414 9462 9481 9492  
 PCINT 7533  
     7533 7624  
 pde\_t 0103  
     0103 0426 0427 0428 0429  
     0430 0431 0432 0433 0436  
     0437 1210 1270 1311 1710  
     1754 1756 1779 1836 1839  
     1842 1903 1918 1953 1982  
     2010 2029 2052 2053 2055  
     2102 2118 2373 6718  
 PDX 0812

0812 1759  
 PDXSHIFT 0827  
     0812 0818 0827 1315  
 peek 9351  
     9351 9375 9390 9394 9406  
     9419 9455 9459 9474 9482  
 PGROUNDNDOWN 0830  
     0830 1784 1785 2125  
 PGROUNDUP 0829  
     0829 1963 1990 3254 6757  
 PGSIZE 0823  
     0823 0829 0830 1310 1766  
     1794 1795 1844 1907 1910  
     1911 1923 1925 1929 1932  
     1964 1971 1972 1991 1994  
     2062 2071 2072 2129 2135  
     2512 2519 3255 3269 3273  
     6758 6760  
 PHYSTOP 0203  
     0203 1237 1831 1845 1846  
     3269  
 picenable 7925  
     0346 4556 7925 8531 8580  
     8642  
 picinit 7932  
     0347 1225 7932  
 picsetmask 7917  
     7917 7927 7983  
 pinit 2441  
     0363 1229 2441  
 pipe 6861  
     0254 0352 0353 0354 3855  
     4405 6081 6122 6159 6861  
     6873 6879 6885 6889 6893  
     6911 6930 6951 8763 8955  
     8956 9908  
 PIPE 8860  
     8860 8953 9236 9527  
 pipealloc 6871  
     0351 6659 6871  
 pipeclose 6911  
     0352 6081 6911  
 pipecmd 8885 9230  
     8885 8917 8954 9230 9232  
     9408 9508 9528  
 piperead 6951  
     0353 6122 6951  
 PIPESIZE 6859  
     6859 6863 6936 6944 6966  
 pipewrite 6930

0354 6159 6930  
 popcli 1666  
     0389 1621 1666 1669 1671  
     1884  
 printint 8226  
     8226 8276 8280  
 proc 2371  
     0255 0358 0378 0434 1205  
     1558 1706 1738 1873 1879  
     2317 2332 2371 2380 2388  
     2406 2426 2427 2432 2454  
     2457 2461 2504 2557 2559  
     2562 2565 2566 2577 2584  
     2590 2591 2592 2601 2602  
     2603 2605 2608 2609 2610  
     2635 2638 2643 2644 2645  
     2651 2653 2658 2661 2662  
     2670 2680 2687 2688 2708  
     2714 2760 2768 2777 2780  
     2785 2803 2808 2817 2818  
     2855 2873 2874 2878 2905  
     2907 2929 2932 2957 2972  
     2981 2987 3007 3011 3032  
     3035 3066 3087 3093 3098  
     3105 3129 3134 3505 3554  
     3556 3558 3601 3609 3610  
     3612 3618 3623 3627 3705  
     3719 3733 3736 3747 3760  
     3891 3893 3902 3903 3957  
     3992 4009 4040 4061 4507  
     5166 5912 6211 6226 6243  
     6244 6296 6618 6620 6664  
     6704 6786 6789 6790 6791  
     6792 6793 6794 6854 6937  
     6957 7311 7406 7417 7418  
     7419 7422 8213 8480 8610  
 procdump 3004  
     0364 3004 8465  
 proghdr 0974  
     0974 6717 9670 9684  
 PTE\_ADDR 0844  
     0844 1761 1928 1996 2019  
     2067 2111  
 PTE\_FLAGS 0845  
     0845 2068  
 PTE\_P 0833  
     0833 1313 1315 1760 1770  
     1789 1791 1995 2018 2065  
     2107  
 PTE\_PS 0840

0840 1313 1315  
 pte\_t 0848  
     0848 1753 1757 1761 1763  
     1782 1921 1984 2031 2056  
     2104  
 PTE\_U 0835  
     0835 1770 1911 1972 2036  
     2109  
 PTE\_W 0834  
     0834 1313 1315 1770 1829  
     1831 1832 1911 1972  
 PTX 0815  
     0815 1772  
 PTXSHIFT 0826  
     0815 0818 0826  
 pushcli 1655  
     0388 1576 1655 1875  
 rcr2 0582  
     0582 3604 3611  
 readeflags 0544  
     0544 1659 1668 2805 7658  
 read\_head 4988  
     4988 5020  
 readi 5652  
     0301 1933 5652 5770 5816  
     6125 6369 6370 6729 6740  
 readsb 5178  
     0287 4963 5178 5234 5321  
 readsect 9710  
     9710 9745  
 readseg 9729  
     9664 9677 9688 9729  
 recover\_from\_log 5018  
     4952 4967 5018  
 REDIR 8859  
     8859 8935 9220 9521  
 redircmd 8876 9214  
     8876 8918 8936 9214 9216  
     9425 9428 9431 9509 9522  
 REG\_ID 7810  
     7810 7860  
 REG\_TABLE 7812  
     7812 7867 7868 7881 7882  
 REG\_VER 7811  
     7811 7859  
 release 1602  
     0387 1602 1605 2464 2470  
     2624 2702 2709 2787 2820  
     2830 2869 2882 2920 2940  
     2944 2958 2973 3040 3057

```

3100 3111 3116 3121 3281
3298 3569 4062 4067 4080
4609 4628 4682 4778 4794
4843 5039 5068 5077 5140
5415 5431 5443 5465 5493
5516 5525 6033 6037 6058
6072 6078 6922 6925 6938
6947 6958 6969 8301 8463
8481 8501 8516
ROOTDEV 0157
0157 2837 2838 5910
ROOTINO 4304
4304 5910
run 3214
3214 3215 3221 3267 3277
3290
runcmd 8911
8911 8925 8942 8948 8950
8962 8969 8980 9132
RUNNING 2368
2368 2417 2779 2803 3092
3623
safestrcpy 7082
0395 2522 2605 3048 3050
6786 7082
sb 5174
0287 4354 4360 4961 4963
4964 4965 5174 5178 5183
5210 5211 5212 5234 5235
5321 5322 5323 5359 5360
5381 5468 7764 7766 7768
sched 2795
0366 2671 2795 2800 2802
2804 2806 2819 2875
scheduler 2758
0365 1267 2308 2758 2780
2808 3065 3093
SCROLLLOCK 8014
8014 8047
SECS 7729
7729 7752
SECTOR_SIZE 4514
4514 4581
SECTSIZE 9662
9662 9723 9736 9739 9744
SEG 0769
0769 1725 1726 1727 1728
1731
SEG16 0773
0773 1876

```

```

SEG_ASM 0660
0660 1190 1191 9634 9635
segdesc 0752
0509 0512 0752 0769 0773
1711 2310
seginit 1716
0423 1223 1255 1716
SEG_KCODE 0741
0741 1150 1725 3522 3523
9603
SEG_KCPU 0743
0743 1731 1734 3466
SEG_KDATA 0742
0742 1154 1726 1878 3463
9608
SEG_NULLASM 0654
0654 1189 9633
SEG_TSS 0746
0746 1876 1877 1880
SEG_UCODE 0744
0744 1727 2514
SEG_UDATA 0745
0745 1728 2515
setbuiltin 9025
9025 9075
SETGATE 0921
0921 3522 3523
setgid 2965
0374 2965 3879 4029 8787
9041 9932 10127
setuid 2951
0373 2951 3878 4019 8786
9035 9931 10119
setupkvm 1837
0426 1837 1859 2060 2509
6734
SHIFT 8008
8008 8036 8037 8185
skipelem 5865
5865 5914
sleep 2853
0367 2714 2853 2856 2859
3864 4065 4679 4781 5033
5036 5463 6942 6961 8485
8779 9924
spinlock 1501
0257 0367 0383 0385 0386
0387 0415 1501 1559 1562
1574 1602 1644 2407 2425
2853 3209 3219 3508 3513

```

```

4510 4527 4725 4730 4903
4938 5167 5313 6009 6013
6857 6862 8208 8221 8606
STA_R 0669 0786
0669 0786 1190 1725 1727
9634
start 1125 8708 9561
1124 1125 1167 1175 1177
4939 4964 4977 4990 5006
5088 5322 8707 8708 9560
9561 9617
startothers 1274
1208 1236 1274
stat 4254
0258 0283 0302 4254 5164
5637 6102 6209 6304 8803
9900 9917 9935 10102
stati 5637
0302 5637 6106
STA_W 0668 0785
0668 0785 1191 1726 1728
1731 9635
STA_X 0665 0782
0665 0782 1190 1725 1727
9634
sti 0563
0563 0565 1673 2764 3070
stosb 0492
0492 0494 7010 9690
stosl 0501
0501 0503 7008
strlen 7101
0396 6767 6768 7101 9029
9032 9038 9053 9085 9122
9373 9942
strncmp 7058 9003
0397 5755 7058 9003 9030
9031 9033 9037 9039 9054
9055 9059 9085
strncpy 7068
0398 5822 7068
STS_IG32 0800
0800 0927
STS_T32A 0797
0797 1876
STS_TG32 0801
0801 0927
sum 7326
7326 7328 7330 7332 7333
7345 7392

```

```

superblock 4312
0259 0287 4312 4961 5174
5178
SVR 7516
7516 7607
switchkvm 1866
0435 1254 1860 1866 2781
3094
switchvum 1873
0434 1873 1882 2566 2778
3091 6794
swtch 3158
0380 2780 2808 3093 3157
3158
syscall 3887
0406 3557 3707 3887 10106
SYSCALL 8753 8760 8761 8762 8763 87
8760 8761 8762 8763 8764
8765 8766 8767 8768 8769
8770 8771 8772 8773 8774
8775 8776 8777 8778 8779
8780 8781 8782 8783 8784
8785 8786 8787 8788
sys_chdir 6601
3779 3818 6601
SYS_chdir 3659
3659 3818 3860
sys_close 6289
3780 3830 6289
SYS_close 3671
3671 3830 3872
sys_date 4102
3801 3832 4102
SYS_date 3673
3673 3832 3874
sys_dup 6251
3781 3819 6251
SYS_dup 3660
3660 3819 3861
sys_exec 6625
3782 3816 6625
SYS_exec 3657
3657 3816 3858 8712
sys_exit 3967
3783 3811 3967
SYS_exit 3652
3652 3811 3853 8717
sys_fork 3961
3784 3810 3961
SYS_fork 3651

```

```

3651 3810 3852
sys_fstat 6301
3785 3817 6301
SYS_fstat 3658
3658 3817 3859
sys_getgid 4001
3803 3834 4001
SYS_getgid 3675
3675 3834 3876
sys_getpid 3990
3786 3820 3990
SYS_getpid 3661
3661 3820 3862
sys_getppid 4007
3804 3835 4007
SYS_getppid 3676
3676 3835 3877
sys_getprocs 4116
3807 3838 4116
SYS_getprocs 3679
3679 3838 3880
sys_getuid 3996
3802 3833 3996
SYS_getuid 3674
3674 3833 3875
SYS_halt 3672
3672 3831 3873
sys_kill 3980
3787 3815 3980
SYS_kill 3656
3656 3815 3857
sys_link 6313
3788 3828 6313
SYS_link 3669
3669 3828 3870
sys_mkdir 6551
3789 3829 6551
SYS_mkdir 3670
3670 3829 3871
sys_mknod 6567
3790 3826 6567
SYS_mknod 3667
3667 3826 3868
sys_open 6501
3791 3824 6501
SYS_open 3665
3665 3824 3866
sys_pipe 6651
3792 3813 6651
SYS_pipe 3654
3654 3813 3855
sys_read 6265
3793 3814 6265
SYS_read 3655
3655 3814 3856
sys_sbrk 4033
3794 3821 4033
SYS_sbrk 3662
3662 3821 3863
sys_setgid 4023
3806 3837 4023
SYS_setgid 3678
3678 3837 3879
sys_setuid 4013
3805 3836 4013
SYS_setuid 3677
3677 3836 3878
sys_sleep 4051
3795 3822 4051
SYS_sleep 3663
3663 3822 3864
sys_unlink 6401
3796 3827 6401
SYS_unlink 3668
3668 3827 3869
sys_uptime 4074
3799 3823 4074
SYS_uptime 3664
3664 3823 3865
sys_wait 3974
3797 3812 3974
SYS_wait 3653
3653 3812 3854
sys_write 6277
3798 3825 6277
SYS_write 3666
3666 3825 3867
taskstate 0851
0851 2309
TDCR 7540
7540 7613
T_DEV 4252
4252 5657 5707 6578
T_DIR 4250
4250 5766 5916 6328 6429
6437 6485 6525 6557 6612
T_FILE 4251
4251 6470 6514
ticks 3514
0413 3514 3567 3568 4059

```

```

4060 4065 4079
tickslock 3513
0415 3513 3525 3566 3569
4058 4062 4065 4067 4078
4080
TICR 7538
7538 7615
TIMER 7530
7530 7614
TIMER_16BIT 8571
8571 8577
TIMER_DIV 8566
8566 8578 8579
TIMER_FREQ 8565
8565 8566
timerinit 8574
0409 1235 8574
TIMER_MODE 8568
8568 8577
TIMER_RATEGEN 8570
8570 8577
TIMER_SEL0 8569
8569 8577
T_IRQ0 3379
3379 3564 3573 3577 3580
3584 3588 3589 3623 7607
7614 7627 7867 7881 7947
7966
TPR 7514
7514 7643
trap 3551
3402 3404 3472 3551 3603
3605 3608
trapframe 0602
0602 2381 2481 3551
trapret 3477
2436 2486 3476 3477
T_SYSCALL 3376
3376 3523 3553 8713 8718
8757
tvinit 3517
0414 1230 3517
uart 8615
8615 8636 8655 8665
uartgetc 8663
8663 8675
uartinit 8618
0418 1228 8618
uartintr 8673
0419 3585 8673
uartputc 8651
0420 8410 8412 8647 8651
UID_DEFAULT 2302
2302 2526
uproc 10000
0260 0372 3029 4119 9902
9933 10000 10060
userinit 2502
0368 1238 2502 2510
uva2ka 2102
0427 2102 2126
V2P 0217
0217 1830 1831
V2P_WO 0220
0220 1036 1046
VER 7513
7513 7623
wait 2678
0369 2678 3854 3976 8762
8833 8949 8973 8974 9133
9824 9907
waitdisk 9701
9701 9713 9722
wakeup 2916
0370 2916 3568 4622 4841
5066 5076 5492 5522 6916
6919 6941 6946 6968 8457
wakeup1 2903
2438 2658 2665 2903 2919
walkpgdir 1754
1754 1787 1926 1992 2033
2063 2106
write_head 5004
5004 5023 5105 5108
writei 5702
0303 5702 5824 6176 6435
6436
write_log 5083
5083 5104
xchg 0569
0569 1266 1583 1619
YEAR 7734
7734 7757
yield 2814
0371 2814 3624

```



```
0100 typedef unsigned int    uint;
0101 typedef unsigned short ushort;
0102 typedef unsigned char  uchar;
0103 typedef uint pde_t;
0104
0105
0106
0107
0108
0109
0110
0111
0112
0113
0114
0115
0116
0117
0118
0119
0120
0121
0122
0123
0124
0125
0126
0127
0128
0129
0130
0131
0132
0133
0134
0135
0136
0137
0138
0139
0140
0141
0142
0143
0144
0145
0146
0147
0148
0149
```

```
0150 #define NPROC          64 // maximum number of processes
0151 #define KSTACKSIZE 4096 // size of per-process kernel stack
0152 #define NCPU           8 // maximum number of CPUs
0153 #define NOFILE         16 // open files per process
0154 #define NFILE          100 // open files per system
0155 #define NINODE          50 // maximum number of active i-nodes
0156 #define NDEV           10 // maximum major device number
0157 #define ROOTDEV         1 // device number of file system root disk
0158 #define MAXARG          32 // max exec arguments
0159 #define MAXOPBLOCKS    10 // max # of blocks any FS op writes
0160 #define LOGSIZE         (MAXOPBLOCKS*3) // max data blocks in on-disk log
0161 #define NBUF            (MAXOPBLOCKS*3) // size of disk block cache
0162 #define FSSIZE          1000 // size of file system in blocks
0163
0164
0165
0166
0167
0168
0169
0170
0171
0172
0173
0174
0175
0176
0177
0178
0179
0180
0181
0182
0183
0184
0185
0186
0187
0188
0189
0190
0191
0192
0193
0194
0195
0196
0197
0198
0199
```

```

0200 // Memory layout
0201
0202 #define EXTMEM 0x100000          // Start of extended memory
0203 #define PHYSTOP 0xE000000      // Top physical memory
0204 #define DEVSPACE 0xFE000000    // Other devices are at high addresses
0205
0206 // Key addresses for address space layout (see kmap in vm.c for layout)
0207 #define KERNBASE 0x80000000     // First kernel virtual address
0208 #define KERNLINK (KERNBASE+EXTMEM) // Address where kernel is linked
0209
0210 #ifndef __ASSEMBLER__
0211
0212 static inline uint v2p(void *a) { return ((uint) (a)) - KERNBASE; }
0213 static inline void *p2v(uint a) { return (void *) ((a) + KERNBASE); }
0214
0215 #endif
0216
0217 #define V2P(a) (((uint) (a)) - KERNBASE)
0218 #define P2V(a) (((void *) (a)) + KERNBASE)
0219
0220 #define V2P_WO(x) ((x) - KERNBASE) // same as V2P, but without casts
0221 #define P2V_WO(x) ((x) + KERNBASE) // same as P2V, but without casts
0222
0223
0224
0225
0226
0227
0228
0229
0230
0231
0232
0233
0234
0235
0236
0237
0238
0239
0240
0241
0242
0243
0244
0245
0246
0247
0248
0249

```

```

0250 struct buf;
0251 struct context;
0252 struct file;
0253 struct inode;
0254 struct pipe;
0255 struct proc;
0256 struct rtcdate;
0257 struct spinlock;
0258 struct stat;
0259 struct superblock;
0260 struct uproc;
0261
0262 // bio.c
0263 void          binit(void);
0264 struct buf*   bread(uint, uint);
0265 void          brelse(struct buf*);
0266 void          bwrite(struct buf*);
0267
0268 // console.c
0269 void          consoleinit(void);
0270 void          cprintf(char*, ...);
0271 void          consoleintr(int (*)(void));
0272 void          panic(char*) __attribute__((noreturn));
0273
0274 // exec.c
0275 int           exec(char*, char**);
0276
0277 // file.c
0278 struct file*  filealloc(void);
0279 void          fileclose(struct file*);
0280 struct file*  filedup(struct file*);
0281 void          fileinit(void);
0282 int           fileread(struct file*, char*, int n);
0283 int           filestat(struct file*, struct stat*);
0284 int           filewrite(struct file*, char*, int n);
0285
0286 // fs.c
0287 void          readsb(int dev, struct superblock *sb);
0288 int           dirlink(struct inode*, char*, uint);
0289 struct inode* dirlookup(struct inode*, char*, uint*);
0290 struct inode* ialloc(uint, short);
0291 struct inode* idup(struct inode*);
0292 void          iinit(int dev);
0293 void          ilock(struct inode*);
0294 void          iput(struct inode*);
0295 void          iunlock(struct inode*);
0296 void          iunlockput(struct inode*);
0297 void          iupdate(struct inode*);
0298 int           namecmp(const char*, const char*);
0299 struct inode* namei(char*);

```

```

0300 struct inode*   nameiparent(char*, char*);
0301 int              readi(struct inode*, char*, uint, uint);
0302 void             stati(struct inode*, struct stat*);
0303 int              writei(struct inode*, char*, uint, uint);
0304
0305 // ide.c
0306 void             ideinit(void);
0307 void             ideintr(void);
0308 void             iderw(struct buf*);
0309
0310 // ioapic.c
0311 void             ioapicenable(int irq, int cpu);
0312 extern uchar     ioapicid;
0313 void             ioapicinit(void);
0314
0315 // kalloc.c
0316 char*           kalloc(void);
0317 void             kfree(char*);
0318 void             kinit1(void*, void*);
0319 void             kinit2(void*, void*);
0320
0321 // kbd.c
0322 void             kbdintr(void);
0323
0324 // lapic.c
0325 void             cmostime(struct rtcdate *r);
0326 int              cpunum(void);
0327 extern volatile uint* lapic;
0328 void             lapiceoi(void);
0329 void             lapicinit(void);
0330 void             lapicstartap(uchar, uint);
0331 void             microdelay(int);
0332
0333 // log.c
0334 void             initlog(int dev);
0335 void             log_write(struct buf*);
0336 void             begin_op();
0337 void             end_op();
0338
0339 // mp.c
0340 extern int        ismp;
0341 int              mpbcpu(void);
0342 void             mpinit(void);
0343 void             mpstartthem(void);
0344
0345 // picirq.c
0346 void             picenable(int);
0347 void             picinit(void);
0348
0349

```

```

0350 // pipe.c
0351 int              pipealloc(struct file**, struct file**);
0352 void             pipeclose(struct pipe*, int);
0353 int              piperead(struct pipe*, char*, int);
0354 int              pipewrite(struct pipe*, char*, int);
0355
0356
0357 // proc.c
0358 struct proc*     copyproc(struct proc*);
0359 void             exit(void);
0360 int              fork(void);
0361 int              growproc(int);
0362 int              kill(int);
0363 void             pinit(void);
0364 void             procdump(void);
0365 void             scheduler(void) __attribute__((noreturn));
0366 void             sched(void);
0367 void             sleep(void*, struct spinlock*);
0368 void             userinit(void);
0369 int              wait(void);
0370 void             wakeup(void*);
0371 void             yield(void);
0372 int              getprocs(int, struct uproc*);
0373 int              setuid(int);
0374 int              setgid(int);
0375 int              getuid();
0376 int              getgid();
0377 int              addtoq();
0378 int              putinQ(struct proc *);
0379 // swtch.S
0380 void             swtch(struct context**, struct context*);
0381
0382 // spinlock.c
0383 void             acquire(struct spinlock*);
0384 void             getcallerpcs(void*, uint*);
0385 int              holding(struct spinlock*);
0386 void             initlock(struct spinlock*, char*);
0387 void             release(struct spinlock*);
0388 void             pushcli(void);
0389 void             popcli(void);
0390
0391 // string.c
0392 int              memcmp(const void*, const void*, uint);
0393 void*            memmove(void*, const void*, uint);
0394 void*            memset(void*, int, uint);
0395 char*            safestrcpy(char*, const char*, int);
0396 int              strlen(const char*);
0397 int              strncmp(const char*, const char*, uint);
0398 char*            strncpy(char*, const char*, int);
0399

```

```

0400 // syscall.c
0401 int      argint(int, int*);
0402 int      argptr(int, char**, int);
0403 int      argstr(int, char**);
0404 int      fetchint(uint, int*);
0405 int      fetchstr(uint, char**);
0406 void      syscall(void);
0407
0408 // timer.c
0409 void      timerinit(void);
0410
0411 // trap.c
0412 void      idtinit(void);
0413 extern uint ticks;
0414 void      tvinit(void);
0415 extern struct spinlock tickslock;
0416
0417 // uart.c
0418 void      uartinit(void);
0419 void      uartintr(void);
0420 void      uartputc(int);
0421
0422 // vm.c
0423 void      seginit(void);
0424 void      kvmalloc(void);
0425 void      vmenable(void);
0426 pde_t*    setupkvm(void);
0427 char*     uva2ka(pde_t*, char*);
0428 int      allocvm(pde_t*, uint, uint);
0429 int      deallocvm(pde_t*, uint, uint);
0430 void      freevm(pde_t*);
0431 void      inituvm(pde_t*, char*, uint);
0432 int      loaduvm(pde_t*, char*, struct inode*, uint, uint);
0433 pde_t*    copyuvm(pde_t*, uint);
0434 void      switchuvm(struct proc*);
0435 void      switchkvm(void);
0436 int      copyout(pde_t*, uint, void*, uint);
0437 void      clearpteu(pde_t *pgdir, char *uva);
0438
0439 // number of elements in fixed-size array
0440 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
0441
0442
0443
0444
0445
0446
0447
0448
0449

```

```

0450 // Routines to let C code use special x86 instructions.
0451
0452 static inline uchar
0453 inb(ushort port)
0454 {
0455     uchar data;
0456
0457     asm volatile("in %1,%0" : "=a" (data) : "d" (port));
0458     return data;
0459 }
0460
0461 static inline void
0462 insl(int port, void *addr, int cnt)
0463 {
0464     asm volatile("cld; rep insl" :
0465         "=D" (addr), "=c" (cnt) :
0466         "d" (port), "0" (addr), "1" (cnt) :
0467         "memory", "cc");
0468 }
0469
0470 static inline void
0471 outb(ushort port, uchar data)
0472 {
0473     asm volatile("out %0,%1" : : "a" (data), "d" (port));
0474 }
0475
0476 static inline void
0477 outw(ushort port, ushort data)
0478 {
0479     asm volatile("out %0,%1" : : "a" (data), "d" (port));
0480 }
0481
0482 static inline void
0483 outsl(int port, const void *addr, int cnt)
0484 {
0485     asm volatile("cld; rep outsl" :
0486         "=S" (addr), "=c" (cnt) :
0487         "d" (port), "0" (addr), "1" (cnt) :
0488         "cc");
0489 }
0490
0491 static inline void
0492 stosb(void *addr, int data, int cnt)
0493 {
0494     asm volatile("cld; rep stosb" :
0495         "=D" (addr), "=c" (cnt) :
0496         "0" (addr), "1" (cnt), "a" (data) :
0497         "memory", "cc");
0498 }
0499

```

```

0500 static inline void
0501 stosl(void *addr, int data, int cnt)
0502 {
0503     asm volatile("cld; rep stosl" :
0504                 "=D" (addr), "=c" (cnt) :
0505                 "0" (addr), "1" (cnt), "a" (data) :
0506                 "memory", "cc");
0507 }
0508
0509 struct segdesc;
0510
0511 static inline void
0512 lgdt(struct segdesc *p, int size)
0513 {
0514     volatile ushort pd[3];
0515
0516     pd[0] = size-1;
0517     pd[1] = (uint)p;
0518     pd[2] = (uint)p >> 16;
0519
0520     asm volatile("lgdt (%0)" : : "r" (pd));
0521 }
0522
0523 struct gatedesc;
0524
0525 static inline void
0526 lidt(struct gatedesc *p, int size)
0527 {
0528     volatile ushort pd[3];
0529
0530     pd[0] = size-1;
0531     pd[1] = (uint)p;
0532     pd[2] = (uint)p >> 16;
0533
0534     asm volatile("lidt (%0)" : : "r" (pd));
0535 }
0536
0537 static inline void
0538 ltr(ushort sel)
0539 {
0540     asm volatile("ltr %0" : : "r" (sel));
0541 }
0542
0543 static inline uint
0544 readeflags(void)
0545 {
0546     uint eflags;
0547     asm volatile("pushfl; popl %0" : "=r" (eflags));
0548     return eflags;
0549 }

```

```

0550 static inline void
0551 loadgs(ushort v)
0552 {
0553     asm volatile("movw %0, %%gs" : : "r" (v));
0554 }
0555
0556 static inline void
0557 cli(void)
0558 {
0559     asm volatile("cli");
0560 }
0561
0562 static inline void
0563 sti(void)
0564 {
0565     asm volatile("sti");
0566 }
0567
0568 static inline uint
0569 xchg(volatile uint *addr, uint newval)
0570 {
0571     uint result;
0572
0573     // The + in "+m" denotes a read-modify-write operand.
0574     asm volatile("lock; xchgl %0, %1" :
0575                 "+m" (*addr), "=a" (result) :
0576                 "1" (newval) :
0577                 "cc");
0578     return result;
0579 }
0580
0581 static inline uint
0582 rcr2(void)
0583 {
0584     uint val;
0585     asm volatile("movl %%cr2,%0" : "=r" (val));
0586     return val;
0587 }
0588
0589 static inline void
0590 lcr3(uint val)
0591 {
0592     asm volatile("movl %0,%%cr3" : : "r" (val));
0593 }
0594
0595
0596
0597
0598
0599

```

```

0600 // Layout of the trap frame built on the stack by the
0601 // hardware and by trapasm.S, and passed to trap().
0602 struct trapframe {
0603     // registers as pushed by pusha
0604     uint edi;
0605     uint esi;
0606     uint ebp;
0607     uint oesp;      // useless & ignored
0608     uint ebx;
0609     uint edx;
0610     uint ecx;
0611     uint eax;
0612
0613     // rest of trap frame
0614     ushort gs;
0615     ushort padding1;
0616     ushort fs;
0617     ushort padding2;
0618     ushort es;
0619     ushort padding3;
0620     ushort ds;
0621     ushort padding4;
0622     uint trapno;
0623
0624     // below here defined by x86 hardware
0625     uint err;
0626     uint eip;
0627     ushort cs;
0628     ushort padding5;
0629     uint eflags;
0630
0631     // below here only when crossing rings, such as from user to kernel
0632     uint esp;
0633     ushort ss;
0634     ushort padding6;
0635 };
0636
0637
0638
0639
0640
0641
0642
0643
0644
0645
0646
0647
0648
0649

```

```

0650 //
0651 // assembler macros to create x86 segments
0652 //
0653
0654 #define SEG_NULLASM                                     \
0655     .word 0, 0;                                         \
0656     .byte 0, 0, 0, 0
0657
0658 // The 0xC0 means the limit is in 4096-byte units
0659 // and (for executable segments) 32-bit mode.
0660 #define SEG_ASM(type,base,lim)                         \
0661     .word (((lim) >> 12) & 0xffff), ((base) & 0xffff); \
0662     .byte (((base) >> 16) & 0xff), (0x90 | (type)),    \
0663         (0xC0 | (((lim) >> 28) & 0xf)), (((base) >> 24) & 0xff)
0664
0665 #define STA_X      0x8      // Executable segment
0666 #define STA_E      0x4      // Expand down (non-executable segments)
0667 #define STA_C      0x4      // Conforming code segment (executable only)
0668 #define STA_W      0x2      // Writeable (non-executable segments)
0669 #define STA_R      0x2      // Readable (executable segments)
0670 #define STA_A      0x1      // Accessed
0671
0672
0673
0674
0675
0676
0677
0678
0679
0680
0681
0682
0683
0684
0685
0686
0687
0688
0689
0690
0691
0692
0693
0694
0695
0696
0697
0698
0699

```

```

0700 // This file contains definitions for the
0701 // x86 memory management unit (MMU).
0702
0703 // Eflags register
0704 #define FL_CF      0x00000001    // Carry Flag
0705 #define FL_PF      0x00000004    // Parity Flag
0706 #define FL_AF      0x00000010    // Auxiliary carry Flag
0707 #define FL_ZF      0x00000040    // Zero Flag
0708 #define FL_SF      0x00000080    // Sign Flag
0709 #define FL_TF      0x00000100    // Trap Flag
0710 #define FL_IF      0x00000200    // Interrupt Enable
0711 #define FL_DF      0x00000400    // Direction Flag
0712 #define FL_OF      0x00000800    // Overflow Flag
0713 #define FL_IOPL_MASK 0x00003000 // I/O Privilege Level bitmask
0714 #define FL_IOPL_0    0x00000000 // IOPL == 0
0715 #define FL_IOPL_1    0x00001000 // IOPL == 1
0716 #define FL_IOPL_2    0x00002000 // IOPL == 2
0717 #define FL_IOPL_3    0x00003000 // IOPL == 3
0718 #define FL_NT      0x00004000    // Nested Task
0719 #define FL_RF      0x00010000    // Resume Flag
0720 #define FL_VM      0x00020000    // Virtual 8086 mode
0721 #define FL_AC      0x00040000    // Alignment Check
0722 #define FL_VIF      0x00080000    // Virtual Interrupt Flag
0723 #define FL_VIP      0x00100000    // Virtual Interrupt Pending
0724 #define FL_ID      0x00200000    // ID flag
0725
0726 // Control Register flags
0727 #define CR0_PE      0x00000001    // Protection Enable
0728 #define CR0_MP      0x00000002    // Monitor coProcessor
0729 #define CR0_EM      0x00000004    // Emulation
0730 #define CR0_TS      0x00000008    // Task Switched
0731 #define CR0_ET      0x00000010    // Extension Type
0732 #define CR0_NE      0x00000020    // Numeric Error
0733 #define CR0_WP      0x00010000    // Write Protect
0734 #define CR0_AM      0x00040000    // Alignment Mask
0735 #define CR0_NW      0x02000000    // Not Writethrough
0736 #define CR0_CD      0x40000000    // Cache Disable
0737 #define CR0_PG      0x80000000    // Paging
0738
0739 #define CR4_PSE      0x00000010    // Page size extension
0740
0741 #define SEG_KCODE 1 // kernel code
0742 #define SEG_KDATA 2 // kernel data+stack
0743 #define SEG_KCPU 3 // kernel per-cpu data
0744 #define SEG_UCODE 4 // user code
0745 #define SEG_UDATA 5 // user data+stack
0746 #define SEG_TSS 6 // this process's task state
0747
0748
0749

```

```

0750 #ifndef __ASSEMBLER__
0751 // Segment Descriptor
0752 struct segdesc {
0753     uint lim_15_0 : 16; // Low bits of segment limit
0754     uint base_15_0 : 16; // Low bits of segment base address
0755     uint base_23_16 : 8; // Middle bits of segment base address
0756     uint type : 4; // Segment type (see STS_constants)
0757     uint s : 1; // 0 = system, 1 = application
0758     uint dpl : 2; // Descriptor Privilege Level
0759     uint p : 1; // Present
0760     uint lim_19_16 : 4; // High bits of segment limit
0761     uint avl : 1; // Unused (available for software use)
0762     uint rsv1 : 1; // Reserved
0763     uint db : 1; // 0 = 16-bit segment, 1 = 32-bit segment
0764     uint g : 1; // Granularity: limit scaled by 4K when set
0765     uint base_31_24 : 8; // High bits of segment base address
0766 };
0767
0768 // Normal segment
0769 #define SEG(type, base, lim, dpl) (struct segdesc) \
0770 { ((lim) >> 12) & 0xffff, (uint)(base) & 0xffff, \
0771   ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
0772   (uint)(lim) >> 28, 0, 0, 1, 1, (uint)(base) >> 24 }
0773 #define SEG16(type, base, lim, dpl) (struct segdesc) \
0774 { (lim) & 0xffff, (uint)(base) & 0xffff, \
0775   ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
0776   (uint)(lim) >> 16, 0, 0, 1, 0, (uint)(base) >> 24 }
0777 #endif
0778
0779 #define DPL_USER 0x3 // User DPL
0780
0781 // Application segment type bits
0782 #define STA_X 0x8 // Executable segment
0783 #define STA_E 0x4 // Expand down (non-executable segments)
0784 #define STA_C 0x4 // Conforming code segment (executable only)
0785 #define STA_W 0x2 // Writeable (non-executable segments)
0786 #define STA_R 0x2 // Readable (executable segments)
0787 #define STA_A 0x1 // Accessed
0788
0789 // System segment type bits
0790 #define STS_T16A 0x1 // Available 16-bit TSS
0791 #define STS_LDT 0x2 // Local Descriptor Table
0792 #define STS_T16B 0x3 // Busy 16-bit TSS
0793 #define STS_CG16 0x4 // 16-bit Call Gate
0794 #define STS_TG 0x5 // Task Gate / Coum Transmissions
0795 #define STS_IG16 0x6 // 16-bit Interrupt Gate
0796 #define STS_TG16 0x7 // 16-bit Trap Gate
0797 #define STS_T32A 0x9 // Available 32-bit TSS
0798 #define STS_T32B 0xB // Busy 32-bit TSS
0799 #define STS_CG32 0xC // 32-bit Call Gate

```

```

0800 #define STS_IG32    0xE    // 32-bit Interrupt Gate
0801 #define STS_TG32    0xF    // 32-bit Trap Gate
0802
0803 // A virtual address 'la' has a three-part structure as follows:
0804 //
0805 // +-----10-----+-----10-----+-----12-----+
0806 // | Page Directory | Page Table | Offset within Page |
0807 // |      Index      |      Index      |                |
0808 // +-----+-----+-----+
0809 // \--- PDX(va) --/ \--- PTX(va) --/
0810
0811 // page directory index
0812 #define PDX(va)      (((uint)(va) >> PDXSHIFT) & 0x3FF)
0813
0814 // page table index
0815 #define PTX(va)      (((uint)(va) >> PTXSHIFT) & 0x3FF)
0816
0817 // construct virtual address from indexes and offset
0818 #define PGADDR(d, t, o) ((uint)((d) << PDXSHIFT | (t) << PTXSHIFT | (o)))
0819
0820 // Page directory and page table constants.
0821 #define NPENTRIES    1024    // # directory entries per page directory
0822 #define NPTENTRIES    1024    // # PTEs per page table
0823 #define PGSIZE        4096    // bytes mapped by a page
0824
0825 #define PGSHIFT        12    // log2(PGSIZE)
0826 #define PTXSHIFT        12    // offset of PTX in a linear address
0827 #define PDXSHIFT        22    // offset of PDX in a linear address
0828
0829 #define PGROUNDUP(sz)  (((sz)+PGSIZE-1) & ~(PGSIZE-1))
0830 #define PGROUNDDOWN(a) (((a)) & ~(PGSIZE-1))
0831
0832 // Page table/directory entry flags.
0833 #define PTE_P          0x001    // Present
0834 #define PTE_W          0x002    // Writeable
0835 #define PTE_U          0x004    // User
0836 #define PTE_PWT        0x008    // Write-Through
0837 #define PTE_PCD        0x010    // Cache-Disable
0838 #define PTE_A          0x020    // Accessed
0839 #define PTE_D          0x040    // Dirty
0840 #define PTE_PS         0x080    // Page Size
0841 #define PTE_MBZ        0x180    // Bits must be zero
0842
0843 // Address in page table or page directory entry
0844 #define PTE_ADDR(pte)  ((uint)(pte) & ~0xFFF)
0845 #define PTE_FLAGS(pte) ((uint)(pte) & 0xFFF)
0846
0847 #ifndef __ASSEMBLER__
0848 typedef uint pte_t;
0849

```

```

0850 // Task state segment format
0851 struct taskstate {
0852     uint link;        // Old ts selector
0853     uint esp0;        // Stack pointers and segment selectors
0854     ushort ss0;       // after an increase in privilege level
0855     ushort padding1;
0856     uint *esp1;
0857     ushort ssl;
0858     ushort padding2;
0859     uint *esp2;
0860     ushort ss2;
0861     ushort padding3;
0862     void *cr3;        // Page directory base
0863     uint *eip;        // Saved state from last task switch
0864     uint eflags;
0865     uint eax;         // More saved state (registers)
0866     uint ecx;
0867     uint edx;
0868     uint ebx;
0869     uint *esp;
0870     uint *ebp;
0871     uint esi;
0872     uint edi;
0873     ushort es;        // Even more saved state (segment selectors)
0874     ushort padding4;
0875     ushort cs;
0876     ushort padding5;
0877     ushort ss;
0878     ushort padding6;
0879     ushort ds;
0880     ushort padding7;
0881     ushort fs;
0882     ushort padding8;
0883     ushort gs;
0884     ushort padding9;
0885     ushort ldt;
0886     ushort padding10;
0887     ushort t;         // Trap on task switch
0888     ushort iomb;      // I/O map base address
0889 };
0890
0891
0892
0893
0894
0895
0896
0897
0898
0899

```



```

0900 // Gate descriptors for interrupts and traps
0901 struct gatedesc {
0902     uint off_15_0 : 16;    // low 16 bits of offset in segment
0903     uint cs : 16;           // code segment selector
0904     uint args : 5;          // # args, 0 for interrupt/trap gates
0905     uint rsv1 : 3;          // reserved(should be zero I guess)
0906     uint type : 4;          // type(STS_{TG,IG32,TG32})
0907     uint s : 1;            // must be 0 (system)
0908     uint dpl : 2;          // descriptor(meaning new) privilege level
0909     uint p : 1;            // Present
0910     uint off_31_16 : 16;    // high bits of offset in segment
0911 };
0912
0913 // Set up a normal interrupt/trap gate descriptor.
0914 // - istrap: 1 for a trap (= exception) gate, 0 for an interrupt gate.
0915 // - interrupt gate clears FL_IF, trap gate leaves FL_IF alone
0916 // - sel: Code segment selector for interrupt/trap handler
0917 // - off: Offset in code segment for interrupt/trap handler
0918 // - dpl: Descriptor Privilege Level -
0919 //       the privilege level required for software to invoke
0920 //       this interrupt/trap gate explicitly using an int instruction.
0921 #define SETGATE(gate, istrap, sel, off, d) \
0922 { \
0923     (gate).off_15_0 = (uint)(off) & 0xffff; \
0924     (gate).cs = (sel); \
0925     (gate).args = 0; \
0926     (gate).rsv1 = 0; \
0927     (gate).type = (istrap) ? STS_TG32 : STS_IG32; \
0928     (gate).s = 0; \
0929     (gate).dpl = (d); \
0930     (gate).p = 1; \
0931     (gate).off_31_16 = (uint)(off) >> 16; \
0932 }
0933
0934 #endif
0935
0936
0937
0938
0939
0940
0941
0942
0943
0944
0945
0946
0947
0948
0949

```

```

0950 // Format of an ELF executable file
0951
0952 #define ELF_MAGIC 0x464C457FU // "\x7FELF" in little endian
0953
0954 // File header
0955 struct elfhdr {
0956     uint magic; // must equal ELF_MAGIC
0957     uchar elf[12];
0958     ushort type;
0959     ushort machine;
0960     uint version;
0961     uint entry;
0962     uint phoff;
0963     uint shoff;
0964     uint flags;
0965     ushort ehsize;
0966     ushort phentsize;
0967     ushort phnum;
0968     ushort shentsize;
0969     ushort shnum;
0970     ushort shstrndx;
0971 };
0972
0973 // Program section header
0974 struct proghdr {
0975     uint type;
0976     uint off;
0977     uint vaddr;
0978     uint paddr;
0979     uint filesz;
0980     uint memsz;
0981     uint flags;
0982     uint align;
0983 };
0984
0985 // Values for Proghdr type
0986 #define ELF_PROG_LOAD 1
0987
0988 // Flag bits for Proghdr flags
0989 #define ELF_PROG_FLAG_EXEC 1
0990 #define ELF_PROG_FLAG_WRITE 2
0991 #define ELF_PROG_FLAG_READ 4
0992
0993
0994
0995
0996
0997
0998
0999

```

```

1000 # Multiboot header, for multiboot boot loaders like GNU Grub.
1001 # http://www.gnu.org/software/grub/manual/multiboot/multiboot.html
1002 #
1003 # Using GRUB 2, you can boot xv6 from a file stored in a
1004 # Linux file system by copying kernel or kernelmemfs to /boot
1005 # and then adding this menu entry:
1006 #
1007 # menuentry "xv6" {
1008 #   insmod ext2
1009 #   set root='(hd0,msdos1)'
1010 #   set kernel='/boot/kernel'
1011 #   echo "Loading ${kernel}..."
1012 #   multiboot ${kernel} ${kernel}
1013 #   boot
1014 # }
1015
1016 #include "asm.h"
1017 #include "memlayout.h"
1018 #include "mmu.h"
1019 #include "param.h"
1020
1021 # Multiboot header. Data to direct multiboot loader.
1022 .p2align 2
1023 .text
1024 .globl multiboot_header
1025 multiboot_header:
1026     #define magic 0x1badb002
1027     #define flags 0
1028     .long magic
1029     .long flags
1030     .long (-magic-flags)
1031
1032 # By convention, the _start symbol specifies the ELF entry point.
1033 # Since we haven't set up virtual memory yet, our entry point is
1034 # the physical address of 'entry'.
1035 .globl _start
1036 _start = V2P_WO(entry)
1037
1038 # Entering xv6 on boot processor, with paging off.
1039 .globl entry
1040 entry:
1041     # Turn on page size extension for 4Mbyte pages
1042     movl    %cr4, %eax
1043     orl     $(CR4_PSE), %eax
1044     movl    %eax, %cr4
1045     # Set page directory
1046     movl    $(V2P_WO(entrypgdir)), %eax
1047     movl    %eax, %cr3
1048     # Turn on paging.
1049     movl    %cr0, %eax

```

```

1050     orl     $(CR0_PG|CR0_WP), %eax
1051     movl    %eax, %cr0
1052
1053     # Set up the stack pointer.
1054     movl    $(stack + KSTACKSIZE), %esp
1055
1056     # Jump to main(), and switch to executing at
1057     # high addresses. The indirect call is needed because
1058     # the assembler produces a PC-relative instruction
1059     # for a direct jump.
1060     mov     $main, %eax
1061     jmp     *%eax
1062
1063 .comm stack, KSTACKSIZE
1064
1065
1066
1067
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079
1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099

```

```

1100 #include "asm.h"
1101 #include "memlayout.h"
1102 #include "mmu.h"
1103
1104 # Each non-boot CPU ("AP") is started up in response to a STARTUP
1105 # IPI from the boot CPU. Section B.4.2 of the Multi-Processor
1106 # Specification says that the AP will start in real mode with CS:IP
1107 # set to XY00:0000, where XY is an 8-bit value sent with the
1108 # STARTUP. Thus this code must start at a 4096-byte boundary.
1109 #
1110 # Because this code sets DS to zero, it must sit
1111 # at an address in the low 2^16 bytes.
1112 #
1113 # Startothers (in main.c) sends the STARTUPs one at a time.
1114 # It copies this code (start) at 0x7000. It puts the address of
1115 # a newly allocated per-core stack in start-4, the address of the
1116 # place to jump to (mpenter) in start-8, and the physical address
1117 # of entrypgdir in start-12.
1118 #
1119 # This code is identical to bootasm.S except:
1120 #   - it does not need to enable A20
1121 #   - it uses the address at start-4, start-8, and start-12
1122
1123 .code16
1124 .globl start
1125 start:
1126     cli
1127
1128     xorw    %ax,%ax
1129     movw    %ax,%ds
1130     movw    %ax,%es
1131     movw    %ax,%ss
1132
1133     lgdt    gdtdesc
1134     movl    %cr0,%eax
1135     orl     $CR0_PE, %eax
1136     movl    %eax,%cr0
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149

```

```

1150     ljmp    $(SEG_KCODE<<3), $(start32)
1151
1152 .code32
1153 start32:
1154     movw    $(SEG_KDATA<<3), %ax
1155     movw    %ax,%ds
1156     movw    %ax,%es
1157     movw    %ax,%ss
1158     movw    $0,%ax
1159     movw    %ax,%fs
1160     movw    %ax,%gs
1161
1162     # Turn on page size extension for 4Mbyte pages
1163     movl    %cr4,%eax
1164     orl     $(CR4_PSE), %eax
1165     movl    %eax,%cr4
1166     # Use enterpgdir as our initial page table
1167     movl    (start-12), %eax
1168     movl    %eax,%cr3
1169     # Turn on paging.
1170     movl    %cr0,%eax
1171     orl     $(CR0_PE|CR0_PG|CR0_WP), %eax
1172     movl    %eax,%cr0
1173
1174     # Switch to the stack allocated by startothers()
1175     movl    (start-4), %esp
1176     # Call mpenter()
1177     call    *(start-8)
1178
1179     movw    $0x8a00, %ax
1180     movw    %ax,%dx
1181     outw    %ax,%dx
1182     movw    $0x8ae0, %ax
1183     outw    %ax,%dx
1184 spin:
1185     jmp     spin
1186
1187 .p2align 2
1188 gdt:
1189     SEG_NULLASM
1190     SEG_ASM(STA_X|STA_R, 0, 0xffffffff)
1191     SEG_ASM(STA_W, 0, 0xffffffff)
1192
1193
1194 gdtdesc:
1195     .word   (gdtdesc - gdt - 1)
1196     .long   gdt
1197
1198
1199

```

```

1200 #include "types.h"
1201 #include "defs.h"
1202 #include "param.h"
1203 #include "memlayout.h"
1204 #include "mmu.h"
1205 #include "proc.h"
1206 #include "x86.h"
1207
1208 static void startothers(void);
1209 static void mpmain(void) __attribute__((noreturn));
1210 extern pde_t *kpgdir;
1211 extern char end[]; // first address after kernel loaded from ELF file
1212
1213 // Bootstrap processor starts running C code here.
1214 // Allocate a real stack and switch to it, first
1215 // doing some setup required for memory allocator to work.
1216 int
1217 main(void)
1218 {
1219     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1220     kvmalloc(); // kernel page table
1221     mpinit(); // collect info about this machine
1222     lapicinit();
1223     seginit(); // set up segments
1224     cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
1225     picinit(); // interrupt controller
1226     ioapicinit(); // another interrupt controller
1227     consoleinit(); // I/O devices & their interrupts
1228     uartinit(); // serial port
1229     pinit(); // process table
1230     tvinit(); // trap vectors
1231     binit(); // buffer cache
1232     fileinit(); // file table
1233     ideinit(); // disk
1234     if(!ismp)
1235         timerinit(); // uniprocessor timer
1236     startothers(); // start other processors
1237     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
1238     userinit(); // first user process
1239     // Finish setting up this processor in mpmain.
1240     mpmain();
1241 }
1242
1243
1244
1245
1246
1247
1248
1249

```

```

1250 // Other CPUs jump here from entryother.S.
1251 static void
1252 mpenter(void)
1253 {
1254     switchkvm();
1255     seginit();
1256     lapicinit();
1257     mpmain();
1258 }
1259
1260 // Common CPU setup code.
1261 static void
1262 mpmain(void)
1263 {
1264     cprintf("cpu%d: starting\n", cpu->id);
1265     idtinit(); // load idt register
1266     xchg(&cpu->started, 1); // tell startothers() we're up
1267     scheduler(); // start running processes
1268 }
1269
1270 pde_t entrypgdir[]; // For entry.S
1271
1272 // Start the non-boot (AP) processors.
1273 static void
1274 startothers(void)
1275 {
1276     extern uchar _binary_entryother_start[], _binary_entryother_size[];
1277     uchar *code;
1278     struct cpu *c;
1279     char *stack;
1280
1281     // Write entry code to unused memory at 0x7000.
1282     // The linker has placed the image of entryother.S in
1283     // _binary_entryother_start.
1284     code = p2v(0x7000);
1285     memmove(code, _binary_entryother_start, (uint)_binary_entryother_size);
1286
1287     for(c = cpus; c < cpus+ncpu; c++){
1288         if(c == cpus+cpunum()) // We've started already.
1289             continue;
1290
1291         // Tell entryother.S what stack to use, where to enter, and what
1292         // pgdir to use. We cannot use kpgdir yet, because the AP processor
1293         // is running in low memory, so we use entrypgdir for the APs too.
1294         stack = kalloc();
1295         *(void**)(code-4) = stack + KSTACKSIZE;
1296         *(void**)(code-8) = mpenter;
1297         *(int**)(code-12) = (void *) v2p(entrypgdir);
1298
1299         lapicstartap(c->id, v2p(code));

```

```

1300 // wait for cpu to finish mpmain()
1301 while(c->started == 0)
1302     ;
1303 }
1304 }
1305
1306 // Boot page table used in entry.S and entryother.S.
1307 // Page directories (and page tables), must start on a page boundary,
1308 // hence the "__aligned__" attribute.
1309 // Use PTE_PS in page directory entry to enable 4Mbyte pages.
1310 __attribute__((__aligned__(PGSIZE)))
1311 pde_t entrypgdir[NPDENTRIES] = {
1312     // Map VA's [0, 4MB) to PA's [0, 4MB)
1313     [0] = (0) | PTE_P | PTE_W | PTE_PS,
1314     // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
1315     [KERNBASE >> PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
1316 };
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349

```

```

1350 // Blank page.
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399

```

1400 // Blank page.

1401  
1402  
1403  
1404  
1405  
1406  
1407  
1408  
1409  
1410  
1411  
1412  
1413  
1414  
1415  
1416  
1417  
1418  
1419  
1420  
1421  
1422  
1423  
1424  
1425  
1426  
1427  
1428  
1429  
1430  
1431  
1432  
1433  
1434  
1435  
1436  
1437  
1438  
1439  
1440  
1441  
1442  
1443  
1444  
1445  
1446  
1447  
1448  
1449

1450 // Blank page.

1451  
1452  
1453  
1454  
1455  
1456  
1457  
1458  
1459  
1460  
1461  
1462  
1463  
1464  
1465  
1466  
1467  
1468  
1469  
1470  
1471  
1472  
1473  
1474  
1475  
1476  
1477  
1478  
1479  
1480  
1481  
1482  
1483  
1484  
1485  
1486  
1487  
1488  
1489  
1490  
1491  
1492  
1493  
1494  
1495  
1496  
1497  
1498  
1499

```

1500 // Mutual exclusion lock.
1501 struct spinlock {
1502     uint locked;      // Is the lock held?
1503
1504     // For debugging:
1505     char *name;        // Name of lock.
1506     struct cpu *cpu;   // The cpu holding the lock.
1507     uint pcs[10];      // The call stack (an array of program counters)
1508                       // that locked the lock.
1509 };
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549

```

```

1550 // Mutual exclusion spin locks.
1551
1552 #include "types.h"
1553 #include "defs.h"
1554 #include "param.h"
1555 #include "x86.h"
1556 #include "memlayout.h"
1557 #include "mmu.h"
1558 #include "proc.h"
1559 #include "spinlock.h"
1560
1561 void
1562 initlock(struct spinlock *lk, char *name)
1563 {
1564     lk->name = name;
1565     lk->locked = 0;
1566     lk->cpu = 0;
1567 }
1568
1569 // Acquire the lock.
1570 // Loops (spins) until the lock is acquired.
1571 // Holding a lock for a long time may cause
1572 // other CPUs to waste time spinning to acquire it.
1573 void
1574 acquire(struct spinlock *lk)
1575 {
1576     pushcli(); // disable interrupts to avoid deadlock.
1577     if(holding(lk))
1578         panic("acquire");
1579
1580     // The xchg is atomic.
1581     // It also serializes, so that reads after acquire are not
1582     // reordered before it.
1583     while(xchg(&lk->locked, 1) != 0)
1584         ;
1585
1586     // Record info about lock acquisition for debugging.
1587     lk->cpu = cpu;
1588     getcallerpcs(&lk, lk->pcs);
1589 }
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599

```

```

1600 // Release the lock.
1601 void
1602 release(struct spinlock *lk)
1603 {
1604     if(!holding(lk))
1605         panic("release");
1606
1607     lk->pcs[0] = 0;
1608     lk->cpu = 0;
1609
1610     // The xchg serializes, so that reads before release are
1611     // not reordered after it. The 1996 PentiumPro manual (Volume 3,
1612     // 7.2) says reads can be carried out speculatively and in
1613     // any order, which implies we need to serialize here.
1614     // But the 2007 Intel 64 Architecture Memory Ordering White
1615     // Paper says that Intel 64 and IA-32 will not move a load
1616     // after a store. So lock->locked = 0 would work here.
1617     // The xchg being asm volatile ensures gcc emits it after
1618     // the above assignments (and after the critical section).
1619     xchg(&lk->locked, 0);
1620
1621     popcli();
1622 }
1623
1624 // Record the current call stack in pcs[] by following the %ebp chain.
1625 void
1626 getcallerpcs(void *v, uint pcs[])
1627 {
1628     uint *ebp;
1629     int i;
1630
1631     ebp = (uint*)v - 2;
1632     for(i = 0; i < 10; i++){
1633         if(ebp == 0 || ebp < (uint*)KERNBASE || ebp == (uint*)0xffffffff)
1634             break;
1635         pcs[i] = ebp[1]; // saved %eip
1636         ebp = (uint*)ebp[0]; // saved %ebp
1637     }
1638     for(; i < 10; i++)
1639         pcs[i] = 0;
1640 }
1641
1642 // Check whether this cpu is holding the lock.
1643 int
1644 holding(struct spinlock *lock)
1645 {
1646     return lock->locked && lock->cpu == cpu;
1647 }
1648
1649

```

```

1650 // Pushcli/popcli are like cli/sti except that they are matched:
1651 // it takes two popcli to undo two pushcli. Also, if interrupts
1652 // are off, then pushcli, popcli leaves them off.
1653
1654 void
1655 pushcli(void)
1656 {
1657     int eflags;
1658
1659     eflags = readeflags();
1660     cli();
1661     if(cpu->ncli++ == 0)
1662         cpu->intena = eflags & FL_IF;
1663 }
1664
1665 void
1666 popcli(void)
1667 {
1668     if(readeflags() & FL_IF)
1669         panic("popcli - interruptible");
1670     if(--cpu->ncli < 0)
1671         panic("popcli");
1672     if(cpu->ncli == 0 && cpu->intena)
1673         sti();
1674 }
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699

```



```

1700 #include "param.h"
1701 #include "types.h"
1702 #include "defs.h"
1703 #include "x86.h"
1704 #include "memlayout.h"
1705 #include "mmu.h"
1706 #include "proc.h"
1707 #include "elf.h"
1708
1709 extern char data[]; // defined by kernel.ld
1710 pde_t *kpgdir; // for use in scheduler()
1711 struct segdesc gdt[NSEGS];
1712
1713 // Set up CPU's kernel segment descriptors.
1714 // Run once on entry on each CPU.
1715 void
1716 seginit(void)
1717 {
1718     struct cpu *c;
1719
1720     // Map "logical" addresses to virtual addresses using identity map.
1721     // Cannot share a CODE descriptor for both kernel and user
1722     // because it would have to have DPL_USR, but the CPU forbids
1723     // an interrupt from CPL=0 to DPL=3.
1724     c = &cpus[cpunum()];
1725     c->gdt[SEG_KCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, 0);
1726     c->gdt[SEG_KDATA] = SEG(STA_W, 0, 0xffffffff, 0);
1727     c->gdt[SEG_UCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, DPL_USER);
1728     c->gdt[SEG_UDATA] = SEG(STA_W, 0, 0xffffffff, DPL_USER);
1729
1730     // Map cpu, and curproc
1731     c->gdt[SEG_KCPU] = SEG(STA_W, &c->cpu, 8, 0);
1732
1733     lgdt(c->gdt, sizeof(c->gdt));
1734     loadgs(SEG_KCPU << 3);
1735
1736     // Initialize cpu-local storage.
1737     cpu = c;
1738     proc = 0;
1739 }
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749

```

```

1750 // Return the address of the PTE in page table pgdir
1751 // that corresponds to virtual address va. If alloc!=0,
1752 // create any required page table pages.
1753 static pte_t *
1754 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1755 {
1756     pde_t *pde;
1757     pte_t *pgtab;
1758
1759     pde = &pgdir[PDX(va)];
1760     if(*pde & PTE_P){
1761         pgtab = (pte_t*)p2v(PTE_ADDR(*pde));
1762     } else {
1763         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1764             return 0;
1765         // Make sure all those PTE_P bits are zero.
1766         memset(pgtab, 0, PGSIZE);
1767         // The permissions here are overly generous, but they can
1768         // be further restricted by the permissions in the page table
1769         // entries, if necessary.
1770         *pde = v2p(pgtab) | PTE_P | PTE_W | PTE_U;
1771     }
1772     return &pgtab[PTX(va)];
1773 }
1774
1775 // Create PTEs for virtual addresses starting at va that refer to
1776 // physical addresses starting at pa. va and size might not
1777 // be page-aligned.
1778 static int
1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN((uint)va) + size - 1;
1786     for(;;){
1787         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
1799

```

```

1800 // There is one page table per process, plus one that's used when
1801 // a CPU is not running any process (kpgdir). The kernel uses the
1802 // current process's page table during system calls and interrupts;
1803 // page protection bits prevent user code from using the kernel's
1804 // mappings.
1805 //
1806 // setupkvm() and exec() set up every page table like this:
1807 //
1808 // 0..KERNBASE: user memory (text+data+stack+heap), mapped to
1809 // phys memory allocated by the kernel
1810 // KERNBASE..KERNBASE+EXTMEM: mapped to 0..EXTMEM (for I/O space)
1811 // KERNBASE+EXTMEM..data: mapped to EXTMEM..V2P(data)
1812 // for the kernel's instructions and r/o data
1813 // data..KERNBASE+PHYSTOP: mapped to V2P(data)..PHYSTOP,
1814 // rw data + free physical memory
1815 // 0xfe000000..0: mapped direct (devices such as ioapic)
1816 //
1817 // The kernel allocates physical memory for its heap and for user memory
1818 // between V2P(end) and the end of physical memory (PHYSTOP)
1819 // (directly addressable from end..P2V(PHYSTOP)).
1820 //
1821 // This table defines the kernel's mappings, which are present in
1822 // every process's page table.
1823 static struct kmap {
1824     void *virt;
1825     uint phys_start;
1826     uint phys_end;
1827     int perm;
1828 } kmap[] = {
1829     { (void*)KERNBASE, 0,          EXTMEM,    PTE_W}, // I/O space
1830     { (void*)KERNLINK, V2P(KERNLINK), 0}, // kern text+rodata
1831     { (void*)data,     V2P(data),   PHYSTOP,   PTE_W}, // kern data+memory
1832     { (void*)DEVSPACE, DEVSPACE,    0,        PTE_W}, // more devices
1833 };
1834
1835 // Set up kernel part of a page table.
1836 pde_t*
1837 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
1845     if (p2v(PHYSTOP) > (void*)DEVSPACE)
1846         panic("PHYSTOP too high");
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849             (uint)k->phys_start, k->perm) < 0)

```

```

1850         return 0;
1851     return pgdir;
1852 }
1853
1854 // Allocate one page table for the machine for the kernel address
1855 // space for scheduler processes.
1856 void
1857 kvmalloc(void)
1858 {
1859     kpgdir = setupkvm();
1860     switchkvm();
1861 }
1862
1863 // Switch h/w page table register to the kernel-only page table,
1864 // for when no process is running.
1865 void
1866 switchkvm(void)
1867 {
1868     lcr3(v2p(kpgdir)); // switch to the kernel page table
1869 }
1870
1871 // Switch TSS and h/w page table to correspond to process p.
1872 void
1873 switchvm(struct proc *p)
1874 {
1875     pushcli();
1876     cpu->gdt[SEG_TSS] = SEG16(STS_T32A, &cpu->ts, sizeof(cpu->ts)-1, 0);
1877     cpu->gdt[SEG_TSS].s = 0;
1878     cpu->ts.ss0 = SEG_KDATA << 3;
1879     cpu->ts.esp0 = (uint)proc->kstack + KSTACKSIZE;
1880     ltr(SEG_TSS << 3);
1881     if(p->pgdir == 0)
1882         panic("switchvm: no pgdir");
1883     lcr3(v2p(p->pgdir)); // switch to new address space
1884     popcli();
1885 }
1886
1887
1888
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899

```

```

1900 // Load the initcode into address 0 of pgdir.
1901 // sz must be less than a page.
1902 void
1903 inituvm(pde_t *pgdir, char *init, uint sz)
1904 {
1905     char *mem;
1906
1907     if(sz >= PGSIZE)
1908         panic("inituvm: more than a page");
1909     mem = kalloc();
1910     memset(mem, 0, PGSIZE);
1911     mappages(pgdir, 0, PGSIZE, v2p(mem), PTE_W|PTE_U);
1912     memmove(mem, init, sz);
1913 }
1914
1915 // Load a program segment into pgdir.  addr must be page-aligned
1916 // and the pages from addr to addr+sz must already be mapped.
1917 int
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint sz)
1919 {
1920     uint i, pa, n;
1921     pte_t *pte;
1922
1923     if((uint) addr % PGSIZE != 0)
1924         panic("loaduvm: addr must be page aligned");
1925     for(i = 0; i < sz; i += PGSIZE){
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1927             panic("loaduvm: address should exist");
1928         pa = PTE_ADDR(*pte);
1929         if(sz - i < PGSIZE)
1930             n = sz - i;
1931         else
1932             n = PGSIZE;
1933         if(readi(ip, p2v(pa), offset+i, n) != n)
1934             return -1;
1935     }
1936     return 0;
1937 }
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949

```

```

1950 // Allocate page tables and physical memory to grow process from oldsz to
1951 // newsz, which need not be page aligned.  Returns new size or 0 on error.
1952 int
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
1955     char *mem;
1956     uint a;
1957
1958     if(newsz >= KERNBASE)
1959         return 0;
1960     if(newsz < oldsz)
1961         return oldsz;
1962
1963     a = PGROUNDUP(oldsz);
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
1966         if(mem == 0){
1967             cprintf("allocuvm out of memory\n");
1968             deallocuvm(pgdir, newsz, oldsz);
1969             return 0;
1970         }
1971         memset(mem, 0, PGSIZE);
1972         mappages(pgdir, (char*)a, PGSIZE, v2p(mem), PTE_W|PTE_U);
1973     }
1974     return newsz;
1975 }
1976
1977 // Deallocate user pages to bring the process size from oldsz to
1978 // newsz.  oldsz and newsz need not be page-aligned, nor does newsz
1979 // need to be less than oldsz.  oldsz can be larger than the actual
1980 // process size.  Returns the new process size.
1981 int
1982 deallocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1983 {
1984     pte_t *pte;
1985     uint a, pa;
1986
1987     if(newsz >= oldsz)
1988         return oldsz;
1989
1990     a = PGROUNDUP(newsz);
1991     for(; a < oldsz; a += PGSIZE){
1992         pte = walkpgdir(pgdir, (char*)a, 0);
1993         if(!pte)
1994             a += (NPENTRIES - 1) * PGSIZE;
1995         else if((*pte & PTE_P) != 0){
1996             pa = PTE_ADDR(*pte);
1997             if(pa == 0)
1998                 panic("kfree");
1999             char *v = p2v(pa);

```

```

2000     kfree(v);
2001     *pte = 0;
2002 }
2003 }
2004 return newsz;
2005 }
2006
2007 // Free a page table and all the physical memory pages
2008 // in the user part.
2009 void
2010 freevm(pde_t *pgdir)
2011 {
2012     uint i;
2013
2014     if(pgdir == 0)
2015         panic("freevm: no pgdir");
2016     deallocvm(pgdir, KERNBASE, 0);
2017     for(i = 0; i < NPENTRIES; i++){
2018         if(pgdir[i] & PTE_P){
2019             char *v = p2v(PTE_ADDR(pgdir[i]));
2020             kfree(v);
2021         }
2022     }
2023     kfree((char*)pgdir);
2024 }
2025
2026 // Clear PTE_U on a page. Used to create an inaccessible
2027 // page beneath the user stack.
2028 void
2029 clearpteu(pde_t *pgdir, char *uva)
2030 {
2031     pte_t *pte;
2032
2033     pte = walkpgdir(pgdir, uva, 0);
2034     if(pte == 0)
2035         panic("clearpteu");
2036     *pte &= ~PTE_U;
2037 }
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049

```

```

2050 // Given a parent process's page table, create a copy
2051 // of it for a child.
2052 pde_t*
2053 copyuvm(pde_t *pgdir, uint sz)
2054 {
2055     pde_t *d;
2056     pte_t *pte;
2057     uint pa, i, flags;
2058     char *mem;
2059
2060     if((d = setupkvm()) == 0)
2061         return 0;
2062     for(i = 0; i < sz; i += PGSIZE){
2063         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
2064             panic("copyuvm: pte should exist");
2065         if(!(*pte & PTE_P))
2066             panic("copyuvm: page not present");
2067         pa = PTE_ADDR(*pte);
2068         flags = PTE_FLAGS(*pte);
2069         if((mem = kalloc()) == 0)
2070             goto bad;
2071         memmove(mem, (char*)p2v(pa), PGSIZE);
2072         if(mappages(d, (void*)i, PGSIZE, v2p(mem), flags) < 0)
2073             goto bad;
2074     }
2075     return d;
2076
2077 bad:
2078     freevm(d);
2079     return 0;
2080 }
2081
2082
2083
2084
2085
2086
2087
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099

```

```

2100 // Map user virtual address to kernel address.
2101 char*
2102 uva2ka(pde_t *pgdir, char *uva)
2103 {
2104     pte_t *pte;
2105
2106     pte = walkpgdir(pgdir, uva, 0);
2107     if((*pte & PTE_P) == 0)
2108         return 0;
2109     if((*pte & PTE_U) == 0)
2110         return 0;
2111     return (char*)p2v(PTE_ADDR(*pte));
2112 }
2113
2114 // Copy len bytes from p to user address va in page table pgdir.
2115 // Most useful when pgdir is not the current page table.
2116 // uva2ka ensures this only works for PTE_U pages.
2117 int
2118 copyout(pde_t *pgdir, uint va, void *p, uint len)
2119 {
2120     char *buf, *pa0;
2121     uint n, va0;
2122
2123     buf = (char*)p;
2124     while(len > 0){
2125         va0 = (uint)PGROUNDDOWN(va);
2126         pa0 = uva2ka(pgdir, (char*)va0);
2127         if(pa0 == 0)
2128             return -1;
2129         n = PGSIZE - (va - va0);
2130         if(n > len)
2131             n = len;
2132         memmove(pa0 + (va - va0), buf, n);
2133         len -= n;
2134         buf += n;
2135         va = va0 + PGSIZE;
2136     }
2137     return 0;
2138 }
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149

```

```

2150 // Blank page.
2151
2152
2153
2154
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199

```

2200 // Blank page.

2201  
2202  
2203  
2204  
2205  
2206  
2207  
2208  
2209  
2210  
2211  
2212  
2213  
2214  
2215  
2216  
2217  
2218  
2219  
2220  
2221  
2222  
2223  
2224  
2225  
2226  
2227  
2228  
2229  
2230  
2231  
2232  
2233  
2234  
2235  
2236  
2237  
2238  
2239  
2240  
2241  
2242  
2243  
2244  
2245  
2246  
2247  
2248  
2249

2250 // Blank page.

2251  
2252  
2253  
2254  
2255  
2256  
2257  
2258  
2259  
2260  
2261  
2262  
2263  
2264  
2265  
2266  
2267  
2268  
2269  
2270  
2271  
2272  
2273  
2274  
2275  
2276  
2277  
2278  
2279  
2280  
2281  
2282  
2283  
2284  
2285  
2286  
2287  
2288  
2289  
2290  
2291  
2292  
2293  
2294  
2295  
2296  
2297  
2298  
2299

```

2300 // Segments in proc->gdt.
2301 #define NSEGS      7
2302 #define UID_DEFAULT 7
2303 #define GID_DEFAULT 5
2304
2305 // Per-CPU state
2306 struct cpu {
2307     uchar id;                    // Local APIC ID; index into cpus[] below
2308     struct context *scheduler;    // swtch() here to enter scheduler
2309     struct taskstate ts;         // Used by x86 to find stack for interrupt
2310     struct segdesc gdt[NSEGS];   // x86 global descriptor table
2311     volatile uint started;       // Has the CPU started?
2312     int ncli;                    // Depth of pushcli nesting.
2313     int intena;                  // Were interrupts enabled before pushcli?
2314
2315     // Cpu-local storage variables; see below
2316     struct cpu *cpu;
2317     struct proc *proc;           // The currently-running process.
2318 };
2319
2320 extern struct cpu cpus[NCPU];
2321 extern int ncpu;
2322
2323 // Per-CPU variables, holding pointers to the
2324 // current cpu and to the current process.
2325 // The asm suffix tells gcc to use "%gs:0" to refer to cpu
2326 // and "%gs:4" to refer to proc.  seginit sets up the
2327 // %gs segment register so that %gs refers to the memory
2328 // holding those two variables in the local cpu's struct cpu.
2329 // This is similar to how thread-local variables are implemented
2330 // in thread libraries such as Linux pthreads.
2331 extern struct cpu *cpu asm("%gs:0"); // &cpus[cpunum()]
2332 extern struct proc *proc asm("%gs:4"); // cpus[cpunum()].proc
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349

```

```

2350 // Saved registers for kernel context switches.
2351 // Don't need to save all the segment registers (%cs, etc),
2352 // because they are constant across kernel contexts.
2353 // Don't need to save %eax, %ecx, %edx, because the
2354 // x86 convention is that the caller has saved them.
2355 // Contexts are stored at the bottom of the stack they
2356 // describe; the stack pointer is the address of the context.
2357 // The layout of the context matches the layout of the stack in swtch.S
2358 // at the "Switch stacks" comment. Switch doesn't save eip explicitly,
2359 // but it is on the stack and allocproc() manipulates it.
2360 struct context {
2361     uint edi;
2362     uint esi;
2363     uint ebx;
2364     uint ebp;
2365     uint eip;
2366 };
2367
2368 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
2369 //no one can use processes in embryo state.
2370 // Per-process state
2371 struct proc {
2372     uint sz;                        // Size of process memory (bytes)
2373     pde_t * pgdir;                 // Page table
2374     char *kstack;                  // Bottom of kernel stack for this process
2375     enum procstate state;          // Process state
2376     int pid;                        // Process ID
2377     uint uid;                       // User ID
2378     uint gid;                       // Group ID
2379     int ppid;                       // Parent process
2380     struct proc *parent;            // Parent process
2381     struct trapframe *tf;           // Trap frame for current syscall
2382     struct context *context;        // swtch() here to run process
2383     void *chan;                     // If non-zero, sleeping on chan
2384     int killed;                     // If non-zero, have been killed
2385     struct file *ofile[NOFILE];    // Open files
2386     struct inode *cwd;              // Current directory
2387     char name[16];                  // Process name (debugging)
2388     struct proc *next;
2389     int priority;
2390 };
2391
2392
2393
2394
2395 // Process memory is laid out contiguously, low addresses first:
2396 //   text
2397 //   original data and bss
2398 //   fixed-size stack
2399 //   expandable heap

```

```

2400 #include "types.h"
2401 #include "defs.h"
2402 #include "param.h"
2403 #include "memlayout.h"
2404 #include "mmu.h"
2405 #include "x86.h"
2406 #include "proc.h"
2407 #include "spinlock.h"
2408 #include "ps.h"
2409
2410 #define NULL 0
2411
2412 static char *states[] = {
2413     [UNUSED]    "UNUSED",
2414     [EMBRYO]    "EMBRYO",
2415     [SLEEPING]  "SLEEPING",
2416     [RUNNABLE]  "RUNNABLE",
2417     [RUNNING]   "RUNNING",
2418     [ZOMBIE]    "ZOMBIE"
2419 };
2420
2421
2422 //uses round robin
2423 //nproc is set to 64
2424 struct {
2425     struct spinlock lock;
2426     struct proc proc[NPROC];
2427     struct proc *pReadyList[2];
2428     // struct proc *pFreeList;
2429     // uint TimeToReset;
2430 } ptable;
2431
2432 static struct proc *initproc;
2433
2434 int nextpid = 1;
2435 extern void forkret(void);
2436 extern void trapret(void);
2437
2438 static void wakeup1(void *chan);
2439
2440 void
2441 pinit(void)
2442 {
2443     initlock(&ptable.lock, "ptable");
2444 }
2445
2446
2447
2448
2449

```

```

2450 // Look in the process table for an UNUSED proc.
2451 // If found, change state to EMBRYO and initialize
2452 // state required to run in the kernel.
2453 // Otherwise return 0.
2454 static struct proc*
2455 allocproc(void)
2456 {
2457     struct proc *p;
2458     char *sp;
2459
2460     acquire(&ptable.lock);
2461     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
2462         if(p->state == UNUSED)
2463             goto found;
2464     release(&ptable.lock);
2465     return 0;
2466
2467 found:
2468     p->state = EMBRYO;
2469     p->pid = nextpid++;
2470     release(&ptable.lock);
2471
2472     // Allocate kernel stack.
2473     if((p->kstack = kalloc()) == 0){
2474         p->state = UNUSED;
2475         return 0;
2476     }
2477     sp = p->kstack + KSTACKSIZE;
2478
2479     // Leave room for trap frame.
2480     sp -= sizeof *p->tf;
2481     p->tf = (struct trapframe*)sp;
2482
2483     // Set up new context to start executing at forkret,
2484     // which returns to trapret.
2485     sp -= 4;
2486     *(uint*)sp = (uint)trapret;
2487
2488     sp -= sizeof *p->context;
2489     p->context = (struct context*)sp;
2490     memset(p->context, 0, sizeof *p->context);
2491     p->context->eip = (uint)forkret;
2492
2493     return p;
2494 }
2495
2496
2497
2498
2499

```



```

2500 // Set up first user process.
2501 void
2502 userinit(void)
2503 {
2504     struct proc *p;
2505     extern char _binary_initcode_start[], _binary_initcode_size[];
2506
2507     p = allocproc();
2508     initproc = p;
2509     if((p->pgdir = setupkvm()) == 0)
2510         panic("userinit: out of memory?");
2511     inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
2512     p->sz = PGSIZE;
2513     memset(p->tf, 0, sizeof(*p->tf));
2514     p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
2515     p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
2516     p->tf->es = p->tf->ds;
2517     p->tf->ss = p->tf->ds;
2518     p->tf->eflags = FL_IF;
2519     p->tf->esp = PGSIZE;
2520     p->tf->eip = 0; // beginning of initcode.S
2521
2522     safestrcpy(p->name, "initcode", sizeof(p->name));
2523     p->cwd = namei("/");
2524
2525     p->state = RUNNABLE;
2526     p->uid = UID_DEFAULT;
2527     p->gid = GID_DEFAULT;
2528
2529     p->ppid = 1;
2530     p->priority = 1;
2531     ptable.pReadyList[1] = p;
2532     p->next = NULL;
2533 }
2534 }
2535
2536
2537
2538
2539
2540
2541
2542
2543
2544
2545
2546
2547
2548
2549

```

```

2550 // Grow current process's memory by n bytes.
2551 // Return 0 on success, -1 on failure.
2552 int
2553 growproc(int n)
2554 {
2555     uint sz;
2556
2557     sz = proc->sz;
2558     if(n > 0){
2559         if((sz = allocuvm(proc->pgdir, sz, sz + n)) == 0)
2560             return -1;
2561     } else if(n < 0){
2562         if((sz = deallocuvm(proc->pgdir, sz, sz + n)) == 0)
2563             return -1;
2564     }
2565     proc->sz = sz;
2566     switchuvm(proc);
2567     return 0;
2568 }
2569
2570 // Create a new process copying p as the parent.
2571 // Sets up stack to return as if from system call.
2572 // Caller must set state of returned proc to RUNNABLE.
2573 int
2574 fork(void)
2575 {
2576     int i, pid;
2577     struct proc *np;
2578
2579     // Allocate process.
2580     if((np = allocproc()) == 0)
2581         return -1;
2582
2583     // Copy process state from p.
2584     if((np->pgdir = copyuvm(proc->pgdir, proc->sz)) == 0){
2585         kfree(np->kstack);
2586         np->kstack = 0;
2587         np->state = UNUSED;
2588         return -1;
2589     }
2590     np->sz = proc->sz;
2591     np->parent = proc;
2592     *np->tf = *proc->tf;
2593
2594     // Clear %eax so that fork returns 0 in the child.
2595     np->tf->eax = 0;
2596
2597
2598
2599

```

```

2600 for(i = 0; i < NOFILE; i++)
2601     if(proc->ofile[i])
2602         np->ofile[i] = filedup(proc->ofile[i]);
2603 np->cwd = idup(proc->cwd);
2604
2605 safestrcpy(np->name, proc->name, sizeof(proc->name));
2606
2607 // acquire(&ptable.lock);
2608 np->uid = proc->uid;
2609 np->gid = proc->gid;
2610 np->ppid = proc->pid;
2611 np->priority = 1;
2612 // release(&ptable.lock);
2613
2614
2615
2616 pid = np->pid;
2617
2618
2619 // lock to force the compiler to emit the np->state write last.
2620 acquire(&ptable.lock);
2621 np->state = RUNNABLE;
2622 putinQ(np);
2623
2624 release(&ptable.lock);
2625
2626 return pid;
2627 }
2628
2629 // Exit the current process. Does not return.
2630 // An exited process remains in the zombie state
2631 // until its parent calls wait() to find out it exited.
2632 void
2633 exit(void)
2634 {
2635     struct proc *p;
2636     int fd;
2637
2638     if(proc == initproc)
2639         panic("init exiting");
2640
2641     // Close all open files.
2642     for(fd = 0; fd < NOFILE; fd++){
2643         if(proc->ofile[fd]){
2644             fileclose(proc->ofile[fd]);
2645             proc->ofile[fd] = 0;
2646         }
2647     }
2648
2649

```

```

2650 begin_op();
2651 iput(proc->cwd);
2652 end_op();
2653 proc->cwd = 0;
2654
2655 acquire(&ptable.lock);
2656
2657 // Parent might be sleeping in wait().
2658 wakeupl(proc->parent);
2659
2660 // Pass abandoned children to init.
2661 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2662     if(p->parent == proc){
2663         p->parent = initproc;
2664         if(p->state == ZOMBIE)
2665             wakeupl(initproc);
2666     }
2667 }
2668
2669 // Jump into the scheduler, never to return.
2670 proc->state = ZOMBIE;
2671 sched();
2672 panic("zombie exit");
2673 }
2674
2675 // Wait for a child process to exit and return its pid.
2676 // Return -1 if this process has no children.
2677 int
2678 wait(void)
2679 {
2680     struct proc *p;
2681     int havekids, pid;
2682
2683     acquire(&ptable.lock);
2684     for(;;){
2685         // Scan through table looking for zombie children.
2686         havekids = 0;
2687         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2688             if(p->parent != proc)
2689                 continue;
2690             havekids = 1;
2691             if(p->state == ZOMBIE){
2692                 // Found one.
2693                 pid = p->pid;
2694                 kfree(p->kstack);
2695                 p->kstack = 0;
2696                 freevm(p->pgdir);
2697                 p->state = UNUSED;
2698                 p->pid = 0;
2699                 p->parent = 0;

```

```

2700     p->name[0] = 0;
2701     p->killed = 0;
2702     release(&ptable.lock);
2703     return pid;
2704 }
2705 }
2706
2707 // No point waiting if we don't have any children.
2708 if(!havekids || proc->killed){
2709     release(&ptable.lock);
2710     return -1;
2711 }
2712
2713 // Wait for children to exit. (See wakeup1 call in proc_exit.)
2714 sleep(proc, &ptable.lock);
2715 }
2716 }
2717
2718
2719
2720
2721
2722
2723
2724
2725
2726
2727
2728
2729
2730
2731
2732
2733
2734
2735
2736
2737
2738
2739
2740
2741
2742
2743
2744
2745
2746
2747
2748
2749

```

```

2750 // Per-CPU process scheduler.
2751 // Each CPU calls scheduler() after setting itself up.
2752 // Scheduler never returns. It loops, doing:
2753 // - choose a process to run
2754 // - switch to start running that process
2755 // - eventually that process transfers control
2756 //   via swtch back to the scheduler.
2757 /*void
2758 scheduler(void)
2759 {
2760     struct proc *p;
2761
2762     for(;;){
2763         // Enable interrupts on this processor.
2764         sti();
2765
2766         // Loop over process table looking for process to run.
2767         acquire(&ptable.lock);
2768         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2769             if(p->state != RUNNABLE)
2770                 continue;
2771
2772             // Switch to chosen process. It is the process's job
2773             // to release ptable.lock and then reacquire it
2774             // before jumping back to us.
2775
2776             proc = p;
2777             switchvm(p);
2778             p->state = RUNNING;
2779             swtch(&cpu->scheduler, proc->context); //context swtch. proc->context r1
2780             switchkvm(); //switch to the correct kernel virtual memory.
2781
2782             // Process is done running for now.
2783             // It should have changed its p->state before coming back.
2784             proc = 0;
2785         }
2786         release(&ptable.lock);
2787     }
2788 }
2789 */
2790 */
2791
2792 // Enter scheduler. Must hold only ptable.lock
2793 // and have changed proc->state.
2794 void
2795 sched(void)
2796 {
2797     int intena;
2798
2799     if(!holding(&ptable.lock))

```

```

2800     panic("sched ptable.lock");
2801     if(cpu->ncli != 1)
2802         panic("sched locks");
2803     if(proc->state == RUNNING)
2804         panic("sched running");
2805     if(readeflags() & FL_IF)
2806         panic("sched interruptible");
2807     intena = cpu->intena;
2808     swtch(&proc->context, cpu->scheduler);
2809     cpu->intena = intena;
2810 }
2811
2812 // Give up the CPU for one scheduling round.
2813 void
2814 yield(void)
2815 {
2816     acquire(&ptable.lock);
2817     proc->state = RUNNABLE;
2818     putinQ(proc);
2819     sched();
2820     release(&ptable.lock);
2821 }
2822
2823 // A fork child's very first scheduling by scheduler()
2824 // will swtch here. "Return" to user space.
2825 void
2826 forkret(void)
2827 {
2828     static int first = 1;
2829     // Still holding ptable.lock from scheduler.
2830     release(&ptable.lock);
2831
2832     if (first) {
2833         // Some initialization functions must be run in the context
2834         // of a regular process (e.g., they call sleep), and thus cannot
2835         // be run from main().
2836         first = 0;
2837         iinit(ROOTDEV);
2838         initlog(ROOTDEV);
2839     }
2840
2841     // Return to "caller", actually trapret (see allocproc).
2842 }
2843
2844
2845
2846
2847
2848
2849

```

```

2850 // Atomically release lock and sleep on chan.
2851 // Reacquires lock when awakened.
2852 void
2853 sleep(void *chan, struct spinlock *lk)
2854 {
2855     if(proc == 0)
2856         panic("sleep");
2857
2858     if(lk == 0)
2859         panic("sleep without lk");
2860
2861     // Must acquire ptable.lock in order to
2862     // change p->state and then call sched.
2863     // Once we hold ptable.lock, we can be
2864     // guaranteed that we won't miss any wakeup
2865     // (wakeup runs with ptable.lock locked),
2866     // so it's okay to release lk.
2867     if(lk != &ptable.lock){
2868         acquire(&ptable.lock);
2869         release(lk);
2870     }
2871
2872     // Go to sleep.
2873     proc->chan = chan;
2874     proc->state = SLEEPING;
2875     sched();
2876
2877     // Tidy up.
2878     proc->chan = 0;
2879
2880     // Reacquire original lock.
2881     if(lk != &ptable.lock){
2882         release(&ptable.lock);
2883         acquire(lk);
2884     }
2885 }
2886
2887
2888
2889
2890
2891
2892
2893
2894
2895
2896
2897
2898
2899

```

```

2900 // Wake up all processes sleeping on chan.
2901 // The ptable lock must be held.
2902 static void
2903 wakeup1(void *chan)
2904 {
2905     struct proc *p;
2906
2907     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2908         if(p->state == SLEEPING && p->chan == chan){
2909             p->state = RUNNABLE;
2910             putinQ(p);
2911         }
2912     }
2913
2914 // Wake up all processes sleeping on chan.
2915 void
2916 wakeup(void *chan)
2917 {
2918     acquire(&ptable.lock);
2919     wakeup1(chan);
2920     release(&ptable.lock);
2921 }
2922
2923 // Kill the process with the given pid.
2924 // Process won't exit until it returns
2925 // to user space (see trap in trap.c).
2926 int
2927 kill(int pid)
2928 {
2929     struct proc *p;
2930
2931     acquire(&ptable.lock);
2932     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2933         if(p->pid == pid){
2934             p->killed = 1;
2935             // Wake process from sleep if necessary.
2936             if(p->state == SLEEPING){
2937                 p->state = RUNNABLE;
2938             }
2939             putinQ(p);
2940             release(&ptable.lock);
2941             return 0;
2942         }
2943     }
2944     release(&ptable.lock);
2945     return -1;
2946 }
2947
2948
2949

```

```

2950 int
2951 setuid(int uid)
2952 {
2953     if(uid<0)
2954         return -1;
2955
2956     acquire(&ptable.lock);
2957     proc->uid = uid;
2958     release(&ptable.lock);
2959
2960     return 0;
2961 }
2962
2963 int
2964 setgid(int gid)
2965 {
2966     if(gid<0)
2967         return -1;
2968
2969     acquire(&ptable.lock);
2970     proc->gid = gid;
2971     release(&ptable.lock);
2972
2973     return 0;
2974 }
2975
2976 int
2977 getuid()
2978 {
2979     return proc->uid;
2980 }
2981
2982 int
2983 getgid()
2984 {
2985     return proc->gid;
2986 }
2987
2988
2989
2990
2991
2992
2993
2994
2995
2996
2997
2998
2999

```

```

3000 // Print a process listing to console. For debugging.
3001 // Runs when user types ^P on console.
3002 // No lock to avoid wedging a stuck machine further.
3003 void
3004 procdump(void)
3005 {
3006     int i;
3007     struct proc *p;
3008     char *state;
3009     uint pc[10];
3010
3011     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
3012         if(p->state == UNUSED)
3013             continue;
3014         if(p->state >= 0 && p->state < NELEM(states) && states[p->state])
3015             state = states[p->state];
3016         else
3017             state = "???";
3018         cprintf("%d %d %d %s %s", p->pid, p->uid, p->gid, state, p->name);
3019         if(p->state == SLEEPING){
3020             getcallerpcs((uint*)p->context->ebp+2, pc);
3021             for(i=0; i<10 && pc[i] != 0; i++)
3022                 cprintf(" %p", pc[i]);
3023         }
3024         cprintf("\n");
3025     }
3026 }
3027
3028 int
3029 getprocs(int max, struct uproc *table)
3030 {
3031     int count=0;
3032     struct proc *p;
3033
3034     acquire(&ptable.lock);
3035     for(p=ptable.proc; p<&ptable.proc[NPROC]; p++){
3036
3037         if(p->state == UNUSED || p->state == ZOMBIE || p->state == EMBRYO)
3038             continue;
3039         if(count >= max){
3040             release(&ptable.lock);
3041             return count;
3042         }
3043
3044         table[count].pid = p->pid;
3045         table[count].uid = p->uid;
3046         table[count].gid = p->gid;
3047         table[count].ppid = p->ppid;
3048         safestrcpy(table[count].state, states[p->state], sizeof(table[count].state));
3049         table[count].size = p->sz;

```

```

3050         safestrcpy(table[count].name, p->name, sizeof(table[count].name));
3051
3052         count = count+1;
3053     }
3054
3055     release(&ptable.lock);
3056     if(max >= count)
3057         return count;
3058
3059     return -1;
3060 }
3061
3062 void scheduler(void){
3063     struct proc *p;
3064     int i;
3065     for(;;){
3066         // Enable interrupts on this processor.
3067         sti();
3068
3069         // Loop over process table looking for process to run.
3070         acquire(&ptable.lock);
3071
3072         p=ptable.pReadyList[0];
3073         i = 0;
3074         if(p == NULL){
3075             p=ptable.pReadyList[1];
3076             i=1;
3077             if(p == NULL){
3078                 p=ptable.pReadyList[2];
3079                 i=2;
3080             }
3081         }
3082
3083         if(p!=NULL){
3084             proc = p;
3085
3086             ptable.pReadyList[i]= ptable.pReadyList[i]->next;
3087             p->next = NULL;
3088             switchvm(p);
3089             p->state = RUNNING;
3090             swtch(&cpu->scheduler, proc->context); //context swtch. proc->context r
3091             switchkvm(); //switch to the correct kernel virtual memory.
3092
3093             // Process is done running for now.
3094             // It should have changed its p->state before coming back.
3095             proc = 0;
3096         }

```

```

3100     release(&ptable.lock);
3101   }
3102 }
3103
3104
3105 int putinQ(struct proc *prc){
3106
3107     acquire(&ptable.lock);
3108
3109     if(prc->priority == 0){
3110         addtoq(&ptable.pReadyList[0], prc);
3111         release(&ptable.lock);
3112         return 0;
3113     }
3114     else if (prc->priority == 1){
3115         addtoq(&ptable.pReadyList[1], prc);
3116         release(&ptable.lock);
3117         return 0;
3118     }
3119     else if(prc->priority == 2){
3120         addtoq(&ptable.pReadyList[2], prc);
3121         release(&ptable.lock);
3122         return 0;
3123     }
3124
3125     return 1;
3126
3127 }
3128
3129 int addtoq(struct proc **p, struct proc *prc){
3130     if(!*p)
3131         *p=prc;
3132     else
3133     {
3134         struct proc *current;
3135         current = *p;
3136         while(current->next != NULL)
3137             current = current->next;
3138
3139         current->next = prc;
3140         prc->next = NULL;
3141     }
3142
3143     return 0;
3144 }
3145
3146
3147
3148
3149

```

```

3150 # Context switch
3151 #
3152 # void swtch(struct context **old, struct context *new);
3153 #
3154 # Save current register context in old
3155 # and then load register context from new.
3156
3157 .globl swtch
3158 swtch:
3159     movl 4(%esp), %eax
3160     movl 8(%esp), %edx
3161
3162     # Save old callee-save registers
3163     pushl %ebp
3164     pushl %ebx
3165     pushl %esi
3166     pushl %edi
3167
3168     # Switch stacks
3169     movl %esp, (%eax)
3170     movl %edx, %esp
3171
3172     # Load new callee-save registers
3173     popl %edi
3174     popl %esi
3175     popl %ebx
3176     popl %ebp
3177     ret
3178
3179
3180
3181
3182
3183
3184
3185
3186
3187
3188
3189
3190
3191
3192
3193
3194
3195
3196
3197
3198
3199

```

```

3200 // Physical memory allocator, intended to allocate
3201 // memory for user processes, kernel stacks, page table pages,
3202 // and pipe buffers. Allocates 4096-byte pages.
3203
3204 #include "types.h"
3205 #include "defs.h"
3206 #include "param.h"
3207 #include "memlayout.h"
3208 #include "mmu.h"
3209 #include "spinlock.h"
3210
3211 void freerange(void *vstart, void *vend);
3212 extern char end[]; // first address after kernel loaded from ELF file
3213
3214 struct run {
3215     struct run *next;
3216 };
3217
3218 struct {
3219     struct spinlock lock;
3220     int use_lock;
3221     struct run *freelist;
3222 } kmem;
3223
3224 // Initialization happens in two phases.
3225 // 1. main() calls kinit1() while still using entrypgdir to place just
3226 // the pages mapped by entrypgdir on free list.
3227 // 2. main() calls kinit2() with the rest of the physical pages
3228 // after installing a full page table that maps them on all cores.
3229 void
3230 kinit1(void *vstart, void *vend)
3231 {
3232     initlock(&kmem.lock, "kmem");
3233     kmem.use_lock = 0;
3234     freerange(vstart, vend);
3235 }
3236
3237 void
3238 kinit2(void *vstart, void *vend)
3239 {
3240     freerange(vstart, vend);
3241     kmem.use_lock = 1;
3242 }
3243
3244
3245
3246
3247
3248
3249

```

```

3250 void
3251 freerange(void *vstart, void *vend)
3252 {
3253     char *p;
3254     p = (char*)PGROUNDUP((uint)vstart);
3255     for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
3256         kfree(p);
3257 }
3258
3259
3260 // Free the page of physical memory pointed at by v,
3261 // which normally should have been returned by a
3262 // call to kalloc(). (The exception is when
3263 // initializing the allocator; see kinit above.)
3264 void
3265 kfree(char *v)
3266 {
3267     struct run *r;
3268
3269     if((uint)v % PGSIZE || v < end || v2p(v) >= PHYSTOP)
3270         panic("kfree");
3271
3272     // Fill with junk to catch dangling refs.
3273     memset(v, 1, PGSIZE);
3274
3275     if(kmem.use_lock)
3276         acquire(&kmem.lock);
3277     r = (struct run*)v;
3278     r->next = kmem.freelist;
3279     kmem.freelist = r;
3280     if(kmem.use_lock)
3281         release(&kmem.lock);
3282 }
3283
3284 // Allocate one 4096-byte page of physical memory.
3285 // Returns a pointer that the kernel can use.
3286 // Returns 0 if the memory cannot be allocated.
3287 char*
3288 kalloc(void)
3289 {
3290     struct run *r;
3291
3292     if(kmem.use_lock)
3293         acquire(&kmem.lock);
3294     r = kmem.freelist;
3295     if(r)
3296         kmem.freelist = r->next;
3297     if(kmem.use_lock)
3298         release(&kmem.lock);
3299     return (char*)r;

```



```

3300 }
3301
3302
3303
3304
3305
3306
3307
3308
3309
3310
3311
3312
3313
3314
3315
3316
3317
3318
3319
3320
3321
3322
3323
3324
3325
3326
3327
3328
3329
3330
3331
3332
3333
3334
3335
3336
3337
3338
3339
3340
3341
3342
3343
3344
3345
3346
3347
3348
3349

```

```

3350 // x86 trap and interrupt constants.
3351
3352 // Processor-defined:
3353 #define T_DIVIDE      0      // divide error
3354 #define T_DEBUG      1      // debug exception
3355 #define T_NMI        2      // non-maskable interrupt
3356 #define T_BRKPT      3      // breakpoint
3357 #define T_OFLOW      4      // overflow
3358 #define T_BOUND      5      // bounds check
3359 #define T_ILLOP      6      // illegal opcode
3360 #define T_DEVICE      7      // device not available
3361 #define T_DBLFLT     8      // double fault
3362 // #define T_COPROC    9      // reserved (not used since 486)
3363 #define T_TSS        10     // invalid task switch segment
3364 #define T_SEGNP      11     // segment not present
3365 #define T_STACK      12     // stack exception
3366 #define T_GPFLT      13     // general protection fault
3367 #define T_PGFLT      14     // page fault
3368 // #define T_RES       15     // reserved
3369 #define T_FPERR      16     // floating point error
3370 #define T_ALIGN      17     // alignment check
3371 #define T_MCHK       18     // machine check
3372 #define T_SIMDERR    19     // SIMD floating point error
3373
3374 // These are arbitrarily chosen, but with care not to overlap
3375 // processor defined exceptions or interrupt vectors.
3376 #define T_SYSCALL     64     // system call
3377 #define T_DEFAULT     500    // catchall
3378
3379 #define T_IRQ0        32     // IRQ 0 corresponds to int T_IRQ
3380
3381 #define IRQ_TIMER      0
3382 #define IRQ_KBD        1
3383 #define IRQ_COM1       4
3384 #define IRQ_IDE       14
3385 #define IRQ_ERROR      19
3386 #define IRQ_SPURIOUS   31
3387
3388
3389
3390
3391
3392
3393
3394
3395
3396
3397
3398
3399

```

```

3400 #!/usr/bin/perl -w
3401
3402 # Generate vectors.S, the trap/interrupt entry points.
3403 # There has to be one entry point per interrupt number
3404 # since otherwise there's no way for trap() to discover
3405 # the interrupt number.
3406
3407 print "# generated by vectors.pl - do not edit\n";
3408 print "# handlers\n";
3409 print ".globl alltraps\n";
3410 for(my $i = 0; $i < 256; $i++){
3411     print ".globl vector$i\n";
3412     print "vector$i:\n";
3413     if(!($i == 8 || ($i >= 10 && $i <= 14) || $i == 17)){
3414         print "    pushl $0\n";
3415     }
3416     print "    pushl $$i\n";
3417     print "    jmp alltraps\n";
3418 }
3419
3420 print "\n# vector table\n";
3421 print ".data\n";
3422 print ".globl vectors\n";
3423 print "vectors:\n";
3424 for(my $i = 0; $i < 256; $i++){
3425     print "    .long vector$i\n";
3426 }
3427
3428 # sample output:
3429 #   # handlers
3430 #   .globl alltraps
3431 #   .globl vector0
3432 #   vector0:
3433 #       pushl $0
3434 #       pushl $0
3435 #       jmp alltraps
3436 #   ...
3437 #
3438 #   # vector table
3439 #   .data
3440 #   .globl vectors
3441 #   vectors:
3442 #       .long vector0
3443 #       .long vector1
3444 #       .long vector2
3445 #   ...
3446
3447
3448
3449

```

```

3450 #include "mmu.h"
3451
3452 # vectors.S sends all traps here.
3453 .globl alltraps
3454 alltraps:
3455     # Build trap frame.
3456     pushl %ds
3457     pushl %es
3458     pushl %fs
3459     pushl %gs
3460     pushal
3461
3462     # Set up data and per-cpu segments.
3463     movw $(SEG_KDATA<<3), %ax
3464     movw %ax, %ds
3465     movw %ax, %es
3466     movw $(SEG_KCPU<<3), %ax
3467     movw %ax, %fs
3468     movw %ax, %gs
3469
3470     # Call trap(tf), where tf=%esp
3471     pushl %esp
3472     call trap
3473     addl $4, %esp
3474
3475     # Return falls through to trapret...
3476 .globl trapret
3477 trapret:
3478     popal
3479     popl %gs
3480     popl %fs
3481     popl %es
3482     popl %ds
3483     addl $0x8, %esp # trapno and errcode
3484     iret
3485
3486
3487
3488
3489
3490
3491
3492
3493
3494
3495
3496
3497
3498
3499

```

```

3500 #include "types.h"
3501 #include "defs.h"
3502 #include "param.h"
3503 #include "memlayout.h"
3504 #include "mmu.h"
3505 #include "proc.h"
3506 #include "x86.h"
3507 #include "traps.h"
3508 #include "spinlock.h"
3509
3510 // Interrupt descriptor table (shared by all CPUs).
3511 struct gatedesc idt[256];
3512 extern uint vectors[]; // in vectors.S: array of 256 entry pointers
3513 struct spinlock tickslock;
3514 uint ticks;
3515
3516 void
3517 tvinit(void)
3518 {
3519     int i;
3520
3521     for(i = 0; i < 256; i++)
3522         SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
3523     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);
3524
3525     initlock(&tickslock, "time");
3526 }
3527
3528 void
3529 idtinit(void)
3530 {
3531     lidt(idt, sizeof(idt));
3532 }
3533
3534
3535
3536
3537
3538
3539
3540
3541
3542
3543
3544
3545
3546
3547
3548
3549

```

```

3550 void
3551 trap(struct trapframe *tf)
3552 {
3553     if(tf->trapno == T_SYSCALL){
3554         if(proc->killed)
3555             exit();
3556         proc->tf = tf;
3557         syscall();
3558         if(proc->killed)
3559             exit();
3560         return;
3561     }
3562
3563     switch(tf->trapno){
3564     case T_IRQ0 + IRQ_TIMER:
3565         if(cpu->id == 0){
3566             acquire(&tickslock);
3567             ticks++;
3568             wakeup(&ticks);
3569             release(&tickslock);
3570         }
3571         lapiceoi();
3572         break;
3573     case T_IRQ0 + IRQ_IDE:
3574         ideintr();
3575         lapiceoi();
3576         break;
3577     case T_IRQ0 + IRQ_IDE+1:
3578         // Bochs generates spurious IDE1 interrupts.
3579         break;
3580     case T_IRQ0 + IRQ_KBD:
3581         kbdintr();
3582         lapiceoi();
3583         break;
3584     case T_IRQ0 + IRQ_COM1:
3585         uartintr();
3586         lapiceoi();
3587         break;
3588     case T_IRQ0 + 7:
3589     case T_IRQ0 + IRQ_SPURIOUS:
3590         cprintf("cpu%d: spurious interrupt at %x:%x\n",
3591             cpu->id, tf->cs, tf->eip);
3592         lapiceoi();
3593         break;
3594
3595
3596
3597
3598
3599

```

```

3600 default:
3601     if(proc == 0 || (tf->cs&3) == 0){
3602         // In kernel, it must be our mistake.
3603         cprintf("unexpected trap %d from cpu %d eip %x (cr2=0x%x)\n",
3604             tf->trapno, cpu->id, tf->eip, rcr2());
3605         panic("trap");
3606     }
3607     // In user space, assume process misbehaved.
3608     cprintf("pid %d %s: trap %d err %d on cpu %d "
3609         "eip 0x%x addr 0x%x--kill proc\n",
3610         proc->pid, proc->name, tf->trapno, tf->err, cpu->id, tf->eip,
3611         rcr2());
3612     proc->killed = 1;
3613 }
3614
3615 // Force process exit if it has been killed and is in user space.
3616 // (If it is still executing in the kernel, let it keep running
3617 // until it gets to the regular system call return.)
3618 if(proc && proc->killed && (tf->cs&3) == DPL_USER)
3619     exit();
3620
3621 // Force process to give up CPU on clock tick.
3622 // If interrupts were on while locks held, would need to check nlock.
3623 if(proc && proc->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER)
3624     yield();
3625
3626 // Check if the process has been killed since we yielded
3627 if(proc && proc->killed && (tf->cs&3) == DPL_USER)
3628     exit();
3629 }
3630
3631
3632
3633
3634
3635
3636
3637
3638
3639
3640
3641
3642
3643
3644
3645
3646
3647
3648
3649

```

```

3650 // System call numbers
3651 #define SYS_fork    1
3652 #define SYS_exit    2
3653 #define SYS_wait    3
3654 #define SYS_pipe    4
3655 #define SYS_read    5
3656 #define SYS_kill    6
3657 #define SYS_exec    7
3658 #define SYS_fstat   8
3659 #define SYS_chdir   9
3660 #define SYS_dup    10
3661 #define SYS_getpid  11
3662 #define SYS_sbrk   12
3663 #define SYS_sleep  13
3664 #define SYS_uptime 14
3665 #define SYS_open   15
3666 #define SYS_write  16
3667 #define SYS_mknod  17
3668 #define SYS_unlink 18
3669 #define SYS_link   19
3670 #define SYS_mkdir  20
3671 #define SYS_close  21
3672 #define SYS_halt   22
3673 #define SYS_date   23
3674 #define SYS_getuid 24
3675 #define SYS_getgid 25
3676 #define SYS_getppid 26
3677 #define SYS_setuid 27
3678 #define SYS_setgid 28
3679 #define SYS_getprocs 29
3680
3681
3682
3683
3684
3685
3686
3687
3688
3689
3690
3691
3692
3693
3694
3695
3696
3697
3698
3699

```

```

3700 #include "types.h"
3701 #include "defs.h"
3702 #include "param.h"
3703 #include "memlayout.h"
3704 #include "mmu.h"
3705 #include "proc.h"
3706 #include "x86.h"
3707 #include "syscall.h"
3708
3709 // User code makes a system call with INT T_SYSCALL.
3710 // System call number in %eax.
3711 // Arguments on the stack, from the user call to the C
3712 // library system call function. The saved user %esp points
3713 // to a saved program counter, and then the first argument.
3714
3715 // Fetch the int at addr from the current process.
3716 int
3717 fetchint(uint addr, int *ip)
3718 {
3719     if(addr >= proc->sz || addr+4 > proc->sz)
3720         return -1;
3721     *ip = *(int*)(addr);
3722     return 0;
3723 }
3724
3725 // Fetch the nul-terminated string at addr from the current process.
3726 // Doesn't actually copy the string - just sets *pp to point at it.
3727 // Returns length of string, not including nul.
3728 int
3729 fetchstr(uint addr, char **pp)
3730 {
3731     char *s, *ep;
3732
3733     if(addr >= proc->sz)
3734         return -1;
3735     *pp = (char*)addr;
3736     ep = (char*)proc->sz;
3737     for(s = *pp; s < ep; s++)
3738         if(*s == 0)
3739             return s - *pp;
3740     return -1;
3741 }
3742
3743 // Fetch the nth 32-bit system call argument.
3744 int
3745 argint(int n, int *ip)
3746 {
3747     return fetchint(proc->tf->esp + 4 + 4*n, ip);
3748 }
3749

```

```

3750 // Fetch the nth word-sized system call argument as a pointer
3751 // to a block of memory of size n bytes. Check that the pointer
3752 // lies within the process address space.
3753 int
3754 argptr(int n, char **pp, int size)
3755 {
3756     int i;
3757
3758     if(argint(n, &i) < 0)
3759         return -1;
3760     if((uint)i >= proc->sz || (uint)i+size > proc->sz)
3761         return -1;
3762     *pp = (char*)i;
3763     return 0;
3764 }
3765
3766 // Fetch the nth word-sized system call argument as a string pointer.
3767 // Check that the pointer is valid and the string is nul-terminated.
3768 // (There is no shared writable memory, so the string can't change
3769 // between this check and being used by the kernel.)
3770 int
3771 argstr(int n, char **pp)
3772 {
3773     int addr;
3774     if(argint(n, &addr) < 0)
3775         return -1;
3776     return fetchstr(addr, pp);
3777 }
3778
3779 extern int sys_chdir(void);
3780 extern int sys_close(void);
3781 extern int sys_dup(void);
3782 extern int sys_exec(void);
3783 extern int sys_exit(void);
3784 extern int sys_fork(void);
3785 extern int sys_fstat(void);
3786 extern int sys_getpid(void);
3787 extern int sys_kill(void);
3788 extern int sys_link(void);
3789 extern int sys_mkdir(void);
3790 extern int sys_mknod(void);
3791 extern int sys_open(void);
3792 extern int sys_pipe(void);
3793 extern int sys_read(void);
3794 extern int sys_sbrk(void);
3795 extern int sys_sleep(void);
3796 extern int sys_unlink(void);
3797 extern int sys_wait(void);
3798 extern int sys_write(void);
3799 extern int sys_uptime(void);

```

```

3800 extern int sys_halt(void);
3801 extern int sys_date(void);
3802 extern int sys_getuid(void);
3803 extern int sys_getgid(void);
3804 extern int sys_getppid(void);
3805 extern int sys_setuid(void);
3806 extern int sys_setgid(void);
3807 extern int sys_getprocs(void);
3808
3809 static int (*syscalls[])(void) = {
3810 [SYS_fork] sys_fork,
3811 [SYS_exit] sys_exit,
3812 [SYS_wait] sys_wait,
3813 [SYS_pipe] sys_pipe,
3814 [SYS_read] sys_read,
3815 [SYS_kill] sys_kill,
3816 [SYS_exec] sys_exec,
3817 [SYS_fstat] sys_fstat,
3818 [SYS_chdir] sys_chdir,
3819 [SYS_dup] sys_dup,
3820 [SYS_getpid] sys_getpid,
3821 [SYS_sbrk] sys_sbrk,
3822 [SYS_sleep] sys_sleep,
3823 [SYS_uptime] sys_uptime,
3824 [SYS_open] sys_open,
3825 [SYS_write] sys_write,
3826 [SYS_mknod] sys_mknod,
3827 [SYS_unlink] sys_unlink,
3828 [SYS_link] sys_link,
3829 [SYS_mkdir] sys_mkdir,
3830 [SYS_close] sys_close,
3831 [SYS_halt] sys_halt,
3832 [SYS_date] sys_date,
3833 [SYS_getuid] sys_getuid,
3834 [SYS_getgid] sys_getgid,
3835 [SYS_getppid] sys_getppid,
3836 [SYS_setuid] sys_setuid,
3837 [SYS_setgid] sys_setgid,
3838 [SYS_getprocs] sys_getprocs,
3839 };
3840
3841
3842
3843
3844
3845
3846
3847
3848
3849

```

```

3850 const char *syscallname[] = {
3851
3852 [SYS_fork] "fork",
3853 [SYS_exit] "exit",
3854 [SYS_wait] "wait",
3855 [SYS_pipe] "pipe",
3856 [SYS_read] "read",
3857 [SYS_kill] "kill",
3858 [SYS_exec] "exec",
3859 [SYS_fstat] "fstat",
3860 [SYS_chdir] "chdir",
3861 [SYS_dup] "dup",
3862 [SYS_getpid] "getpid",
3863 [SYS_sbrk] "sbrk",
3864 [SYS_sleep] "sleep",
3865 [SYS_uptime] "uptime",
3866 [SYS_open] "open",
3867 [SYS_write] "write",
3868 [SYS_mknod] "mknod",
3869 [SYS_unlink] "unlink",
3870 [SYS_link] "link",
3871 [SYS_mkdir] "mkdir",
3872 [SYS_close] "close",
3873 [SYS_halt] "halt",
3874 [SYS_date] "date",
3875 [SYS_getuid] "getuid",
3876 [SYS_getgid] "getgid",
3877 [SYS_getppid] "getppid",
3878 [SYS_setuid] "setuid",
3879 [SYS_setgid] "setgid",
3880 [SYS_getprocs] "getprocs",
3881
3882
3883 };
3884
3885
3886 void
3887 syscall(void)
3888 {
3889     int num;
3890
3891     num = proc->tf->eax;
3892     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
3893         proc->tf->eax = syscalls[num]();
3894
3895         // cprintf("%s %d %d %d %d %d %d %d\n", syscallname[num], proc->tf->eax, pr
3896
3897
3898
3899

```

```

3900 } else {
3901     cprintf("%d %s: unknown sys call %d\n",
3902             proc->pid, proc->name, num);
3903     proc->tf->eax = -1;
3904 }
3905 }
3906
3907
3908
3909
3910
3911
3912
3913
3914
3915
3916
3917
3918
3919
3920
3921
3922
3923
3924
3925
3926
3927
3928
3929
3930
3931
3932
3933
3934
3935
3936
3937
3938
3939
3940
3941
3942
3943
3944
3945
3946
3947
3948
3949

```

```

3950 #include "types.h"
3951 #include "x86.h"
3952 #include "defs.h"
3953 #include "date.h"
3954 #include "param.h"
3955 #include "memlayout.h"
3956 #include "mmu.h"
3957 #include "proc.h"
3958 #include "ps.h"
3959
3960 int
3961 sys_fork(void)
3962 {
3963     return fork();
3964 }
3965
3966 int
3967 sys_exit(void)
3968 {
3969     exit();
3970     return 0; // not reached
3971 }
3972
3973 int
3974 sys_wait(void)
3975 {
3976     return wait();
3977 }
3978
3979 int
3980 sys_kill(void)
3981 {
3982     int pid;
3983
3984     if(argint(0, &pid) < 0)
3985         return -1;
3986     return kill(pid);
3987 }
3988
3989 int
3990 sys_getpid(void)
3991 {
3992     return proc->pid;
3993 }
3994
3995 int
3996 sys_getuid(void)
3997 {
3998     return getuid();
3999 }

```

```

4000 int
4001 sys_getgid(void)
4002 {
4003     return getgid();
4004 }
4005
4006 int
4007 sys_getppid(void)
4008 {
4009     return proc->ppid;
4010 }
4011
4012 int
4013 sys_setuid(void)
4014 {
4015     int uid;
4016     if(argint(0, &uid) < 0)
4017         return -1;
4018
4019     return setuid(uid);
4020 }
4021
4022 int
4023 sys_setgid(void)
4024 {
4025     int gid;
4026     if(argint(0, &gid) < 0)
4027         return -1;
4028
4029     return setgid(gid);
4030 }
4031
4032 int
4033 sys_sbrk(void)
4034 {
4035     int addr;
4036     int n;
4037
4038     if(argint(0, &n) < 0)
4039         return -1;
4040     addr = proc->sz;
4041     if(growproc(n) < 0)
4042         return -1;
4043     return addr;
4044 }
4045
4046
4047
4048
4049

```

```

4050 int
4051 sys_sleep(void)
4052 {
4053     int n;
4054     uint ticks0;
4055
4056     if(argint(0, &n) < 0)
4057         return -1;
4058     acquire(&tickslock);
4059     ticks0 = ticks;
4060     while(ticks - ticks0 < n){
4061         if(proc->killed){
4062             release(&tickslock);
4063             return -1;
4064         }
4065         sleep(&ticks, &tickslock);
4066     }
4067     release(&tickslock);
4068     return 0;
4069 }
4070
4071 // return how many clock tick interrupts have occurred
4072 // since start.
4073 int
4074 sys_uptime(void)
4075 {
4076     uint xticks;
4077
4078     acquire(&tickslock);
4079     xticks = ticks;
4080     release(&tickslock);
4081     return xticks;
4082 }
4083
4084 //Turn of the computer
4085 int sys_halt(void){
4086     cprintf("Shutting down ...\n");
4087     outw(0xB004, 0x0 | 0x2000);
4088     return 0;
4089 }
4090
4091
4092
4093
4094
4095
4096
4097
4098
4099

```



```
4100 //Date
4101 int
4102 sys_date(void)
4103 {
4104     struct rtcdate *d;
4105     if(argptr(0, (void*)&d, sizeof(*d)) < 0)
4106         return -1;
4107     cmostime(d);
4108     return 0;
4109 }
4110 //getprocs
4111 int
4112 sys_getprocs(void)
4113 {
4114     int max;
4115     struct uproc *table;
4116     if(argint(0, &max) < 0)
4117         return -1;
4118     if(argptr(1, (void*)&table, sizeof(*table)) < 0)
4119         return -1;
4120     return getprocs(max, table);
4121 }
4122
4123
4124
4125
4126
4127
4128
4129
4130
4131
4132
4133
4134
4135
4136
4137
4138
4139
4140
4141
4142
4143
4144
4145
4146
4147
4148
4149
```

```
4150 struct buf {
4151     int flags;
4152     uint dev;
4153     uint blockno;
4154     struct buf *prev; // LRU cache list
4155     struct buf *next;
4156     struct buf *qnext; // disk queue
4157     uchar data[BSIZE];
4158 };
4159 #define B_BUSY 0x1 // buffer is locked by some process
4160 #define B_VALID 0x2 // buffer has been read from disk
4161 #define B_DIRTY 0x4 // buffer needs to be written to disk
4162
4163
4164
4165
4166
4167
4168
4169
4170
4171
4172
4173
4174
4175
4176
4177
4178
4179
4180
4181
4182
4183
4184
4185
4186
4187
4188
4189
4190
4191
4192
4193
4194
4195
4196
4197
4198
4199
```

```
4200 #define O_RDONLY 0x000
4201 #define O_WRONLY 0x001
4202 #define O_RDWR 0x002
4203 #define O_CREATE 0x200
4204
4205
4206
4207
4208
4209
4210
4211
4212
4213
4214
4215
4216
4217
4218
4219
4220
4221
4222
4223
4224
4225
4226
4227
4228
4229
4230
4231
4232
4233
4234
4235
4236
4237
4238
4239
4240
4241
4242
4243
4244
4245
4246
4247
4248
4249
```

```
4250 #define T_DIR 1 // Directory
4251 #define T_FILE 2 // File
4252 #define T_DEV 3 // Device
4253
4254 struct stat {
4255     short type; // Type of file
4256     int dev; // File system's disk device
4257     uint ino; // Inode number
4258     short nlink; // Number of links to file
4259     uint size; // Size of file in bytes
4260 };
4261
4262
4263
4264
4265
4266
4267
4268
4269
4270
4271
4272
4273
4274
4275
4276
4277
4278
4279
4280
4281
4282
4283
4284
4285
4286
4287
4288
4289
4290
4291
4292
4293
4294
4295
4296
4297
4298
4299
```

```

4300 // On-disk file system format.
4301 // Both the kernel and user programs use this header file.
4302
4303
4304 #define ROOTINO 1 // root i-number
4305 #define BSIZE 512 // block size
4306
4307 // Disk layout:
4308 // [ boot block | super block | log | inode blocks | free bit map | data blocks ]
4309 //
4310 // mkfs computes the super block and builds an initial file system. The super block
4311 // the disk layout:
4312 struct superblock {
4313     uint size; // Size of file system image (blocks)
4314     uint nblocks; // Number of data blocks
4315     uint ninodes; // Number of inodes.
4316     uint nlog; // Number of log blocks
4317     uint logstart; // Block number of first log block
4318     uint inodestart; // Block number of first inode block
4319     uint bmapstart; // Block number of first free map block
4320 };
4321
4322 #define NDIRECT 12
4323 #define NINDIRECT (BSIZE / sizeof(uint))
4324 #define MAXFILE (NDIRECT + NINDIRECT)
4325
4326 // On-disk inode structure
4327 struct dinode {
4328     short type; // File type
4329     short major; // Major device number (T_DEV only)
4330     short minor; // Minor device number (T_DEV only)
4331     short nlink; // Number of links to inode in file system
4332     uint size; // Size of file (bytes)
4333     uint addrs[NDIRECT+1]; // Data block addresses
4334 };
4335
4336
4337
4338
4339
4340
4341
4342
4343
4344
4345
4346
4347
4348
4349

```

```

4350 // Inodes per block.
4351 #define IPB (BSIZE / sizeof(struct dinode))
4352
4353 // Block containing inode i
4354 #define IBLOCK(i, sb) ((i) / IPB + sb.inodestart)
4355
4356 // Bitmap bits per block
4357 #define BPB (BSIZE*8)
4358
4359 // Block of free map containing bit for block b
4360 #define BBLOCK(b, sb) (b/BPB + sb.bmapstart)
4361
4362 // Directory is a file containing a sequence of dirent structures.
4363 #define DIRSIZ 14
4364
4365 struct dirent {
4366     ushort inum;
4367     char name[DIRSIZ];
4368 };
4369
4370
4371
4372
4373
4374
4375
4376
4377
4378
4379
4380
4381
4382
4383
4384
4385
4386
4387
4388
4389
4390
4391
4392
4393
4394
4395
4396
4397
4398
4399

```

```
4400 struct file {
4401     enum { FD_NONE, FD_PIPE, FD_INODE } type;
4402     int ref; // reference count
4403     char readable;
4404     char writable;
4405     struct pipe *pipe;
4406     struct inode *ip;
4407     uint off;
4408 };
4409
4410
4411 // in-memory copy of an inode
4412 struct inode {
4413     uint dev;           // Device number
4414     uint inum;          // Inode number
4415     int ref;            // Reference count
4416     int flags;          // I_BUSY, I_VALID
4417
4418     short type;         // copy of disk inode
4419     short major;
4420     short minor;
4421     short nlink;
4422     uint size;
4423     uint addrs[NDIRECT+1];
4424 };
4425 #define I_BUSY 0x1
4426 #define I_VALID 0x2
4427
4428 // table mapping major device number to
4429 // device functions
4430 struct devsw {
4431     int (*read)(struct inode*, char*, int);
4432     int (*write)(struct inode*, char*, int);
4433 };
4434
4435 extern struct devsw devsw[];
4436
4437 #define CONSOLE 1
4438
4439
4440
4441
4442
4443
4444
4445
4446
4447
4448
4449
```

```
4450 // Blank page.
4451
4452
4453
4454
4455
4456
4457
4458
4459
4460
4461
4462
4463
4464
4465
4466
4467
4468
4469
4470
4471
4472
4473
4474
4475
4476
4477
4478
4479
4480
4481
4482
4483
4484
4485
4486
4487
4488
4489
4490
4491
4492
4493
4494
4495
4496
4497
4498
4499
```

```

4500 // Simple PIO-based (non-DMA) IDE driver code.
4501
4502 #include "types.h"
4503 #include "defs.h"
4504 #include "param.h"
4505 #include "memlayout.h"
4506 #include "mmu.h"
4507 #include "proc.h"
4508 #include "x86.h"
4509 #include "traps.h"
4510 #include "spinlock.h"
4511 #include "fs.h"
4512 #include "buf.h"
4513
4514 #define SECTOR_SIZE 512
4515 #define IDE_BSY 0x80
4516 #define IDE_DRDY 0x40
4517 #define IDE_DF 0x20
4518 #define IDE_ERR 0x01
4519
4520 #define IDE_CMD_READ 0x20
4521 #define IDE_CMD_WRITE 0x30
4522
4523 // idequeue points to the buf now being read/written to the disk.
4524 // idequeue->qnext points to the next buf to be processed.
4525 // You must hold idelock while manipulating queue.
4526
4527 static struct spinlock idelock;
4528 static struct buf *idequeue;
4529
4530 static int havdiskl;
4531 static void idestart(struct buf*);
4532
4533 // Wait for IDE disk to become ready.
4534 static int
4535 idewait(int checkerr)
4536 {
4537     int r;
4538     while(((r = inb(0x1f7)) & (IDE_BSY|IDE_DRDY)) != IDE_DRDY)
4539         ;
4540     if(checkerr && (r & (IDE_DF|IDE_ERR)) != 0)
4541         return -1;
4542     return 0;
4543 }
4544
4545
4546
4547
4548
4549

```

```

4550 void
4551 ideinit(void)
4552 {
4553     int i;
4554
4555     initlock(&idelock, "ide");
4556     picenable(IRQ_IDE);
4557     ioapicenable(IRQ_IDE, ncpu - 1);
4558     idewait(0);
4559
4560     // Check if disk 1 is present
4561     outb(0x1f6, 0xe0 | (1<<4));
4562     for(i=0; i<1000; i++){
4563         if(inb(0x1f7) != 0){
4564             havdiskl = 1;
4565             break;
4566         }
4567     }
4568
4569     // Switch back to disk 0.
4570     outb(0x1f6, 0xe0 | (0<<4));
4571 }
4572
4573 // Start the request for b. Caller must hold idelock.
4574 static void
4575 idestart(struct buf *b)
4576 {
4577     if(b == 0)
4578         panic("idestart");
4579     if(b->blockno >= FSSIZE)
4580         panic("incorrect blockno");
4581     int sector_per_block = BSIZE/SECTOR_SIZE;
4582     int sector = b->blockno * sector_per_block;
4583
4584     if (sector_per_block > 7) panic("idestart");
4585
4586     idewait(0);
4587     outb(0x3f6, 0); // generate interrupt
4588     outb(0x1f2, sector_per_block); // number of sectors
4589     outb(0x1f3, sector & 0xff);
4590     outb(0x1f4, (sector >> 8) & 0xff);
4591     outb(0x1f5, (sector >> 16) & 0xff);
4592     outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((sector>>24)&0x0f));
4593     if(b->flags & B_DIRTY){
4594         outb(0x1f7, IDE_CMD_WRITE);
4595         outsl(0x1f0, b->data, BSIZE/4);
4596     } else {
4597         outb(0x1f7, IDE_CMD_READ);
4598     }
4599 }

```

```

4600 // Interrupt handler.
4601 void
4602 ideintr(void)
4603 {
4604     struct buf *b;
4605
4606     // First queued buffer is the active request.
4607     acquire(&idelock);
4608     if((b = idequeue) == 0){
4609         release(&idelock);
4610         // cprintf("spurious IDE interrupt\n");
4611         return;
4612     }
4613     idequeue = b->qnext;
4614
4615     // Read data if needed.
4616     if(!(b->flags & B_DIRTY) && idewait(1) >= 0)
4617         insl(0x1f0, b->data, BSIZE/4);
4618
4619     // Wake process waiting for this buf.
4620     b->flags |= B_VALID;
4621     b->flags &= ~B_DIRTY;
4622     wakeup(b);
4623
4624     // Start disk on next buf in queue.
4625     if(idequeue != 0)
4626         idestart(idequeue);
4627
4628     release(&idelock);
4629 }
4630
4631
4632
4633
4634
4635
4636
4637
4638
4639
4640
4641
4642
4643
4644
4645
4646
4647
4648
4649

```

```

4650 // Sync buf with disk.
4651 // If B_DIRTY is set, write buf to disk, clear B_DIRTY, set B_VALID.
4652 // Else if B_VALID is not set, read buf from disk, set B_VALID.
4653 void
4654 iderw(struct buf *b)
4655 {
4656     struct buf **pp;
4657
4658     if(!(b->flags & B_BUSY))
4659         panic("iderw: buf not busy");
4660     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
4661         panic("iderw: nothing to do");
4662     if(b->dev != 0 && !havedisk1)
4663         panic("iderw: ide disk 1 not present");
4664
4665     acquire(&idelock);
4666
4667     // Append b to idequeue.
4668     b->qnext = 0;
4669     for(pp=&idequeue; *pp; pp=(*pp)->qnext)
4670         ;
4671     *pp = b;
4672
4673     // Start disk if necessary.
4674     if(idequeue == b)
4675         idestart(b);
4676
4677     // Wait for request to finish.
4678     while((b->flags & (B_VALID|B_DIRTY)) != B_VALID){
4679         sleep(b, &idelock);
4680     }
4681
4682     release(&idelock);
4683 }
4684
4685
4686
4687
4688
4689
4690
4691
4692
4693
4694
4695
4696
4697
4698
4699

```

```

4700 // Buffer cache.
4701 //
4702 // The buffer cache is a linked list of buf structures holding
4703 // cached copies of disk block contents. Caching disk blocks
4704 // in memory reduces the number of disk reads and also provides
4705 // a synchronization point for disk blocks used by multiple processes.
4706 //
4707 // Interface:
4708 // * To get a buffer for a particular disk block, call bread.
4709 // * After changing buffer data, call bwrite to write it to disk.
4710 // * When done with the buffer, call brelse.
4711 // * Do not use the buffer after calling brelse.
4712 // * Only one process at a time can use a buffer,
4713 //   so do not keep them longer than necessary.
4714 //
4715 // The implementation uses three state flags internally:
4716 // * B_BUSY: the block has been returned from bread
4717 //   and has not been passed back to brelse.
4718 // * B_VALID: the buffer data has been read from the disk.
4719 // * B_DIRTY: the buffer data has been modified
4720 //   and needs to be written to disk.
4721
4722 #include "types.h"
4723 #include "defs.h"
4724 #include "param.h"
4725 #include "spinlock.h"
4726 #include "fs.h"
4727 #include "buf.h"
4728
4729 struct {
4730   struct spinlock lock;
4731   struct buf buf[NBUF];
4732
4733   // Linked list of all buffers, through prev/next.
4734   // head.next is most recently used.
4735   struct buf head;
4736 } bcache;
4737
4738 void
4739 binit(void)
4740 {
4741   struct buf *b;
4742
4743   initlock(&bcache.lock, "bcache");
4744
4745
4746
4747
4748
4749

```

```

4750 // Create linked list of buffers
4751 bcache.head.prev = &bcache.head;
4752 bcache.head.next = &bcache.head;
4753 for(b = bcache.buf; b < bcache.buf+NBUF; b++){
4754   b->next = bcache.head.next;
4755   b->prev = &bcache.head;
4756   b->dev = -1;
4757   bcache.head.next->prev = b;
4758   bcache.head.next = b;
4759 }
4760 }
4761
4762 // Look through buffer cache for block on device dev.
4763 // If not found, allocate a buffer.
4764 // In either case, return B_BUSY buffer.
4765 static struct buf*
4766 bget(uint dev, uint blockno)
4767 {
4768   struct buf *b;
4769
4770   acquire(&bcache.lock);
4771
4772   loop:
4773   // Is the block already cached?
4774   for(b = bcache.head.next; b != &bcache.head; b = b->next){
4775     if(b->dev == dev && b->blockno == blockno){
4776       if(!(b->flags & B_BUSY)){
4777         b->flags |= B_BUSY;
4778         release(&bcache.lock);
4779         return b;
4780       }
4781       sleep(b, &bcache.lock);
4782       goto loop;
4783     }
4784   }
4785
4786   // Not cached; recycle some non-busy and clean buffer.
4787   // "clean" because B_DIRTY and !B_BUSY means log.c
4788   // hasn't yet committed the changes to the buffer.
4789   for(b = bcache.head.prev; b != &bcache.head; b = b->prev){
4790     if((b->flags & B_BUSY) == 0 && (b->flags & B_DIRTY) == 0){
4791       b->dev = dev;
4792       b->blockno = blockno;
4793       b->flags = B_BUSY;
4794       release(&bcache.lock);
4795       return b;
4796     }
4797   }
4798   panic("bget: no buffers");
4799 }

```

```

4800 // Return a B_BUSY buf with the contents of the indicated block.
4801 struct buf*
4802 bread(uint dev, uint blockno)
4803 {
4804     struct buf *b;
4805
4806     b = bget(dev, blockno);
4807     if(!(b->flags & B_VALID)) {
4808         iderw(b);
4809     }
4810     return b;
4811 }
4812
4813 // Write b's contents to disk. Must be B_BUSY.
4814 void
4815 bwrite(struct buf *b)
4816 {
4817     if((b->flags & B_BUSY) == 0)
4818         panic("bwrite");
4819     b->flags |= B_DIRTY;
4820     iderw(b);
4821 }
4822
4823 // Release a B_BUSY buffer.
4824 // Move to the head of the MRU list.
4825 void
4826 brelse(struct buf *b)
4827 {
4828     if((b->flags & B_BUSY) == 0)
4829         panic("brelse");
4830
4831     acquire(&bcache.lock);
4832
4833     b->next->prev = b->prev;
4834     b->prev->next = b->next;
4835     b->next = bcache.head.next;
4836     b->prev = &bcache.head;
4837     bcache.head.next->prev = b;
4838     bcache.head.next = b;
4839
4840     b->flags &= ~B_BUSY;
4841     wakeup(b);
4842
4843     release(&bcache.lock);
4844 }
4845
4846
4847
4848
4849

```

```

4850 // Blank page.
4851
4852
4853
4854
4855
4856
4857
4858
4859
4860
4861
4862
4863
4864
4865
4866
4867
4868
4869
4870
4871
4872
4873
4874
4875
4876
4877
4878
4879
4880
4881
4882
4883
4884
4885
4886
4887
4888
4889
4890
4891
4892
4893
4894
4895
4896
4897
4898
4899

```



```

4900 #include "types.h"
4901 #include "defs.h"
4902 #include "param.h"
4903 #include "spinlock.h"
4904 #include "fs.h"
4905 #include "buf.h"
4906
4907 // Simple logging that allows concurrent FS system calls.
4908 //
4909 // A log transaction contains the updates of multiple FS system
4910 // calls. The logging system only commits when there are
4911 // no FS system calls active. Thus there is never
4912 // any reasoning required about whether a commit might
4913 // write an uncommitted system call's updates to disk.
4914 //
4915 // A system call should call begin_op()/end_op() to mark
4916 // its start and end. Usually begin_op() just increments
4917 // the count of in-progress FS system calls and returns.
4918 // But if it thinks the log is close to running out, it
4919 // sleeps until the last outstanding end_op() commits.
4920 //
4921 // The log is a physical re-do log containing disk blocks.
4922 // The on-disk log format:
4923 //   header block, containing block #s for block A, B, C, ...
4924 //   block A
4925 //   block B
4926 //   block C
4927 //   ...
4928 // Log appends are synchronous.
4929
4930 // Contents of the header block, used for both the on-disk header block
4931 // and to keep track in memory of logged block# before commit.
4932 struct logheader {
4933   int n;
4934   int block[LOGSIZE];
4935 };
4936
4937 struct log {
4938   struct spinlock lock;
4939   int start;
4940   int size;
4941   int outstanding; // how many FS sys calls are executing.
4942   int committing;  // in commit(), please wait.
4943   int dev;
4944   struct logheader lh;
4945 };
4946
4947
4948
4949

```

```

4950 struct log log;
4951
4952 static void recover_from_log(void);
4953 static void commit();
4954
4955 void
4956 initlog(int dev)
4957 {
4958   if (sizeof(struct logheader) >= BSIZE)
4959     panic("initlog: too big logheader");
4960
4961   struct superblock sb;
4962   initlock(&log.lock, "log");
4963   readsb(dev, &sb);
4964   log.start = sb.logstart;
4965   log.size = sb.nlog;
4966   log.dev = dev;
4967   recover_from_log();
4968 }
4969
4970 // Copy committed blocks from log to their home location
4971 static void
4972 install_trans(void)
4973 {
4974   int tail;
4975
4976   for (tail = 0; tail < log.lh.n; tail++) {
4977     struct buf *lbuf = bread(log.dev, log.start+tail+1); // read log block
4978     struct buf *dbuf = bread(log.dev, log.lh.block[tail]); // read dst
4979     memmove(dbuf->data, lbuf->data, BSIZE); // copy block to dst
4980     bwrite(dbuf); // write dst to disk
4981     brelse(lbuf);
4982     brelse(dbuf);
4983   }
4984 }
4985
4986 // Read the log header from disk into the in-memory log header
4987 static void
4988 read_head(void)
4989 {
4990   struct buf *buf = bread(log.dev, log.start);
4991   struct logheader *lh = (struct logheader *) (buf->data);
4992   int i;
4993   log.lh.n = lh->n;
4994   for (i = 0; i < log.lh.n; i++) {
4995     log.lh.block[i] = lh->block[i];
4996   }
4997   brelse(buf);
4998 }
4999

```

```

5000 // Write in-memory log header to disk.
5001 // This is the true point at which the
5002 // current transaction commits.
5003 static void
5004 write_head(void)
5005 {
5006     struct buf *buf = bread(log.dev, log.start);
5007     struct logheader *hb = (struct logheader *) (buf->data);
5008     int i;
5009     hb->n = log.lh.n;
5010     for (i = 0; i < log.lh.n; i++) {
5011         hb->block[i] = log.lh.block[i];
5012     }
5013     bwrite(buf);
5014     brelse(buf);
5015 }
5016
5017 static void
5018 recover_from_log(void)
5019 {
5020     read_head();
5021     install_trans(); // if committed, copy from log to disk
5022     log.lh.n = 0;
5023     write_head(); // clear the log
5024 }
5025
5026 // called at the start of each FS system call.
5027 void
5028 begin_op(void)
5029 {
5030     acquire(&log.lock);
5031     while(1){
5032         if(log.committing){
5033             sleep(&log, &log.lock);
5034         } else if(log.lh.n + (log.outstanding+1)*MAXOPBLOCKS > LOGSIZE){
5035             // this op might exhaust log space; wait for commit.
5036             sleep(&log, &log.lock);
5037         } else {
5038             log.outstanding += 1;
5039             release(&log.lock);
5040             break;
5041         }
5042     }
5043 }
5044
5045
5046
5047
5048
5049

```

```

5050 // called at the end of each FS system call.
5051 // commits if this was the last outstanding operation.
5052 void
5053 end_op(void)
5054 {
5055     int do_commit = 0;
5056
5057     acquire(&log.lock);
5058     log.outstanding -= 1;
5059     if(log.committing)
5060         panic("log.committing");
5061     if(log.outstanding == 0){
5062         do_commit = 1;
5063         log.committing = 1;
5064     } else {
5065         // begin_op() may be waiting for log space.
5066         wakeup(&log);
5067     }
5068     release(&log.lock);
5069
5070     if(do_commit){
5071         // call commit w/o holding locks, since not allowed
5072         // to sleep with locks.
5073         commit();
5074         acquire(&log.lock);
5075         log.committing = 0;
5076         wakeup(&log);
5077         release(&log.lock);
5078     }
5079 }
5080
5081 // Copy modified blocks from cache to log.
5082 static void
5083 write_log(void)
5084 {
5085     int tail;
5086
5087     for (tail = 0; tail < log.lh.n; tail++) {
5088         struct buf *to = bread(log.dev, log.start+tail+1); // log block
5089         struct buf *from = bread(log.dev, log.lh.block[tail]); // cache block
5090         memmove(to->data, from->data, BSIZE);
5091         bwrite(to); // write the log
5092         brelse(from);
5093         brelse(to);
5094     }
5095 }
5096
5097
5098
5099

```

```

5100 static void
5101 commit()
5102 {
5103     if (log.lh.n > 0) {
5104         write_log(); // Write modified blocks from cache to log
5105         write_head(); // Write header to disk -- the real commit
5106         install_trans(); // Now install writes to home locations
5107         log.lh.n = 0;
5108         write_head(); // Erase the transaction from the log
5109     }
5110 }
5111
5112 // Caller has modified b->data and is done with the buffer.
5113 // Record the block number and pin in the cache with B_DIRTY.
5114 // commit()/write_log() will do the disk write.
5115 //
5116 // log_write() replaces bwrite(); a typical use is:
5117 //   bp = bread(...)
5118 //   modify bp->data[]
5119 //   log_write(bp)
5120 //   brelse(bp)
5121 void
5122 log_write(struct buf *b)
5123 {
5124     int i;
5125
5126     if (log.lh.n >= LOGSIZE || log.lh.n >= log.size - 1)
5127         panic("too big a transaction");
5128     if (log.outstanding < 1)
5129         panic("log_write outside of trans");
5130
5131     acquire(&log.lock);
5132     for (i = 0; i < log.lh.n; i++) {
5133         if (log.lh.block[i] == b->blockno) // log absorbtion
5134             break;
5135     }
5136     log.lh.block[i] = b->blockno;
5137     if (i == log.lh.n)
5138         log.lh.n++;
5139     b->flags |= B_DIRTY; // prevent eviction
5140     release(&log.lock);
5141 }
5142
5143
5144
5145
5146
5147
5148
5149

```

```

5150 // File system implementation. Five layers:
5151 //   + Blocks: allocator for raw disk blocks.
5152 //   + Log: crash recovery for multi-step updates.
5153 //   + Files: inode allocator, reading, writing, metadata.
5154 //   + Directories: inode with special contents (list of other inodes!)
5155 //   + Names: paths like /usr/rtn/xv6/fs.c for convenient naming.
5156 //
5157 // This file contains the low-level file system manipulation
5158 // routines. The (higher-level) system call implementations
5159 // are in sysfile.c.
5160
5161 #include "types.h"
5162 #include "defs.h"
5163 #include "param.h"
5164 #include "stat.h"
5165 #include "mmu.h"
5166 #include "proc.h"
5167 #include "spinlock.h"
5168 #include "fs.h"
5169 #include "buf.h"
5170 #include "file.h"
5171
5172 #define min(a, b) ((a) < (b) ? (a) : (b))
5173 static void itrunc(struct inode*);
5174 struct superblock sb; // there should be one per dev, but we run with one
5175
5176 // Read the super block.
5177 void
5178 readsb(int dev, struct superblock *sb)
5179 {
5180     struct buf *bp;
5181
5182     bp = bread(dev, 1);
5183     memmove(sb, bp->data, sizeof(*sb));
5184     brelse(bp);
5185 }
5186
5187 // Zero a block.
5188 static void
5189 bzero(int dev, int bno)
5190 {
5191     struct buf *bp;
5192
5193     bp = bread(dev, bno);
5194     memset(bp->data, 0, BSIZE);
5195     log_write(bp);
5196     brelse(bp);
5197 }
5198
5199

```

```

5200 // Blocks.
5201
5202 // Allocate a zeroed disk block.
5203 static uint
5204 balloc(uint dev)
5205 {
5206     int b, bi, m;
5207     struct buf *bp;
5208
5209     bp = 0;
5210     for(b = 0; b < sb.size; b += BPB){
5211         bp = bread(dev, BBLOCK(b, sb));
5212         for(bi = 0; bi < BPB && b + bi < sb.size; bi++){
5213             m = 1 << (bi % 8);
5214             if((bp->data[bi/8] & m) == 0){ // Is block free?
5215                 bp->data[bi/8] |= m; // Mark block in use.
5216                 log_write(bp);
5217                 brelse(bp);
5218                 bzero(dev, b + bi);
5219                 return b + bi;
5220             }
5221         }
5222         brelse(bp);
5223     }
5224     panic("balloc: out of blocks");
5225 }
5226
5227 // Free a disk block.
5228 static void
5229 bfree(int dev, uint b)
5230 {
5231     struct buf *bp;
5232     int bi, m;
5233
5234     readsb(dev, &sb);
5235     bp = bread(dev, BBLOCK(b, sb));
5236     bi = b % BPB;
5237     m = 1 << (bi % 8);
5238     if((bp->data[bi/8] & m) == 0)
5239         panic("freeing free block");
5240     bp->data[bi/8] &= ~m;
5241     log_write(bp);
5242     brelse(bp);
5243 }
5244
5245
5246
5247
5248
5249

```

```

5250 // Inodes.
5251 //
5252 // An inode describes a single unnamed file.
5253 // The inode disk structure holds metadata: the file's type,
5254 // its size, the number of links referring to it, and the
5255 // list of blocks holding the file's content.
5256 //
5257 // The inodes are laid out sequentially on disk at
5258 // sb.startinode. Each inode has a number, indicating its
5259 // position on the disk.
5260 //
5261 // The kernel keeps a cache of in-use inodes in memory
5262 // to provide a place for synchronizing access
5263 // to inodes used by multiple processes. The cached
5264 // inodes include book-keeping information that is
5265 // not stored on disk: ip->ref and ip->flags.
5266 //
5267 // An inode and its in-memory representative go through a
5268 // sequence of states before they can be used by the
5269 // rest of the file system code.
5270 //
5271 // * Allocation: an inode is allocated if its type (on disk)
5272 //   is non-zero. ialloc() allocates, iput() frees if
5273 //   the link count has fallen to zero.
5274 //
5275 // * Referencing in cache: an entry in the inode cache
5276 //   is free if ip->ref is zero. Otherwise ip->ref tracks
5277 //   the number of in-memory pointers to the entry (open
5278 //   files and current directories). iget() to find or
5279 //   create a cache entry and increment its ref, iput()
5280 //   to decrement ref.
5281 //
5282 // * Valid: the information (type, size, &c) in an inode
5283 //   cache entry is only correct when the I_VALID bit
5284 //   is set in ip->flags. ilock() reads the inode from
5285 //   the disk and sets I_VALID, while iput() clears
5286 //   I_VALID if ip->ref has fallen to zero.
5287 //
5288 // * Locked: file system code may only examine and modify
5289 //   the information in an inode and its content if it
5290 //   has first locked the inode. The I_BUSY flag indicates
5291 //   that the inode is locked. ilock() sets I_BUSY,
5292 //   while iunlock clears it.
5293 //
5294 // Thus a typical sequence is:
5295 //   ip = iget(dev, inum)
5296 //   ilock(ip)
5297 //   ... examine and modify ip->xxx ...
5298 //   iunlock(ip)
5299 //   iput(ip)

```

```

5300 //
5301 // ilock() is separate from iget() so that system calls can
5302 // get a long-term reference to an inode (as for an open file)
5303 // and only lock it for short periods (e.g., in read()).
5304 // The separation also helps avoid deadlock and races during
5305 // pathname lookup. iget() increments ip->ref so that the inode
5306 // stays cached and pointers to it remain valid.
5307 //
5308 // Many internal file system functions expect the caller to
5309 // have locked the inodes involved; this lets callers create
5310 // multi-step atomic operations.
5311
5312 struct {
5313   struct spinlock lock;
5314   struct inode inode[NINODE];
5315 } icache;
5316
5317 void
5318 iinit(int dev)
5319 {
5320   initlock(&icache.lock, "icache");
5321   readsb(dev, &sb);
5322   cprintf("sb: size %d nblocks %d ninodes %d nlog %d logstart %d inodestart %d\n",
5323           sb.nblocks, sb.ninodes, sb.nlog, sb.logstart, sb.inodestart, sb.bmap);
5324 }
5325
5326 static struct inode* iget(uint dev, uint inum);
5327
5328
5329
5330
5331
5332
5333
5334
5335
5336
5337
5338
5339
5340
5341
5342
5343
5344
5345
5346
5347
5348
5349

```

```

5350 // Allocate a new inode with the given type on device dev.
5351 // A free inode has a type of zero.
5352 struct inode*
5353 ialloc(uint dev, short type)
5354 {
5355   int inum;
5356   struct buf *bp;
5357   struct dinode *dip;
5358
5359   for(inum = 1; inum < sb.ninodes; inum++){
5360     bp = bread(dev, IBLOCK(inum, sb));
5361     dip = (struct dinode*)bp->data + inum%IPB;
5362     if(dip->type == 0){ // a free inode
5363       memset(dip, 0, sizeof(*dip));
5364       dip->type = type;
5365       log_write(bp); // mark it allocated on the disk
5366       brelse(bp);
5367       return iget(dev, inum);
5368     }
5369     brelse(bp);
5370   }
5371   panic("ialloc: no inodes");
5372 }
5373
5374 // Copy a modified in-memory inode to disk.
5375 void
5376 iupdate(struct inode *ip)
5377 {
5378   struct buf *bp;
5379   struct dinode *dip;
5380
5381   bp = bread(ip->dev, IBLOCK(ip->inum, sb));
5382   dip = (struct dinode*)bp->data + ip->inum%IPB;
5383   dip->type = ip->type;
5384   dip->major = ip->major;
5385   dip->minor = ip->minor;
5386   dip->nlink = ip->nlink;
5387   dip->size = ip->size;
5388   memmove(dip->addrs, ip->addrs, sizeof(ip->addrs));
5389   log_write(bp);
5390   brelse(bp);
5391 }
5392
5393
5394
5395
5396
5397
5398
5399

```

```

5400 // Find the inode with number inum on device dev
5401 // and return the in-memory copy. Does not lock
5402 // the inode and does not read it from disk.
5403 static struct inode*
5404 iget(uint dev, uint inum)
5405 {
5406     struct inode *ip, *empty;
5407
5408     acquire(&icache.lock);
5409
5410     // Is the inode already cached?
5411     empty = 0;
5412     for(ip = &icache.inode[0]; ip < &icache.inode[NINODE]; ip++){
5413         if(ip->ref > 0 && ip->dev == dev && ip->inum == inum){
5414             ip->ref++;
5415             release(&icache.lock);
5416             return ip;
5417         }
5418         if(empty == 0 && ip->ref == 0)    // Remember empty slot.
5419             empty = ip;
5420     }
5421
5422     // Recycle an inode cache entry.
5423     if(empty == 0)
5424         panic("iget: no inodes");
5425
5426     ip = empty;
5427     ip->dev = dev;
5428     ip->inum = inum;
5429     ip->ref = 1;
5430     ip->flags = 0;
5431     release(&icache.lock);
5432
5433     return ip;
5434 }
5435
5436 // Increment reference count for ip.
5437 // Returns ip to enable ip = idup(ip1) idiom.
5438 struct inode*
5439 idup(struct inode *ip)
5440 {
5441     acquire(&icache.lock);
5442     ip->ref++;
5443     release(&icache.lock);
5444     return ip;
5445 }
5446
5447
5448
5449

```

```

5450 // Lock the given inode.
5451 // Reads the inode from disk if necessary.
5452 void
5453 ilock(struct inode *ip)
5454 {
5455     struct buf *bp;
5456     struct dinode *dip;
5457
5458     if(ip == 0 || ip->ref < 1)
5459         panic("ilock");
5460
5461     acquire(&icache.lock);
5462     while(ip->flags & I_BUSY)
5463         sleep(ip, &icache.lock);
5464     ip->flags |= I_BUSY;
5465     release(&icache.lock);
5466
5467     if(!(ip->flags & I_VALID)){
5468         bp = bread(ip->dev, IBLOCK(ip->inum, sb));
5469         dip = (struct dinode*)bp->data + ip->inum%IPB;
5470         ip->type = dip->type;
5471         ip->major = dip->major;
5472         ip->minor = dip->minor;
5473         ip->nlink = dip->nlink;
5474         ip->size = dip->size;
5475         memmove(ip->addrs, dip->addrs, sizeof(ip->addrs));
5476         brelse(bp);
5477         ip->flags |= I_VALID;
5478         if(ip->type == 0)
5479             panic("ilock: no type");
5480     }
5481 }
5482
5483 // Unlock the given inode.
5484 void
5485 iunlock(struct inode *ip)
5486 {
5487     if(ip == 0 || !(ip->flags & I_BUSY) || ip->ref < 1)
5488         panic("iunlock");
5489
5490     acquire(&icache.lock);
5491     ip->flags &= ~I_BUSY;
5492     wakeup(ip);
5493     release(&icache.lock);
5494 }
5495
5496
5497
5498
5499

```

```

5500 // Drop a reference to an in-memory inode.
5501 // If that was the last reference, the inode cache entry can
5502 // be recycled.
5503 // If that was the last reference and the inode has no links
5504 // to it, free the inode (and its content) on disk.
5505 // All calls to iput() must be inside a transaction in
5506 // case it has to free the inode.
5507 void
5508 iput(struct inode *ip)
5509 {
5510     acquire(&icache.lock);
5511     if(ip->ref == 1 && (ip->flags & I_INVALID) && ip->nlink == 0){
5512         // inode has no links and no other references: truncate and free.
5513         if(ip->flags & I_BUSY)
5514             panic("iput busy");
5515         ip->flags |= I_BUSY;
5516         release(&icache.lock);
5517         itrunc(ip);
5518         ip->type = 0;
5519         iupdate(ip);
5520         acquire(&icache.lock);
5521         ip->flags = 0;
5522         wakeup(ip);
5523     }
5524     ip->ref--;
5525     release(&icache.lock);
5526 }
5527
5528 // Common idiom: unlock, then put.
5529 void
5530 iunlockput(struct inode *ip)
5531 {
5532     iunlock(ip);
5533     iput(ip);
5534 }
5535
5536
5537
5538
5539
5540
5541
5542
5543
5544
5545
5546
5547
5548
5549

```

```

5550 // Inode content
5551 //
5552 // The content (data) associated with each inode is stored
5553 // in blocks on the disk. The first NDIRECT block numbers
5554 // are listed in ip->addrs[]. The next NINDIRECT blocks are
5555 // listed in block ip->addrs[NDIRECT].
5556
5557 // Return the disk block address of the nth block in inode ip.
5558 // If there is no such block, bmap allocates one.
5559 static uint
5560 bmap(struct inode *ip, uint bn)
5561 {
5562     uint addr, *a;
5563     struct buf *bp;
5564
5565     if(bn < NDIRECT){
5566         if((addr = ip->addrs[bn]) == 0)
5567             ip->addrs[bn] = addr = balloc(ip->dev);
5568         return addr;
5569     }
5570     bn -= NDIRECT;
5571
5572     if(bn < NINDIRECT){
5573         // Load indirect block, allocating if necessary.
5574         if((addr = ip->addrs[NDIRECT]) == 0)
5575             ip->addrs[NDIRECT] = addr = balloc(ip->dev);
5576         bp = bread(ip->dev, addr);
5577         a = (uint*)bp->data;
5578         if((addr = a[bn]) == 0){
5579             a[bn] = addr = balloc(ip->dev);
5580             log_write(bp);
5581         }
5582         brelse(bp);
5583         return addr;
5584     }
5585
5586     panic("bmap: out of range");
5587 }
5588
5589
5590
5591
5592
5593
5594
5595
5596
5597
5598
5599

```

```

5600 // Truncate inode (discard contents).
5601 // Only called when the inode has no links
5602 // to it (no directory entries referring to it)
5603 // and has no in-memory reference to it (is
5604 // not an open file or current directory).
5605 static void
5606 itrunc(struct inode *ip)
5607 {
5608     int i, j;
5609     struct buf *bp;
5610     uint *a;
5611
5612     for(i = 0; i < NDIRECT; i++){
5613         if(ip->addrs[i]){
5614             bfree(ip->dev, ip->addrs[i]);
5615             ip->addrs[i] = 0;
5616         }
5617     }
5618
5619     if(ip->addrs[NDIRECT]){
5620         bp = bread(ip->dev, ip->addrs[NDIRECT]);
5621         a = (uint*)bp->data;
5622         for(j = 0; j < NINDIRECT; j++){
5623             if(a[j])
5624                 bfree(ip->dev, a[j]);
5625         }
5626         brelse(bp);
5627         bfree(ip->dev, ip->addrs[NDIRECT]);
5628         ip->addrs[NDIRECT] = 0;
5629     }
5630
5631     ip->size = 0;
5632     iupdate(ip);
5633 }
5634
5635 // Copy stat information from inode.
5636 void
5637 stati(struct inode *ip, struct stat *st)
5638 {
5639     st->dev = ip->dev;
5640     st->ino = ip->inum;
5641     st->type = ip->type;
5642     st->nlink = ip->nlink;
5643     st->size = ip->size;
5644 }
5645
5646
5647
5648
5649

```

```

5650 // Read data from inode.
5651 int
5652 readi(struct inode *ip, char *dst, uint off, uint n)
5653 {
5654     uint tot, m;
5655     struct buf *bp;
5656
5657     if(ip->type == T_DEV){
5658         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)
5659             return -1;
5660         return devsw[ip->major].read(ip, dst, n);
5661     }
5662
5663     if(off > ip->size || off + n < off)
5664         return -1;
5665     if(off + n > ip->size)
5666         n = ip->size - off;
5667
5668     for(tot=0; tot<n; tot+=m, off+=m, dst+=m){
5669         bp = bread(ip->dev, bmap(ip, off/BSIZE));
5670         m = min(n - tot, BSIZE - off%BSIZE);
5671         memmove(dst, bp->data + off%BSIZE, m);
5672         brelse(bp);
5673     }
5674     return n;
5675 }
5676
5677
5678
5679
5680
5681
5682
5683
5684
5685
5686
5687
5688
5689
5690
5691
5692
5693
5694
5695
5696
5697
5698
5699

```



```

5700 // Write data to inode.
5701 int
5702 writei(struct inode *ip, char *src, uint off, uint n)
5703 {
5704     uint tot, m;
5705     struct buf *bp;
5706
5707     if(ip->type == T_DEV){
5708         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].write)
5709             return -1;
5710         return devsw[ip->major].write(ip, src, n);
5711     }
5712
5713     if(off > ip->size || off + n < off)
5714         return -1;
5715     if(off + n > MAXFILE*BSIZE)
5716         return -1;
5717
5718     for(tot=0; tot<n; tot+=m, off+=m, src+=m){
5719         bp = bread(ip->dev, bmap(ip, off/BSIZE));
5720         m = min(n - tot, BSIZE - off%BSIZE);
5721         memmove(bp->data + off%BSIZE, src, m);
5722         log_write(bp);
5723         brelse(bp);
5724     }
5725
5726     if(n > 0 && off > ip->size){
5727         ip->size = off;
5728         iupdate(ip);
5729     }
5730     return n;
5731 }
5732
5733
5734
5735
5736
5737
5738
5739
5740
5741
5742
5743
5744
5745
5746
5747
5748
5749

```

```

5750 // Directories
5751
5752 int
5753 namecmp(const char *s, const char *t)
5754 {
5755     return strncmp(s, t, DIRSIZ);
5756 }
5757
5758 // Look for a directory entry in a directory.
5759 // If found, set *poff to byte offset of entry.
5760 struct inode*
5761 dirlookup(struct inode *dp, char *name, uint *poff)
5762 {
5763     uint off, inum;
5764     struct dirent de;
5765
5766     if(dp->type != T_DIR)
5767         panic("dirlookup not DIR");
5768
5769     for(off = 0; off < dp->size; off += sizeof(de)){
5770         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
5771             panic("dirlink read");
5772         if(de.inum == 0)
5773             continue;
5774         if(namecmp(name, de.name) == 0){
5775             // entry matches path element
5776             if(poff)
5777                 *poff = off;
5778             inum = de.inum;
5779             return iget(dp->dev, inum);
5780         }
5781     }
5782
5783     return 0;
5784 }
5785
5786
5787
5788
5789
5790
5791
5792
5793
5794
5795
5796
5797
5798
5799

```

```

5800 // Write a new directory entry (name, inum) into the directory dp.
5801 int
5802 dirlink(struct inode *dp, char *name, uint inum)
5803 {
5804     int off;
5805     struct dirent de;
5806     struct inode *ip;
5807
5808     // Check that name is not present.
5809     if((ip = dirlookup(dp, name, 0)) != 0){
5810         iput(ip);
5811         return -1;
5812     }
5813
5814     // Look for an empty dirent.
5815     for(off = 0; off < dp->size; off += sizeof(de)){
5816         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
5817             panic("dirlink read");
5818         if(de.inum == 0)
5819             break;
5820     }
5821
5822     strncpy(de.name, name, DIRSIZ);
5823     de.inum = inum;
5824     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
5825         panic("dirlink");
5826
5827     return 0;
5828 }
5829
5830
5831
5832
5833
5834
5835
5836
5837
5838
5839
5840
5841
5842
5843
5844
5845
5846
5847
5848
5849

```

```

5850 // Paths
5851
5852 // Copy the next path element from path into name.
5853 // Return a pointer to the element following the copied one.
5854 // The returned path has no leading slashes,
5855 // so the caller can check *path=='\0' to see if the name is the last one.
5856 // If no name to remove, return 0.
5857 //
5858 // Examples:
5859 //   skipelem("a/bb/c", name) = "bb/c", setting name = "a"
5860 //   skipelem("///a//bb", name) = "bb", setting name = "a"
5861 //   skipelem("a", name) = "", setting name = "a"
5862 //   skipelem("", name) = skipelem("///", name) = 0
5863 //
5864 static char*
5865 skipelem(char *path, char *name)
5866 {
5867     char *s;
5868     int len;
5869
5870     while(*path == '/')
5871         path++;
5872     if(*path == 0)
5873         return 0;
5874     s = path;
5875     while(*path != '/' && *path != 0)
5876         path++;
5877     len = path - s;
5878     if(len >= DIRSIZ)
5879         memmove(name, s, DIRSIZ);
5880     else {
5881         memmove(name, s, len);
5882         name[len] = 0;
5883     }
5884     while(*path == '/')
5885         path++;
5886     return path;
5887 }
5888
5889
5890
5891
5892
5893
5894
5895
5896
5897
5898
5899

```

```

5900 // Look up and return the inode for a path name.
5901 // If parent != 0, return the inode for the parent and copy the final
5902 // path element into name, which must have room for DIRSIZ bytes.
5903 // Must be called inside a transaction since it calls iput().
5904 static struct inode*
5905 nameex(char *path, int nameparent, char *name)
5906 {
5907     struct inode *ip, *next;
5908
5909     if(*path == '/')
5910         ip = iget(ROOTDEV, ROOTINO);
5911     else
5912         ip = idup(proc->cwd);
5913
5914     while((path = skipelem(path, name)) != 0){
5915         ilock(ip);
5916         if(ip->type != T_DIR){
5917             iunlockput(ip);
5918             return 0;
5919         }
5920         if(nameparent && *path == '\0'){
5921             // Stop one level early.
5922             iunlock(ip);
5923             return ip;
5924         }
5925         if((next = dirlookup(ip, name, 0)) == 0){
5926             iunlockput(ip);
5927             return 0;
5928         }
5929         iunlockput(ip);
5930         ip = next;
5931     }
5932     if(nameparent){
5933         iput(ip);
5934         return 0;
5935     }
5936     return ip;
5937 }
5938
5939 struct inode*
5940 namei(char *path)
5941 {
5942     char name[DIRSIZ];
5943     return nameex(path, 0, name);
5944 }
5945
5946
5947
5948
5949

```

```

5950 struct inode*
5951 nameparent(char *path, char *name)
5952 {
5953     return nameex(path, 1, name);
5954 }
5955
5956
5957
5958
5959
5960
5961
5962
5963
5964
5965
5966
5967
5968
5969
5970
5971
5972
5973
5974
5975
5976
5977
5978
5979
5980
5981
5982
5983
5984
5985
5986
5987
5988
5989
5990
5991
5992
5993
5994
5995
5996
5997
5998
5999

```

```

6000 //
6001 // File descriptors
6002 //
6003
6004 #include "types.h"
6005 #include "defs.h"
6006 #include "param.h"
6007 #include "fs.h"
6008 #include "file.h"
6009 #include "spinlock.h"
6010
6011 struct devsw devsw[NDEV];
6012 struct {
6013     struct spinlock lock;
6014     struct file file[NFILE];
6015 } ftable;
6016
6017 void
6018 fileinit(void)
6019 {
6020     initlock(&ftable.lock, "ftable");
6021 }
6022
6023 // Allocate a file structure.
6024 struct file*
6025 filealloc(void)
6026 {
6027     struct file *f;
6028
6029     acquire(&ftable.lock);
6030     for(f = ftable.file; f < ftable.file + NFILE; f++){
6031         if(f->ref == 0){
6032             f->ref = 1;
6033             release(&ftable.lock);
6034             return f;
6035         }
6036     }
6037     release(&ftable.lock);
6038     return 0;
6039 }
6040
6041
6042
6043
6044
6045
6046
6047
6048
6049

```

```

6050 // Increment ref count for file f.
6051 struct file*
6052 filedup(struct file *f)
6053 {
6054     acquire(&ftable.lock);
6055     if(f->ref < 1)
6056         panic("filedup");
6057     f->ref++;
6058     release(&ftable.lock);
6059     return f;
6060 }
6061
6062 // Close file f. (Decrement ref count, close when reaches 0.)
6063 void
6064 fileclose(struct file *f)
6065 {
6066     struct file ff;
6067
6068     acquire(&ftable.lock);
6069     if(f->ref < 1)
6070         panic("fileclose");
6071     if(--f->ref > 0){
6072         release(&ftable.lock);
6073         return;
6074     }
6075     ff = *f;
6076     f->ref = 0;
6077     f->type = FD_NONE;
6078     release(&ftable.lock);
6079
6080     if(ff.type == FD_PIPE)
6081         pipeclose(ff.pipe, ff.writable);
6082     else if(ff.type == FD_INODE){
6083         begin_op();
6084         iput(ff.ip);
6085         end_op();
6086     }
6087 }
6088
6089
6090
6091
6092
6093
6094
6095
6096
6097
6098
6099

```

```

6100 // Get metadata about file f.
6101 int
6102 filestat(struct file *f, struct stat *st)
6103 {
6104     if(f->type == FD_INODE){
6105         ilock(f->ip);
6106         stati(f->ip, st);
6107         iunlock(f->ip);
6108         return 0;
6109     }
6110     return -1;
6111 }
6112
6113 // Read from file f.
6114 int
6115 fileread(struct file *f, char *addr, int n)
6116 {
6117     int r;
6118
6119     if(f->readable == 0)
6120         return -1;
6121     if(f->type == FD_PIPE)
6122         return piperead(f->pipe, addr, n);
6123     if(f->type == FD_INODE){
6124         ilock(f->ip);
6125         if((r = readi(f->ip, addr, f->off, n)) > 0)
6126             f->off += r;
6127         iunlock(f->ip);
6128         return r;
6129     }
6130     panic("fileread");
6131 }
6132
6133
6134
6135
6136
6137
6138
6139
6140
6141
6142
6143
6144
6145
6146
6147
6148
6149

```

```

6150 // Write to file f.
6151 int
6152 filewrite(struct file *f, char *addr, int n)
6153 {
6154     int r;
6155
6156     if(f->writable == 0)
6157         return -1;
6158     if(f->type == FD_PIPE)
6159         return pipewrite(f->pipe, addr, n);
6160     if(f->type == FD_INODE){
6161         // write a few blocks at a time to avoid exceeding
6162         // the maximum log transaction size, including
6163         // i-node, indirect block, allocation blocks,
6164         // and 2 blocks of slop for non-aligned writes.
6165         // this really belongs lower down, since writei()
6166         // might be writing a device like the console.
6167         int max = ((LOGSIZE-1-1-2) / 2) * 512;
6168         int i = 0;
6169         while(i < n){
6170             int nl = n - i;
6171             if(nl > max)
6172                 nl = max;
6173
6174             begin_op();
6175             ilock(f->ip);
6176             if ((r = writei(f->ip, addr + i, f->off, nl)) > 0)
6177                 f->off += r;
6178             iunlock(f->ip);
6179             end_op();
6180
6181             if(r < 0)
6182                 break;
6183             if(r != nl)
6184                 panic("short filewrite");
6185             i += r;
6186         }
6187         return i == n ? n : -1;
6188     }
6189     panic("filewrite");
6190 }
6191
6192
6193
6194
6195
6196
6197
6198
6199

```

```

6200 //
6201 // File-system system calls.
6202 // Mostly argument checking, since we don't trust
6203 // user code, and calls into file.c and fs.c.
6204 //
6205
6206 #include "types.h"
6207 #include "defs.h"
6208 #include "param.h"
6209 #include "stat.h"
6210 #include "mmu.h"
6211 #include "proc.h"
6212 #include "fs.h"
6213 #include "file.h"
6214 #include "fcntl.h"
6215
6216 // Fetch the nth word-sized system call argument as a file descriptor
6217 // and return both the descriptor and the corresponding struct file.
6218 static int
6219 argfd(int n, int *pfd, struct file **pf)
6220 {
6221     int fd;
6222     struct file *f;
6223
6224     if(argint(n, &fd) < 0)
6225         return -1;
6226     if(fd < 0 || fd >= NOFILE || (f=proc->ofile[fd]) == 0)
6227         return -1;
6228     if(pfd)
6229         *pfd = fd;
6230     if(pf)
6231         *pf = f;
6232     return 0;
6233 }
6234
6235 // Allocate a file descriptor for the given file.
6236 // Takes over file reference from caller on success.
6237 static int
6238 fdalloc(struct file *f)
6239 {
6240     int fd;
6241
6242     for(fd = 0; fd < NOFILE; fd++){
6243         if(proc->ofile[fd] == 0){
6244             proc->ofile[fd] = f;
6245             return fd;
6246         }
6247     }
6248     return -1;
6249 }

```

```

6250 int
6251 sys_dup(void)
6252 {
6253     struct file *f;
6254     int fd;
6255
6256     if(argfd(0, 0, &f) < 0)
6257         return -1;
6258     if((fd=fdalloc(f)) < 0)
6259         return -1;
6260     filedup(f);
6261     return fd;
6262 }
6263
6264 int
6265 sys_read(void)
6266 {
6267     struct file *f;
6268     int n;
6269     char *p;
6270
6271     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
6272         return -1;
6273     return fileread(f, p, n);
6274 }
6275
6276 int
6277 sys_write(void)
6278 {
6279     struct file *f;
6280     int n;
6281     char *p;
6282
6283     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
6284         return -1;
6285     return filewrite(f, p, n);
6286 }
6287
6288 int
6289 sys_close(void)
6290 {
6291     int fd;
6292     struct file *f;
6293
6294     if(argfd(0, &fd, &f) < 0)
6295         return -1;
6296     proc->ofile[fd] = 0;
6297     fileclose(f);
6298     return 0;
6299 }

```

```

6300 int
6301 sys_fstat(void)
6302 {
6303     struct file *f;
6304     struct stat *st;
6305
6306     if(argfd(0, 0, &f) < 0 || argptr(1, (void*)&st, sizeof(*st)) < 0)
6307         return -1;
6308     return filestat(f, st);
6309 }
6310
6311 // Create the path new as a link to the same inode as old.
6312 int
6313 sys_link(void)
6314 {
6315     char name[DIRSIZ], *new, *old;
6316     struct inode *dp, *ip;
6317
6318     if(argstr(0, &old) < 0 || argstr(1, &new) < 0)
6319         return -1;
6320
6321     begin_op();
6322     if((ip = namei(old)) == 0){
6323         end_op();
6324         return -1;
6325     }
6326
6327     ilock(ip);
6328     if(ip->type == T_DIR){
6329         iunlockput(ip);
6330         end_op();
6331         return -1;
6332     }
6333
6334     ip->nlink++;
6335     iupdate(ip);
6336     iunlock(ip);
6337
6338     if((dp = nameiparent(new, name)) == 0)
6339         goto bad;
6340     ilock(dp);
6341     if(dp->dev != ip->dev || dirlink(dp, name, ip->inum) < 0){
6342         iunlockput(dp);
6343         goto bad;
6344     }
6345     iunlockput(dp);
6346     iput(ip);
6347
6348     end_op();
6349

```

```

6350     return 0;
6351
6352 bad:
6353     ilock(ip);
6354     ip->nlink--;
6355     iupdate(ip);
6356     iunlockput(ip);
6357     end_op();
6358     return -1;
6359 }
6360
6361 // Is the directory dp empty except for "." and ".." ?
6362 static int
6363 isdirempty(struct inode *dp)
6364 {
6365     int off;
6366     struct dirent de;
6367
6368     for(off=2*sizeof(de); off<dp->size; off+=sizeof(de)){
6369         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
6370             panic("isdirempty: readi");
6371         if(de.inum != 0)
6372             return 0;
6373     }
6374     return 1;
6375 }
6376
6377
6378
6379
6380
6381
6382
6383
6384
6385
6386
6387
6388
6389
6390
6391
6392
6393
6394
6395
6396
6397
6398
6399

```

```

6400 int
6401 sys_unlink(void)
6402 {
6403     struct inode *ip, *dp;
6404     struct dirent de;
6405     char name[DIRSIZ], *path;
6406     uint off;
6407
6408     if(argstr(0, &path) < 0)
6409         return -1;
6410
6411     begin_op();
6412     if((dp = nameiparent(path, name)) == 0){
6413         end_op();
6414         return -1;
6415     }
6416
6417     ilock(dp);
6418
6419     // Cannot unlink "." or "..".
6420     if(namecmp(name, ".") == 0 || namecmp(name, "..") == 0)
6421         goto bad;
6422
6423     if((ip = dirlookup(dp, name, &off)) == 0)
6424         goto bad;
6425     ilock(ip);
6426
6427     if(ip->nlink < 1)
6428         panic("unlink: nlink < 1");
6429     if(ip->type == T_DIR && !isdirempty(ip)){
6430         iunlockput(ip);
6431         goto bad;
6432     }
6433
6434     memset(&de, 0, sizeof(de));
6435     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
6436         panic("unlink: writei");
6437     if(ip->type == T_DIR){
6438         dp->nlink--;
6439         iupdate(dp);
6440     }
6441     iunlockput(dp);
6442
6443     ip->nlink--;
6444     iupdate(ip);
6445     iunlockput(ip);
6446
6447     end_op();
6448
6449     return 0;

```

```

6450 bad:
6451     iunlockput(dp);
6452     end_op();
6453     return -1;
6454 }
6455
6456 static struct inode*
6457 create(char *path, short type, short major, short minor)
6458 {
6459     uint off;
6460     struct inode *ip, *dp;
6461     char name[DIRSIZ];
6462
6463     if((dp = nameiparent(path, name)) == 0)
6464         return 0;
6465     ilock(dp);
6466
6467     if((ip = dirlookup(dp, name, &off)) != 0){
6468         iunlockput(dp);
6469         ilock(ip);
6470         if(type == T_FILE && ip->type == T_FILE)
6471             return ip;
6472         iunlockput(ip);
6473         return 0;
6474     }
6475
6476     if((ip = ialloc(dp->dev, type)) == 0)
6477         panic("create: ialloc");
6478
6479     ilock(ip);
6480     ip->major = major;
6481     ip->minor = minor;
6482     ip->nlink = 1;
6483     iupdate(ip);
6484
6485     if(type == T_DIR){ // Create . and .. entries.
6486         dp->nlink++; // for ".."
6487         iupdate(dp);
6488         // No ip->nlink++ for ".": avoid cyclic ref count.
6489         if(dirlink(ip, ".", ip->inum) < 0 || dirlink(ip, "..", dp->inum) < 0)
6490             panic("create dots");
6491     }
6492
6493     if(dirlink(dp, name, ip->inum) < 0)
6494         panic("create: dirlink");
6495
6496     iunlockput(dp);
6497
6498     return ip;
6499 }

```



```

6500 int
6501 sys_open(void)
6502 {
6503     char *path;
6504     int fd, omode;
6505     struct file *f;
6506     struct inode *ip;
6507
6508     if(argstr(0, &path) < 0 || argint(1, &omode) < 0)
6509         return -1;
6510
6511     begin_op();
6512
6513     if(omode & O_CREATE){
6514         ip = create(path, T_FILE, 0, 0);
6515         if(ip == 0){
6516             end_op();
6517             return -1;
6518         }
6519     } else {
6520         if((ip = namei(path)) == 0){
6521             end_op();
6522             return -1;
6523         }
6524         ilock(ip);
6525         if(ip->type == T_DIR && omode != O_RDONLY){
6526             iunlockput(ip);
6527             end_op();
6528             return -1;
6529         }
6530     }
6531
6532     if((f = filealloc()) == 0 || (fd = fdalloc(f)) < 0){
6533         if(f)
6534             fileclose(f);
6535         iunlockput(ip);
6536         end_op();
6537         return -1;
6538     }
6539     iunlock(ip);
6540     end_op();
6541
6542     f->type = FD_INODE;
6543     f->ip = ip;
6544     f->off = 0;
6545     f->readable = !(omode & O_WRONLY);
6546     f->writable = (omode & O_WRONLY) || (omode & O_RDWR);
6547     return fd;
6548 }
6549

```

```

6550 int
6551 sys_mkdir(void)
6552 {
6553     char *path;
6554     struct inode *ip;
6555
6556     begin_op();
6557     if(argstr(0, &path) < 0 || (ip = create(path, T_DIR, 0, 0)) == 0){
6558         end_op();
6559         return -1;
6560     }
6561     iunlockput(ip);
6562     end_op();
6563     return 0;
6564 }
6565
6566 int
6567 sys_mknod(void)
6568 {
6569     struct inode *ip;
6570     char *path;
6571     int len;
6572     int major, minor;
6573
6574     begin_op();
6575     if((len=argstr(0, &path)) < 0 ||
6576         argint(1, &major) < 0 ||
6577         argint(2, &minor) < 0 ||
6578         (ip = create(path, T_DEV, major, minor)) == 0){
6579         end_op();
6580         return -1;
6581     }
6582     iunlockput(ip);
6583     end_op();
6584     return 0;
6585 }
6586
6587
6588
6589
6590
6591
6592
6593
6594
6595
6596
6597
6598
6599

```

```

6600 int
6601 sys_chdir(void)
6602 {
6603     char *path;
6604     struct inode *ip;
6605
6606     begin_op();
6607     if(argstr(0, &path) < 0 || (ip = namei(path)) == 0){
6608         end_op();
6609         return -1;
6610     }
6611     ilock(ip);
6612     if(ip->type != T_DIR){
6613         iunlockput(ip);
6614         end_op();
6615         return -1;
6616     }
6617     iunlock(ip);
6618     iput(proc->cwd);
6619     end_op();
6620     proc->cwd = ip;
6621     return 0;
6622 }
6623
6624 int
6625 sys_exec(void)
6626 {
6627     char *path, *argv[MAXARG];
6628     int i;
6629     uint uargv, uarg;
6630
6631     if(argstr(0, &path) < 0 || argint(1, (int*)&uargv) < 0){
6632         return -1;
6633     }
6634     memset(argv, 0, sizeof(argv));
6635     for(i=0;; i++){
6636         if(i >= NELEM(argv))
6637             return -1;
6638         if(fetchint(uargv+4*i, (int*)&uarg) < 0)
6639             return -1;
6640         if(uarg == 0){
6641             argv[i] = 0;
6642             break;
6643         }
6644         if(fetchstr(uarg, &argv[i]) < 0)
6645             return -1;
6646     }
6647     return exec(path, argv);
6648 }
6649

```

```

6650 int
6651 sys_pipe(void)
6652 {
6653     int *fd;
6654     struct file *rf, *wf;
6655     int fd0, fd1;
6656
6657     if(argptr(0, (void*)&fd, 2*sizeof(fd[0])) < 0)
6658         return -1;
6659     if(pipealloc(&rf, &wf) < 0)
6660         return -1;
6661     fd0 = -1;
6662     if((fd0 = fdalloc(rf)) < 0 || (fd1 = fdalloc(wf)) < 0){
6663         if(fd0 >= 0)
6664             proc->ofile[fd0] = 0;
6665         fileclose(rf);
6666         fileclose(wf);
6667         return -1;
6668     }
6669     fd[0] = fd0;
6670     fd[1] = fd1;
6671     return 0;
6672 }
6673
6674
6675
6676
6677
6678
6679
6680
6681
6682
6683
6684
6685
6686
6687
6688
6689
6690
6691
6692
6693
6694
6695
6696
6697
6698
6699

```

```

6700 #include "types.h"
6701 #include "param.h"
6702 #include "memlayout.h"
6703 #include "mmu.h"
6704 #include "proc.h"
6705 #include "defs.h"
6706 #include "x86.h"
6707 #include "elf.h"
6708
6709 int
6710 exec(char *path, char **argv)
6711 {
6712     char *s, *last;
6713     int i, off;
6714     uint argc, sz, sp, ustack[3+MAXARG+1];
6715     struct elfhdr elf;
6716     struct inode *ip;
6717     struct proghdr ph;
6718     pde_t *pgdir, *oldpgdir;
6719
6720     begin_op();
6721     if((ip = namei(path)) == 0){
6722         end_op();
6723         return -1;
6724     }
6725     ilock(ip);
6726     pgdir = 0;
6727
6728     // Check ELF header
6729     if(readi(ip, (char*)&elf, 0, sizeof(elf)) < sizeof(elf))
6730         goto bad;
6731     if(elf.magic != ELF_MAGIC)
6732         goto bad;
6733
6734     if((pgdir = setupkvm()) == 0)
6735         goto bad;
6736
6737     // Load program into memory.
6738     sz = 0;
6739     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
6740         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
6741             goto bad;
6742         if(ph.type != ELF_PROG_LOAD)
6743             continue;
6744         if(ph.memsz < ph.filesz)
6745             goto bad;
6746         if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
6747             goto bad;
6748         if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
6749             goto bad;

```

```

6750     }
6751     iunlockput(ip);
6752     end_op();
6753     ip = 0;
6754
6755     // Allocate two pages at the next page boundary.
6756     // Make the first inaccessible. Use the second as the user stack.
6757     sz = PGROUNDUP(sz);
6758     if((sz = allocuvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
6759         goto bad;
6760     clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
6761     sp = sz;
6762
6763     // Push argument strings, prepare rest of stack in ustack.
6764     for(argc = 0; argv[argc]; argc++) {
6765         if(argc >= MAXARG)
6766             goto bad;
6767         sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
6768         if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
6769             goto bad;
6770         ustack[3+argc] = sp;
6771     }
6772     ustack[3+argc] = 0;
6773
6774     ustack[0] = 0xffffffff; // fake return PC
6775     ustack[1] = argc;
6776     ustack[2] = sp - (argc+1)*4; // argv pointer
6777
6778     sp -= (3+argc+1) * 4;
6779     if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
6780         goto bad;
6781
6782     // Save program name for debugging.
6783     for(last=s=path; *s; s++)
6784         if(*s == '/')
6785             last = s+1;
6786     safestrcpy(proc->name, last, sizeof(proc->name));
6787
6788     // Commit to the user image.
6789     oldpgdir = proc->pgdir;
6790     proc->pgdir = pgdir;
6791     proc->sz = sz;
6792     proc->tf->eip = elf.entry; // main
6793     proc->tf->esp = sp;
6794     switchuvm(proc);
6795     freevm(oldpgdir);
6796     return 0;
6797
6798
6799

```

```

6800 bad:
6801   if(pgdir)
6802       freevm(pgdir);
6803   if(ip){
6804       iunlockput(ip);
6805       end_op();
6806   }
6807   return -1;
6808 }
6809
6810
6811
6812
6813
6814
6815
6816
6817
6818
6819
6820
6821
6822
6823
6824
6825
6826
6827
6828
6829
6830
6831
6832
6833
6834
6835
6836
6837
6838
6839
6840
6841
6842
6843
6844
6845
6846
6847
6848
6849

```

```

6850 #include "types.h"
6851 #include "defs.h"
6852 #include "param.h"
6853 #include "mmu.h"
6854 #include "proc.h"
6855 #include "fs.h"
6856 #include "file.h"
6857 #include "spinlock.h"
6858
6859 #define PIPESIZE 512
6860
6861 struct pipe {
6862     struct spinlock lock;
6863     char data[PIPESIZE];
6864     uint nread;    // number of bytes read
6865     uint nwrite;   // number of bytes written
6866     int readopen;  // read fd is still open
6867     int writeopen; // write fd is still open
6868 };
6869
6870 int
6871 pipealloc(struct file **f0, struct file **f1)
6872 {
6873     struct pipe *p;
6874
6875     p = 0;
6876     *f0 = *f1 = 0;
6877     if((*f0 = filealloc()) == 0 || (*f1 = filealloc()) == 0)
6878         goto bad;
6879     if((p = (struct pipe*)kalloc()) == 0)
6880         goto bad;
6881     p->readopen = 1;
6882     p->writeopen = 1;
6883     p->nwrite = 0;
6884     p->nread = 0;
6885     initlock(&p->lock, "pipe");
6886     (*f0)->type = FD_PIPE;
6887     (*f0)->readable = 1;
6888     (*f0)->writable = 0;
6889     (*f0)->pipe = p;
6890     (*f1)->type = FD_PIPE;
6891     (*f1)->readable = 0;
6892     (*f1)->writable = 1;
6893     (*f1)->pipe = p;
6894     return 0;
6895
6896
6897
6898
6899

```

```

6900 bad:
6901     if(p)
6902         kfree((char*)p);
6903     if(*f0)
6904         fileclose(*f0);
6905     if(*f1)
6906         fileclose(*f1);
6907     return -1;
6908 }
6909
6910 void
6911 pipeclose(struct pipe *p, int writable)
6912 {
6913     acquire(&p->lock);
6914     if(writable){
6915         p->writeopen = 0;
6916         wakeup(&p->nread);
6917     } else {
6918         p->readopen = 0;
6919         wakeup(&p->nwrite);
6920     }
6921     if(p->readopen == 0 && p->writeopen == 0){
6922         release(&p->lock);
6923         kfree((char*)p);
6924     } else
6925         release(&p->lock);
6926 }
6927
6928
6929 int
6930 pipewrite(struct pipe *p, char *addr, int n)
6931 {
6932     int i;
6933
6934     acquire(&p->lock);
6935     for(i = 0; i < n; i++){
6936         while(p->nwrite == p->nread + PIPESIZE){
6937             if(p->readopen == 0 || proc->killed){
6938                 release(&p->lock);
6939                 return -1;
6940             }
6941             wakeup(&p->nread);
6942             sleep(&p->nwrite, &p->lock);
6943         }
6944         p->data[p->nwrite++ % PIPESIZE] = addr[i];
6945     }
6946     wakeup(&p->nread);
6947     release(&p->lock);
6948     return n;
6949 }

```

```

6950 int
6951 piperead(struct pipe *p, char *addr, int n)
6952 {
6953     int i;
6954
6955     acquire(&p->lock);
6956     while(p->nread == p->nwrite && p->writeopen){
6957         if(proc->killed){
6958             release(&p->lock);
6959             return -1;
6960         }
6961         sleep(&p->nread, &p->lock);
6962     }
6963     for(i = 0; i < n; i++){
6964         if(p->nread == p->nwrite)
6965             break;
6966         addr[i] = p->data[p->nread++ % PIPESIZE];
6967     }
6968     wakeup(&p->nwrite);
6969     release(&p->lock);
6970     return i;
6971 }
6972
6973
6974
6975
6976
6977
6978
6979
6980
6981
6982
6983
6984
6985
6986
6987
6988
6989
6990
6991
6992
6993
6994
6995
6996
6997
6998
6999

```

```

7000 #include "types.h"
7001 #include "x86.h"
7002
7003 void*
7004 memset(void *dst, int c, uint n)
7005 {
7006     if ((int)dst%4 == 0 && n%4 == 0){
7007         c &= 0xFF;
7008         stosl(dst, (c<<24)|(c<<16)|(c<<8)|c, n/4);
7009     } else
7010         stosb(dst, c, n);
7011     return dst;
7012 }
7013
7014 int
7015 memcmp(const void *v1, const void *v2, uint n)
7016 {
7017     const uchar *s1, *s2;
7018
7019     s1 = v1;
7020     s2 = v2;
7021     while(n-- > 0){
7022         if(*s1 != *s2)
7023             return *s1 - *s2;
7024         s1++, s2++;
7025     }
7026
7027     return 0;
7028 }
7029
7030 void*
7031 memmove(void *dst, const void *src, uint n)
7032 {
7033     const char *s;
7034     char *d;
7035
7036     s = src;
7037     d = dst;
7038     if(s < d && s + n > d){
7039         s += n;
7040         d += n;
7041         while(n-- > 0)
7042             *--d = *--s;
7043     } else
7044         while(n-- > 0)
7045             *d++ = *s++;
7046
7047     return dst;
7048 }
7049

```

```

7050 // memcpy exists to placate GCC. Use memmove.
7051 void*
7052 memcpy(void *dst, const void *src, uint n)
7053 {
7054     return memmove(dst, src, n);
7055 }
7056
7057 int
7058 strncmp(const char *p, const char *q, uint n)
7059 {
7060     while(n > 0 && *p && *p == *q)
7061         n--, p++, q++;
7062     if(n == 0)
7063         return 0;
7064     return (uchar)*p - (uchar)*q;
7065 }
7066
7067 char*
7068 strncpy(char *s, const char *t, int n)
7069 {
7070     char *os;
7071
7072     os = s;
7073     while(n-- > 0 && (*s++ = *t++) != 0)
7074         ;
7075     while(n-- > 0)
7076         *s++ = 0;
7077     return os;
7078 }
7079
7080 // Like strncpy but guaranteed to NUL-terminate.
7081 char*
7082 safestrcpy(char *s, const char *t, int n)
7083 {
7084     char *os;
7085
7086     os = s;
7087     if(n <= 0)
7088         return os;
7089     while(--n > 0 && (*s++ = *t++) != 0)
7090         ;
7091     *s = 0;
7092     return os;
7093 }
7094
7095
7096
7097
7098
7099

```

```

7100 int
7101 strlen(const char *s)
7102 {
7103     int n;
7104
7105     for(n = 0; s[n]; n++)
7106         ;
7107     return n;
7108 }
7109
7110
7111
7112
7113
7114
7115
7116
7117
7118
7119
7120
7121
7122
7123
7124
7125
7126
7127
7128
7129
7130
7131
7132
7133
7134
7135
7136
7137
7138
7139
7140
7141
7142
7143
7144
7145
7146
7147
7148
7149

```

```

7150 // See MultiProcessor Specification Version 1.[14]
7151
7152 struct mp {                // floating pointer
7153     uchar signature[4];    // "_MP_"
7154     void *physaddr;        // phys addr of MP config table
7155     uchar length;          // 1
7156     uchar specrev;         // [14]
7157     uchar checksum;        // all bytes must add up to 0
7158     uchar type;            // MP system config type
7159     uchar imcrp;
7160     uchar reserved[3];
7161 };
7162
7163 struct mpconf {            // configuration table header
7164     uchar signature[4];    // "PCMP"
7165     ushort length;         // total table length
7166     uchar version;         // [14]
7167     uchar checksum;        // all bytes must add up to 0
7168     uchar product[20];     // product id
7169     uint *oemtable;        // OEM table pointer
7170     ushort oemlength;      // OEM table length
7171     ushort entry;          // entry count
7172     uint *lapicaddr;       // address of local APIC
7173     ushort xlength;        // extended table length
7174     uchar xchecksum;       // extended table checksum
7175     uchar reserved;
7176 };
7177
7178 struct mpproc {            // processor table entry
7179     uchar type;            // entry type (0)
7180     uchar apicid;          // local APIC id
7181     uchar version;         // local APIC verison
7182     uchar flags;           // CPU flags
7183     #define MPBOOT 0x02    // This proc is the bootstrap processor.
7184     uchar signature[4];    // CPU signature
7185     uint feature;          // feature flags from CPUID instruction
7186     uchar reserved[8];
7187 };
7188
7189 struct mpioapic {          // I/O APIC table entry
7190     uchar type;            // entry type (2)
7191     uchar apicno;          // I/O APIC id
7192     uchar version;         // I/O APIC version
7193     uchar flags;           // I/O APIC flags
7194     uint *addr;            // I/O APIC address
7195 };
7196
7197
7198
7199

```

```
7200 // Table entry types
7201 #define MPPROC    0x00 // One per processor
7202 #define MPBUS      0x01 // One per bus
7203 #define MPIOAPIC   0x02 // One per I/O APIC
7204 #define MPIOINTR   0x03 // One per bus interrupt source
7205 #define MPLINTR    0x04 // One per system interrupt source
7206
7207
7208
7209
7210
7211
7212
7213
7214
7215
7216
7217
7218
7219
7220
7221
7222
7223
7224
7225
7226
7227
7228
7229
7230
7231
7232
7233
7234
7235
7236
7237
7238
7239
7240
7241
7242
7243
7244
7245
7246
7247
7248
7249
```

```
7250 // Blank page.
7251
7252
7253
7254
7255
7256
7257
7258
7259
7260
7261
7262
7263
7264
7265
7266
7267
7268
7269
7270
7271
7272
7273
7274
7275
7276
7277
7278
7279
7280
7281
7282
7283
7284
7285
7286
7287
7288
7289
7290
7291
7292
7293
7294
7295
7296
7297
7298
7299
```



```

7300 // Multiprocessor support
7301 // Search memory for MP description structures.
7302 // http://developer.intel.com/design/pentium/datashts/24201606.pdf
7303
7304 #include "types.h"
7305 #include "defs.h"
7306 #include "param.h"
7307 #include "memlayout.h"
7308 #include "mp.h"
7309 #include "x86.h"
7310 #include "mmu.h"
7311 #include "proc.h"
7312
7313 struct cpu cpus[NCPU];
7314 static struct cpu *bcpu;
7315 int ismp;
7316 int ncpu;
7317 uchar ioapicid;
7318
7319 int
7320 mpbcpu(void)
7321 {
7322     return bcpu-cpus;
7323 }
7324
7325 static uchar
7326 sum(uchar *addr, int len)
7327 {
7328     int i, sum;
7329
7330     sum = 0;
7331     for(i=0; i<len; i++)
7332         sum += addr[i];
7333     return sum;
7334 }
7335
7336 // Look for an MP structure in the len bytes at addr.
7337 static struct mp*
7338 mpsearch1(uint a, int len)
7339 {
7340     uchar *e, *p, *addr;
7341
7342     addr = p2v(a);
7343     e = addr+len;
7344     for(p = addr; p < e; p += sizeof(struct mp))
7345         if(memcmp(p, "_MP_", 4) == 0 && sum(p, sizeof(struct mp)) == 0)
7346             return (struct mp*)p;
7347     return 0;
7348 }
7349

```

```

7350 // Search for the MP Floating Pointer Structure, which according to the
7351 // spec is in one of the following three locations:
7352 // 1) in the first KB of the EBDA;
7353 // 2) in the last KB of system base memory;
7354 // 3) in the BIOS ROM between 0xE0000 and 0xFFFFF.
7355 static struct mp*
7356 mpsearch(void)
7357 {
7358     uchar *bda;
7359     uint p;
7360     struct mp *mp;
7361
7362     bda = (uchar *) P2V(0x400);
7363     if((p = ((bda[0x0F]<<8) | bda[0x0E]) << 4)){
7364         if((mp = mpsearch1(p, 1024)))
7365             return mp;
7366     } else {
7367         p = ((bda[0x14]<<8) | bda[0x13])*1024;
7368         if((mp = mpsearch1(p-1024, 1024)))
7369             return mp;
7370     }
7371     return mpsearch1(0xF0000, 0x10000);
7372 }
7373
7374 // Search for an MP configuration table. For now,
7375 // don't accept the default configurations (physaddr == 0).
7376 // Check for correct signature, calculate the checksum and,
7377 // if correct, check the version.
7378 // To do: check extended table checksum.
7379 static struct mpconf*
7380 mpconfig(struct mp **pmp)
7381 {
7382     struct mpconf *conf;
7383     struct mp *mp;
7384
7385     if((mp = mpsearch()) == 0 || mp->physaddr == 0)
7386         return 0;
7387     conf = (struct mpconf*) p2v((uint) mp->physaddr);
7388     if(memcmp(conf, "PCMP", 4) != 0)
7389         return 0;
7390     if(conf->version != 1 && conf->version != 4)
7391         return 0;
7392     if(sum((uchar*)conf, conf->length) != 0)
7393         return 0;
7394     *pmp = mp;
7395     return conf;
7396 }
7397
7398
7399

```

```

7400 void
7401 mpinit(void)
7402 {
7403     uchar *p, *e;
7404     struct mp *mp;
7405     struct mpconf *conf;
7406     struct mpproc *proc;
7407     struct mpioapic *ioapic;
7408
7409     bcpu = &cpus[0];
7410     if((conf = mpconfig(&mp)) == 0)
7411         return;
7412     ismp = 1;
7413     lapic = (uint*)conf->lapicaddr;
7414     for(p=(uchar*)(conf+1), e=(uchar*)conf+conf->length; p<e; ){
7415         switch(*p){
7416             case MPPROC:
7417                 proc = (struct mpproc*)p;
7418                 if(ncpu != proc->apicid){
7419                     cprintf("mpinit: ncpu=%d apicid=%d\n", ncpu, proc->apicid);
7420                     ismp = 0;
7421                 }
7422                 if(proc->flags & MPBOOT)
7423                     bcpu = &cpus[ncpu];
7424                 cpus[ncpu].id = ncpu;
7425                 ncpu++;
7426                 p += sizeof(struct mpproc);
7427                 continue;
7428             case MPIOAPIC:
7429                 ioapic = (struct mpioapic*)p;
7430                 ioapicid = ioapic->apicno;
7431                 p += sizeof(struct mpioapic);
7432                 continue;
7433             case MPBUS:
7434             case MPIOINTR:
7435             case MPLINTR:
7436                 p += 8;
7437                 continue;
7438             default:
7439                 cprintf("mpinit: unknown config type %x\n", *p);
7440                 ismp = 0;
7441         }
7442     }
7443     if(!ismp){
7444         // Didn't like what we found; fall back to no MP.
7445         ncpu = 1;
7446         lapic = 0;
7447         ioapicid = 0;
7448         return;
7449     }

```

```

7450     if(mp->imcrp){
7451         // Bochs doesn't support IMCR, so this doesn't run on Bochs.
7452         // But it would on real hardware.
7453         outb(0x22, 0x70); // Select IMCR
7454         outb(0x23, inb(0x23) | 1); // Mask external interrupts.
7455     }
7456 }
7457
7458
7459
7460
7461
7462
7463
7464
7465
7466
7467
7468
7469
7470
7471
7472
7473
7474
7475
7476
7477
7478
7479
7480
7481
7482
7483
7484
7485
7486
7487
7488
7489
7490
7491
7492
7493
7494
7495
7496
7497
7498
7499

```

```

7500 // The local APIC manages internal (non-I/O) interrupts.
7501 // See Chapter 8 & Appendix C of Intel processor manual volume 3.
7502
7503 #include "types.h"
7504 #include "defs.h"
7505 #include "date.h"
7506 #include "memlayout.h"
7507 #include "traps.h"
7508 #include "mmu.h"
7509 #include "x86.h"
7510
7511 // Local APIC registers, divided by 4 for use as uint[] indices.
7512 #define ID      (0x0020/4) // ID
7513 #define VER     (0x0030/4) // Version
7514 #define TPR     (0x0080/4) // Task Priority
7515 #define EOI     (0x00B0/4) // EOI
7516 #define SVR     (0x00F0/4) // Spurious Interrupt Vector
7517 #define ENABLE  0x00000100 // Unit Enable
7518 #define ESR     (0x0280/4) // Error Status
7519 #define ICRLO   (0x0300/4) // Interrupt Command
7520 #define INIT    0x00000500 // INIT/RESET
7521 #define STARTUP 0x00000600 // Startup IPI
7522 #define DELIVS  0x00001000 // Delivery status
7523 #define ASSERT  0x00004000 // Assert interrupt (vs deassert)
7524 #define DEASSERT 0x00000000
7525 #define LEVEL   0x00008000 // Level triggered
7526 #define BCAST   0x00080000 // Send to all APICs, including self.
7527 #define BUSY    0x00001000
7528 #define FIXED    0x00000000
7529 #define ICRHI   (0x0310/4) // Interrupt Command [63:32]
7530 #define TIMER   (0x0320/4) // Local Vector Table 0 (TIMER)
7531 #define X1      0x0000000B // divide counts by 1
7532 #define PERIODIC 0x00020000 // Periodic
7533 #define PCINT   (0x0340/4) // Performance Counter LVT
7534 #define LINT0   (0x0350/4) // Local Vector Table 1 (LINT0)
7535 #define LINT1   (0x0360/4) // Local Vector Table 2 (LINT1)
7536 #define ERROR   (0x0370/4) // Local Vector Table 3 (ERROR)
7537 #define MASKED  0x00010000 // Interrupt masked
7538 #define TICC    (0x0380/4) // Timer Initial Count
7539 #define TCCR    (0x0390/4) // Timer Current Count
7540 #define TDCR    (0x03E0/4) // Timer Divide Configuration
7541
7542 volatile uint *lapic; // Initialized in mp.c
7543
7544 static void
7545 lapicw(int index, int value)
7546 {
7547     lapic[index] = value;
7548     lapic[ID]; // wait for write to finish, by reading
7549 }

```

```

7550
7551
7552
7553
7554
7555
7556
7557
7558
7559
7560
7561
7562
7563
7564
7565
7566
7567
7568
7569
7570
7571
7572
7573
7574
7575
7576
7577
7578
7579
7580
7581
7582
7583
7584
7585
7586
7587
7588
7589
7590
7591
7592
7593
7594
7595
7596
7597
7598
7599

```

```

7600 void
7601 lapicinit(void)
7602 {
7603     if(!lapic)
7604         return;
7605
7606     // Enable local APIC; set spurious interrupt vector.
7607     lapicw(SVR, ENABLE | (T_IRQ0 + IRQ_SPURIOUS));
7608
7609     // The timer repeatedly counts down at bus frequency
7610     // from lapic[TICR] and then issues an interrupt.
7611     // If xv6 cared more about precise timekeeping,
7612     // TICR would be calibrated using an external time source.
7613     lapicw(TDCR, X1);
7614     lapicw(TIMER, PERIODIC | (T_IRQ0 + IRQ_TIMER));
7615     lapicw(TICR, 10000000);
7616
7617     // Disable logical interrupt lines.
7618     lapicw(LINT0, MASKED);
7619     lapicw(LINT1, MASKED);
7620
7621     // Disable performance counter overflow interrupts
7622     // on machines that provide that interrupt entry.
7623     if(((lapic[VER]>>16) & 0xFF) >= 4)
7624         lapicw(PCINT, MASKED);
7625
7626     // Map error interrupt to IRQ_ERROR.
7627     lapicw(ERROR, T_IRQ0 + IRQ_ERROR);
7628
7629     // Clear error status register (requires back-to-back writes).
7630     lapicw(ESR, 0);
7631     lapicw(ESR, 0);
7632
7633     // Ack any outstanding interrupts.
7634     lapicw(EOI, 0);
7635
7636     // Send an Init Level De-Assert to synchronise arbitration ID's.
7637     lapicw(ICRHI, 0);
7638     lapicw(ICRLO, BCAST | INIT | LEVEL);
7639     while(lapic[ICRLO] & DELIVS)
7640         ;
7641
7642     // Enable interrupts on the APIC (but not on the processor).
7643     lapicw(TPR, 0);
7644 }
7645
7646
7647
7648
7649

```

```

7650 int
7651 cpunum(void)
7652 {
7653     // Cannot call cpu when interrupts are enabled:
7654     // result not guaranteed to last long enough to be used!
7655     // Would prefer to panic but even printing is chancy here:
7656     // almost everything, including cprintf and panic, calls cpu,
7657     // often indirectly through acquire and release.
7658     if(readeflags() & FL_IF){
7659         static int n;
7660         if(n++ == 0)
7661             cprintf("cpu called from %x with interrupts enabled\n",
7662                     __builtin_return_address(0));
7663     }
7664
7665     if(lapic)
7666         return lapic[ID]>>24;
7667     return 0;
7668 }
7669
7670 // Acknowledge interrupt.
7671 void
7672 lapiceoi(void)
7673 {
7674     if(lapic)
7675         lapicw(EOI, 0);
7676 }
7677
7678 // Spin for a given number of microseconds.
7679 // On real hardware would want to tune this dynamically.
7680 void
7681 microdelay(int us)
7682 {
7683 }
7684
7685 #define CMOS_PORT    0x70
7686 #define CMOS_RETURN  0x71
7687
7688 // Start additional processor running entry code at addr.
7689 // See Appendix B of MultiProcessor Specification.
7690 void
7691 lapicstartap(uchar apicid, uint addr)
7692 {
7693     int i;
7694     ushort *wrv;
7695
7696     // "The BSP must initialize CMOS shutdown code to 0AH
7697     // and the warm reset vector (DWORD based at 40:67) to point at
7698     // the AP startup code prior to the [universal startup algorithm]."
7699     outb(CMOS_PORT, 0xF); // offset 0xF is shutdown code

```

```

7700 outb(CMOS_PORT+1, 0x0A);
7701 wrv = (ushort*)P2V((0x40<<4 | 0x67)); // Warm reset vector
7702 wrv[0] = 0;
7703 wrv[1] = addr >> 4;
7704
7705 // "Universal startup algorithm."
7706 // Send INIT (level-triggered) interrupt to reset other CPU.
7707 lapicw(ICRHI, apicid<<24);
7708 lapicw(ICRLO, INIT | LEVEL | ASSERT);
7709 microdelay(200);
7710 lapicw(ICRLO, INIT | LEVEL);
7711 microdelay(100); // should be 10ms, but too slow in Bochs!
7712
7713 // Send startup IPI (twice!) to enter code.
7714 // Regular hardware is supposed to only accept a STARTUP
7715 // when it is in the halted state due to an INIT. So the second
7716 // should be ignored, but it is part of the official Intel algorithm.
7717 // Bochs complains about the second one. Too bad for Bochs.
7718 for(i = 0; i < 2; i++){
7719     lapicw(ICRHI, apicid<<24);
7720     lapicw(ICRLO, STARTUP | (addr>>12));
7721     microdelay(200);
7722 }
7723 }
7724
7725 #define CMOS_STATA 0x0a
7726 #define CMOS_STATB 0x0b
7727 #define CMOS_UIP   (1 << 7) // RTC update in progress
7728
7729 #define SECS 0x00
7730 #define MINS 0x02
7731 #define HOURS 0x04
7732 #define DAY 0x07
7733 #define MONTH 0x08
7734 #define YEAR 0x09
7735
7736 static uint cmos_read(uint reg)
7737 {
7738     outb(CMOS_PORT, reg);
7739     microdelay(200);
7740
7741     return inb(CMOS_RETURN);
7742 }
7743
7744
7745
7746
7747
7748
7749

```

```

7750 static void fill_rtcddate(struct rtcdate *r)
7751 {
7752     r->second = cmos_read(SECS);
7753     r->minute = cmos_read(MINS);
7754     r->hour   = cmos_read(HOURS);
7755     r->day    = cmos_read(DAY);
7756     r->month  = cmos_read(MONTH);
7757     r->year   = cmos_read(YEAR);
7758 }
7759
7760 // qemu seems to use 24-hour GWT and the values are BCD encoded
7761 void cmostime(struct rtcdate *r)
7762 {
7763     struct rtcdate t1, t2;
7764     int sb, bcd;
7765
7766     sb = cmos_read(CMOS_STATB);
7767
7768     bcd = (sb & (1 << 2)) == 0;
7769
7770     // make sure CMOS doesn't modify time while we read it
7771     for (;;) {
7772         fill_rtcddate(&t1);
7773         if (cmos_read(CMOS_STATA) & CMOS_UIP)
7774             continue;
7775         fill_rtcddate(&t2);
7776         if (memcmp(&t1, &t2, sizeof(t1)) == 0)
7777             break;
7778     }
7779
7780     // convert
7781     if (bcd) {
7782 #define CONV(x) ((t1.x >> 4) * 10) + (t1.x & 0xf)
7783         CONV(second);
7784         CONV(minute);
7785         CONV(hour );
7786         CONV(day );
7787         CONV(month );
7788         CONV(year );
7789 #undef CONV
7790     }
7791
7792     *r = t1;
7793     r->year += 2000;
7794 }
7795
7796
7797
7798
7799

```

```

7800 // The I/O APIC manages hardware interrupts for an SMP system.
7801 // http://www.intel.com/design/chipsets/datashts/29056601.pdf
7802 // See also picirq.c.
7803
7804 #include "types.h"
7805 #include "defs.h"
7806 #include "traps.h"
7807
7808 #define IOAPIC 0xFEC00000 // Default physical address of IO APIC
7809
7810 #define REG_ID 0x00 // Register index: ID
7811 #define REG_VER 0x01 // Register index: version
7812 #define REG_TABLE 0x10 // Redirection table base
7813
7814 // The redirection table starts at REG_TABLE and uses
7815 // two registers to configure each interrupt.
7816 // The first (low) register in a pair contains configuration bits.
7817 // The second (high) register contains a bitmask telling which
7818 // CPUs can serve that interrupt.
7819 #define INT_DISABLED 0x00010000 // Interrupt disabled
7820 #define INT_LEVEL 0x00008000 // Level-triggered (vs edge-)
7821 #define INT_ACTIVELOW 0x00002000 // Active low (vs high)
7822 #define INT_LOGICAL 0x00000800 // Destination is CPU id (vs APIC ID)
7823
7824 volatile struct ioapic *ioapic;
7825
7826 // IO APIC MMIO structure: write reg, then read or write data.
7827 struct ioapic {
7828     uint reg;
7829     uint pad[3];
7830     uint data;
7831 };
7832
7833 static uint
7834 ioapicread(int reg)
7835 {
7836     ioapic->reg = reg;
7837     return ioapic->data;
7838 }
7839
7840 static void
7841 ioapicwrite(int reg, uint data)
7842 {
7843     ioapic->reg = reg;
7844     ioapic->data = data;
7845 }
7846
7847
7848
7849

```

```

7850 void
7851 ioapicinit(void)
7852 {
7853     int i, id, maxintr;
7854
7855     if(!ismp)
7856         return;
7857
7858     ioapic = (volatile struct ioapic*)IOAPIC;
7859     maxintr = (ioapicread(REG_VER) >> 16) & 0xFF;
7860     id = ioapicread(REG_ID) >> 24;
7861     if(id != ioapicid)
7862         cprintf("ioapicinit: id isn't equal to ioapicid; not a MP\n");
7863
7864     // Mark all interrupts edge-triggered, active high, disabled,
7865     // and not routed to any CPUs.
7866     for(i = 0; i <= maxintr; i++){
7867         ioapicwrite(REG_TABLE+2*i, INT_DISABLED | (T_IRQ0 + i));
7868         ioapicwrite(REG_TABLE+2*i+1, 0);
7869     }
7870 }
7871
7872 void
7873 ioapicenable(int irq, int cpunum)
7874 {
7875     if(!ismp)
7876         return;
7877
7878     // Mark interrupt edge-triggered, active high,
7879     // enabled, and routed to the given cpunum,
7880     // which happens to be that cpu's APIC ID.
7881     ioapicwrite(REG_TABLE+2*irq, T_IRQ0 + irq);
7882     ioapicwrite(REG_TABLE+2*irq+1, cpunum << 24);
7883 }
7884
7885
7886
7887
7888
7889
7890
7891
7892
7893
7894
7895
7896
7897
7898
7899

```

```

7900 // Intel 8259A programmable interrupt controllers.
7901
7902 #include "types.h"
7903 #include "x86.h"
7904 #include "traps.h"
7905
7906 // I/O Addresses of the two programmable interrupt controllers
7907 #define IO_PIC1      0x20    // Master (IRQs 0-7)
7908 #define IO_PIC2      0xA0    // Slave (IRQs 8-15)
7909
7910 #define IRQ_SLAVE     2      // IRQ at which slave connects to master
7911
7912 // Current IRQ mask.
7913 // Initial IRQ mask has interrupt 2 enabled (for slave 8259A).
7914 static ushort irqmask = 0xFFFF & ~(1<<IRQ_SLAVE);
7915
7916 static void
7917 picsetmask(ushort mask)
7918 {
7919     irqmask = mask;
7920     outb(IO_PIC1+1, mask);
7921     outb(IO_PIC2+1, mask >> 8);
7922 }
7923
7924 void
7925 picenable(int irq)
7926 {
7927     picsetmask(irqmask & ~(1<<irq));
7928 }
7929
7930 // Initialize the 8259A interrupt controllers.
7931 void
7932 picinit(void)
7933 {
7934     // mask all interrupts
7935     outb(IO_PIC1+1, 0xFF);
7936     outb(IO_PIC2+1, 0xFF);
7937
7938     // Set up master (8259A-1)
7939
7940     // ICW1: 0001g0hi
7941     //   g: 0 = edge triggering, 1 = level triggering
7942     //   h: 0 = cascaded PICs, 1 = master only
7943     //   i: 0 = no ICW4, 1 = ICW4 required
7944     outb(IO_PIC1, 0x11);
7945
7946     // ICW2: Vector offset
7947     outb(IO_PIC1+1, T_IRQ0);
7948
7949

```

```

7950 // ICW3: (master PIC) bit mask of IR lines connected to slaves
7951 //        (slave PIC) 3-bit # of slave's connection to master
7952 outb(IO_PIC1+1, 1<<IRQ_SLAVE);
7953
7954 // ICW4: 000nbmap
7955 //   n: 1 = special fully nested mode
7956 //   b: 1 = buffered mode
7957 //   m: 0 = slave PIC, 1 = master PIC
7958 //        (ignored when b is 0, as the master/slave role
7959 //        can be hardwired).
7960 //   a: 1 = Automatic EOI mode
7961 //   p: 0 = MCS-80/85 mode, 1 = intel x86 mode
7962 outb(IO_PIC1+1, 0x3);
7963
7964 // Set up slave (8259A-2)
7965 outb(IO_PIC2, 0x11); // ICW1
7966 outb(IO_PIC2+1, T_IRQ0 + 8); // ICW2
7967 outb(IO_PIC2+1, IRQ_SLAVE); // ICW3
7968 // NB Automatic EOI mode doesn't tend to work on the slave.
7969 // Linux source code says it's "to be investigated".
7970 outb(IO_PIC2+1, 0x3); // ICW4
7971
7972 // OCW3: 0ef01prs
7973 //   ef: 0x = NOP, 10 = clear specific mask, 11 = set specific mask
7974 //   p: 0 = no polling, 1 = polling mode
7975 //   rs: 0x = NOP, 10 = read IRR, 11 = read ISR
7976 outb(IO_PIC1, 0x68); // clear specific mask
7977 outb(IO_PIC1, 0x0a); // read IRR by default
7978
7979 outb(IO_PIC2, 0x68); // OCW3
7980 outb(IO_PIC2, 0x0a); // OCW3
7981
7982 if(irqmask != 0xFFFF)
7983     picsetmask(irqmask);
7984 }
7985
7986
7987
7988
7989
7990
7991
7992
7993
7994
7995
7996
7997
7998
7999

```

```

8000 // PC keyboard interface constants
8001
8002 #define KBSTATP      0x64    // kbd controller status port(I)
8003 #define KBS_DIB      0x01    // kbd data in buffer
8004 #define KBDATAP      0x60    // kbd data port(I)
8005
8006 #define NO            0
8007
8008 #define SHIFT         (1<<0)
8009 #define CTL           (1<<1)
8010 #define ALT           (1<<2)
8011
8012 #define CAPSLOCK      (1<<3)
8013 #define NUMLOCK       (1<<4)
8014 #define SCROLLLOCK   (1<<5)
8015
8016 #define E0ESC         (1<<6)
8017
8018 // Special keycodes
8019 #define KEY_HOME      0xE0
8020 #define KEY_END       0xE1
8021 #define KEY_UP        0xE2
8022 #define KEY_DN        0xE3
8023 #define KEY_LF        0xE4
8024 #define KEY_RT        0xE5
8025 #define KEY_PGUP      0xE6
8026 #define KEY_PGDN      0xE7
8027 #define KEY_INS       0xE8
8028 #define KEY_DEL       0xE9
8029
8030 // C('A') == Control-A
8031 #define C(x) (x - '@')
8032
8033 static uchar shiftcode[256] =
8034 {
8035     [0x1D] CTL,
8036     [0x2A] SHIFT,
8037     [0x36] SHIFT,
8038     [0x38] ALT,
8039     [0x9D] CTL,
8040     [0xB8] ALT
8041 };
8042
8043 static uchar togglecode[256] =
8044 {
8045     [0x3A] CAPSLOCK,
8046     [0x45] NUMLOCK,
8047     [0x46] SCROLLLOCK
8048 };
8049

```

```

8050 static uchar normalmap[256] =
8051 {
8052     NO,    0x1B, '1', '2', '3', '4', '5', '6', // 0x00
8053     '7', '8', '9', '0', '-', '=', '\b', '\t',
8054     'q', 'w', 'e', 'r', 't', 'y', 'u', 'i', // 0x10
8055     'o', 'p', '[', ']', '\n', NO, 'a', 's',
8056     'd', 'f', 'g', 'h', 'j', 'k', 'l', ';', // 0x20
8057     '\'', ',', NO, '\\', 'z', 'x', 'c', 'v',
8058     'b', 'n', 'm', '.', '/', NO, '*', // 0x30
8059     NO, ' ', NO, NO, NO, NO, NO, NO,
8060     NO, NO, NO, NO, NO, NO, NO, '7', // 0x40
8061     '8', '9', '-', '4', '5', '6', '+', '1',
8062     '2', '3', '0', '.', NO, NO, NO, NO, // 0x50
8063     [0x9C] '\n', // KP_Enter
8064     [0xB5] '/', // KP_Div
8065     [0xC8] KEY_UP, [0xD0] KEY_DN,
8066     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
8067     [0xCB] KEY_LF, [0xCD] KEY_RT,
8068     [0x97] KEY_HOME, [0xCF] KEY_END,
8069     [0xD2] KEY_INS, [0xD3] KEY_DEL
8070 };
8071
8072 static uchar shiftmap[256] =
8073 {
8074     NO,    033, '!', '@', '#', '$', '%', '^', // 0x00
8075     '&', '*', '(', ')', '_', '+', '\b', '\t',
8076     'Q', 'W', 'E', 'R', 'T', 'Y', 'U', 'I', // 0x10
8077     'O', 'P', '{', '}', '\n', NO, 'A', 'S',
8078     'D', 'F', 'G', 'H', 'J', 'K', 'L', ';', // 0x20
8079     '"', '~', NO, '|', 'Z', 'X', 'C', 'V',
8080     'B', 'N', 'M', '<', '>', '?', NO, '*', // 0x30
8081     NO, ' ', NO, NO, NO, NO, NO, NO,
8082     NO, NO, NO, NO, NO, NO, NO, '7', // 0x40
8083     '8', '9', '-', '4', '5', '6', '+', '1',
8084     '2', '3', '0', '.', NO, NO, NO, NO, // 0x50
8085     [0x9C] '\n', // KP_Enter
8086     [0xB5] '/', // KP_Div
8087     [0xC8] KEY_UP, [0xD0] KEY_DN,
8088     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
8089     [0xCB] KEY_LF, [0xCD] KEY_RT,
8090     [0x97] KEY_HOME, [0xCF] KEY_END,
8091     [0xD2] KEY_INS, [0xD3] KEY_DEL
8092 };
8093
8094
8095
8096
8097
8098
8099

```



```

8100 static uchar ctlmap[256] =
8101 {
8102     NO,      NO,      NO,      NO,      NO,      NO,      NO,      NO,
8103     NO,      NO,      NO,      NO,      NO,      NO,      NO,      NO,
8104     C('Q'),  C('W'),  C('E'),  C('R'),  C('T'),  C('Y'),  C('U'),  C('I'),
8105     C('O'),  C('P'),  NO,      NO,      '\r',  NO,      C('A'),  C('S'),
8106     C('D'),  C('F'),  C('G'),  C('H'),  C('J'),  C('K'),  C('L'),  NO,
8107     NO,      NO,      NO,      C('\n'), C('Z'),  C('X'),  C('C'),  C('V'),
8108     C('B'),  C('N'),  C('M'),  NO,      NO,      C('/'), NO,      NO,
8109     [0x9C] '\r',      // KP_Enter
8110     [0xB5] C('/'),    // KP_Div
8111     [0xC8] KEY_UP,    [0xD0] KEY_DN,
8112     [0xC9] KEY_PGUP,  [0xD1] KEY_PGDN,
8113     [0xCB] KEY_LF,    [0xCD] KEY_RT,
8114     [0x97] KEY_HOME,  [0xCF] KEY_END,
8115     [0xD2] KEY_INS,   [0xD3] KEY_DEL
8116 };
8117
8118
8119
8120
8121
8122
8123
8124
8125
8126
8127
8128
8129
8130
8131
8132
8133
8134
8135
8136
8137
8138
8139
8140
8141
8142
8143
8144
8145
8146
8147
8148
8149

```

```

8150 #include "types.h"
8151 #include "x86.h"
8152 #include "defs.h"
8153 #include "kbd.h"
8154
8155 int
8156 kbdgetc(void)
8157 {
8158     static uint shift;
8159     static uchar *charcode[4] = {
8160         normalmap, shiftmap, ctlmap, ctlmap
8161     };
8162     uint st, data, c;
8163
8164     st = inb(KBSTATP);
8165     if((st & KBS_DIB) == 0)
8166         return -1;
8167     data = inb(KBDATAP);
8168
8169     if(data == 0xE0){
8170         shift |= E0ESC;
8171         return 0;
8172     } else if(data & 0x80){
8173         // Key released
8174         data = (shift & E0ESC ? data : data & 0x7F);
8175         shift &= ~(shiftcode[data] | E0ESC);
8176         return 0;
8177     } else if(shift & E0ESC){
8178         // Last character was an E0 escape; or with 0x80
8179         data |= 0x80;
8180         shift &= ~E0ESC;
8181     }
8182
8183     shift |= shiftcode[data];
8184     shift ^= togglecode[data];
8185     c = charcode[shift & (CTL | SHIFT)][data];
8186     if(shift & CAPSLOCK){
8187         if('a' <= c && c <= 'z')
8188             c += 'A' - 'a';
8189         else if('A' <= c && c <= 'Z')
8190             c += 'a' - 'A';
8191     }
8192     return c;
8193 }
8194
8195 void
8196 kbdintr(void)
8197 {
8198     consoleintr(kbdgetc);
8199 }

```

```

8200 // Console input and output.
8201 // Input is from the keyboard or serial port.
8202 // Output is written to the screen and serial port.
8203
8204 #include "types.h"
8205 #include "defs.h"
8206 #include "param.h"
8207 #include "traps.h"
8208 #include "spinlock.h"
8209 #include "fs.h"
8210 #include "file.h"
8211 #include "memlayout.h"
8212 #include "mmu.h"
8213 #include "proc.h"
8214 #include "x86.h"
8215
8216 static void consputc(int);
8217
8218 static int panicked = 0;
8219
8220 static struct {
8221   struct spinlock lock;
8222   int locking;
8223 } cons;
8224
8225 static void
8226 printint(int xx, int base, int sign)
8227 {
8228   static char digits[] = "0123456789abcdef";
8229   char buf[16];
8230   int i;
8231   uint x;
8232
8233   if(sign && (sign = xx < 0))
8234     x = -xx;
8235   else
8236     x = xx;
8237
8238   i = 0;
8239   do{
8240     buf[i++] = digits[x % base];
8241   }while((x /= base) != 0);
8242
8243   if(sign)
8244     buf[i++] = '-';
8245
8246   while(--i >= 0)
8247     consputc(buf[i]);
8248 }
8249

```

```

8250 // Print to the console. only understands %d, %x, %p, %s.
8251 void
8252 cprintf(char *fmt, ...)
8253 {
8254   int i, c, locking;
8255   uint *argp;
8256   char *s;
8257
8258   locking = cons.locking;
8259   if(locking)
8260     acquire(&cons.lock);
8261
8262   if (fmt == 0)
8263     panic("null fmt");
8264
8265   argp = (uint*)(void*)&fmt + 1;
8266   for(i = 0; (c = fmt[i] & 0xff) != 0; i++){
8267     if(c != '%'){
8268       consputc(c);
8269       continue;
8270     }
8271     c = fmt[++i] & 0xff;
8272     if(c == 0)
8273       break;
8274     switch(c){
8275     case 'd':
8276       printint(*argp++, 10, 1);
8277       break;
8278     case 'x':
8279     case 'p':
8280       printint(*argp++, 16, 0);
8281       break;
8282     case 's':
8283       if((s = (char*)*argp++) == 0)
8284         s = "(null)";
8285       for(; *s; s++)
8286         consputc(*s);
8287       break;
8288     case '%':
8289       consputc('%');
8290       break;
8291     default:
8292       // Print unknown % sequence to draw attention.
8293       consputc('%');
8294       consputc(c);
8295       break;
8296     }
8297   }
8298 }
8299

```

```

8300  if(locking)
8301      release(&cons.lock);
8302  }
8303
8304  void
8305  panic(char *s)
8306  {
8307      int i;
8308      uint pcs[10];
8309
8310      cli();
8311      cons.locking = 0;
8312      cprintf("cpu%d: panic: ", cpu->id);
8313      cprintf(s);
8314      cprintf("\n");
8315      getcallerpcs(&s, pcs);
8316      for(i=0; i<10; i++)
8317          cprintf(" %p", pcs[i]);
8318      panicked = 1; // freeze other CPU
8319      for(;;)
8320          ;
8321  }
8322
8323
8324
8325
8326
8327
8328
8329
8330
8331
8332
8333
8334
8335
8336
8337
8338
8339
8340
8341
8342
8343
8344
8345
8346
8347
8348
8349

```

```

8350  #define BACKSPACE 0x100
8351  #define CRTPORT 0x3d4
8352  static ushort *crt = (ushort*)P2V(0xb8000); // CGA memory
8353
8354  static void
8355  cgaputc(int c)
8356  {
8357      int pos;
8358
8359      // Cursor position: col + 80*row.
8360      outb(CRTPORT, 14);
8361      pos = inb(CRTPORT+1) << 8;
8362      outb(CRTPORT, 15);
8363      pos |= inb(CRTPORT+1);
8364
8365      if(c == '\n')
8366          pos += 80 - pos%80;
8367      else if(c == BACKSPACE){
8368          if(pos > 0) --pos;
8369      } else
8370          crt[pos++] = (c&0xff) | 0x0700; // black on white
8371
8372      if(pos < 0 || pos > 25*80)
8373          panic("pos under/overflow");
8374
8375      if((pos/80) >= 24){ // Scroll up.
8376          memmove(crt, crt+80, sizeof(crt[0])*23*80);
8377          pos -= 80;
8378          memset(crt+pos, 0, sizeof(crt[0])*(24*80 - pos));
8379      }
8380
8381      outb(CRTPORT, 14);
8382      outb(CRTPORT+1, pos>>8);
8383      outb(CRTPORT, 15);
8384      outb(CRTPORT+1, pos);
8385      crt[pos] = ' ' | 0x0700;
8386  }
8387
8388
8389
8390
8391
8392
8393
8394
8395
8396
8397
8398
8399

```

```

8400 void
8401 consputc(int c)
8402 {
8403     if(panicked){
8404         cli();
8405         for(;;)
8406             ;
8407     }
8408
8409     if(c == BACKSPACE){
8410         uartputc('\b'); uartputc(' '); uartputc('\b');
8411     } else
8412         uartputc(c);
8413     cgaputc(c);
8414 }
8415
8416 #define INPUT_BUF 128
8417 struct {
8418     char buf[INPUT_BUF];
8419     uint r; // Read index
8420     uint w; // Write index
8421     uint e; // Edit index
8422 } input;
8423
8424 #define C(x) ((x)-'@') // Control-x
8425
8426 void
8427 consoleintr(int (*getc)(void))
8428 {
8429     int c, doprocdump = 0;
8430
8431     acquire(&cons.lock);
8432     while((c = getc()) >= 0){
8433         switch(c){
8434             case C('P'): // Process listing.
8435                 doprocdump = 1; // procdump() locks cons.lock indirectly; invoke later
8436                 break;
8437             case C('U'): // Kill line.
8438                 while(input.e != input.w &&
8439                     input.buf[(input.e-1) % INPUT_BUF] != '\n'){
8440                     input.e--;
8441                     consputc(BACKSPACE);
8442                 }
8443                 break;
8444             case C('H'): case '\x7f': // Backspace
8445                 if(input.e != input.w){
8446                     input.e--;
8447                     consputc(BACKSPACE);
8448                 }
8449                 break;

```

```

8450     default:
8451         if(c != 0 && input.e-input.r < INPUT_BUF){
8452             c = (c == '\r') ? '\n' : c;
8453             input.buf[input.e++ % INPUT_BUF] = c;
8454             consputc(c);
8455             if(c == '\n' || c == C('D') || input.e == input.r+INPUT_BUF){
8456                 input.w = input.e;
8457                 wakeup(&input.r);
8458             }
8459         }
8460         break;
8461     }
8462 }
8463 release(&cons.lock);
8464 if(doprocdump) {
8465     procdump(); // now call procdump() wo. cons.lock held
8466 }
8467 }
8468
8469 int
8470 consoleread(struct inode *ip, char *dst, int n)
8471 {
8472     uint target;
8473     int c;
8474
8475     iunlock(ip);
8476     target = n;
8477     acquire(&cons.lock);
8478     while(n > 0){
8479         while(input.r == input.w){
8480             if(proc->killed){
8481                 release(&cons.lock);
8482                 ilock(ip);
8483                 return -1;
8484             }
8485             sleep(&input.r, &cons.lock);
8486         }
8487         c = input.buf[input.r++ % INPUT_BUF];
8488         if(c == C('D')){ // EOF
8489             if(n < target){
8490                 // Save ^D for next time, to make sure
8491                 // caller gets a 0-byte result.
8492                 input.r--;
8493             }
8494             break;
8495         }
8496         *dst++ = c;
8497         --n;
8498         if(c == '\n')
8499             break;

```

```

8500 }
8501 release(&cons.lock);
8502 ilock(ip);
8503
8504 return target - n;
8505 }
8506
8507 int
8508 consolewrite(struct inode *ip, char *buf, int n)
8509 {
8510     int i;
8511
8512     iunlock(ip);
8513     acquire(&cons.lock);
8514     for(i = 0; i < n; i++)
8515         consputc(buf[i] & 0xff);
8516     release(&cons.lock);
8517     ilock(ip);
8518
8519     return n;
8520 }
8521
8522 void
8523 consoleinit(void)
8524 {
8525     initlock(&cons.lock, "console");
8526
8527     devsw[CONSOLE].write = consolewrite;
8528     devsw[CONSOLE].read = consleread;
8529     cons.locking = 1;
8530
8531     picenable(IRQ_KBD);
8532     ioapicenable(IRQ_KBD, 0);
8533 }
8534
8535
8536
8537
8538
8539
8540
8541
8542
8543
8544
8545
8546
8547
8548
8549

```

```

8550 // Intel 8253/8254/82C54 Programmable Interval Timer (PIT).
8551 // Only used on uniprocessors;
8552 // SMP machines use the local APIC timer.
8553
8554 #include "types.h"
8555 #include "defs.h"
8556 #include "traps.h"
8557 #include "x86.h"
8558
8559 #define IO_TIMER1      0x040          // 8253 Timer #1
8560
8561 // Frequency of all three count-down timers;
8562 // (TIMER_FREQ/freq) is the appropriate count
8563 // to generate a frequency of freq Hz.
8564
8565 #define TIMER_FREQ      1193182
8566 #define TIMER_DIV(x)    ((TIMER_FREQ+(x)/2)/(x))
8567
8568 #define TIMER_MODE      (IO_TIMER1 + 3) // timer mode port
8569 #define TIMER_SEL0      0x00          // select counter 0
8570 #define TIMER_RATEGEN    0x04          // mode 2, rate generator
8571 #define TIMER_16BIT      0x30          // r/w counter 16 bits, LSB first
8572
8573 void
8574 timerinit(void)
8575 {
8576     // Interrupt 100 times/sec.
8577     outb(TIMER_MODE, TIMER_SEL0 | TIMER_RATEGEN | TIMER_16BIT);
8578     outb(IO_TIMER1, TIMER_DIV(100) % 256);
8579     outb(IO_TIMER1, TIMER_DIV(100) / 256);
8580     picenable(IRQ_TIMER);
8581 }
8582
8583
8584
8585
8586
8587
8588
8589
8590
8591
8592
8593
8594
8595
8596
8597
8598
8599

```

```

8600 // Intel 8250 serial port (UART).
8601
8602 #include "types.h"
8603 #include "defs.h"
8604 #include "param.h"
8605 #include "traps.h"
8606 #include "spinlock.h"
8607 #include "fs.h"
8608 #include "file.h"
8609 #include "mmu.h"
8610 #include "proc.h"
8611 #include "x86.h"
8612
8613 #define COM1    0x3f8
8614
8615 static int uart;    // is there a uart?
8616
8617 void
8618 uartinit(void)
8619 {
8620     char *p;
8621
8622     // Turn off the FIFO
8623     outb(COM1+2, 0);
8624
8625     // 9600 baud, 8 data bits, 1 stop bit, parity off.
8626     outb(COM1+3, 0x80);    // Unlock divisor
8627     outb(COM1+0, 115200/9600);
8628     outb(COM1+1, 0);
8629     outb(COM1+3, 0x03);    // Lock divisor, 8 data bits.
8630     outb(COM1+4, 0);
8631     outb(COM1+1, 0x01);    // Enable receive interrupts.
8632
8633     // If status is 0xFF, no serial port.
8634     if(inb(COM1+5) == 0xFF)
8635         return;
8636     uart = 1;
8637
8638     // Acknowledge pre-existing interrupt conditions;
8639     // enable interrupts.
8640     inb(COM1+2);
8641     inb(COM1+0);
8642     picenable(IRQ_COM1);
8643     ioapicenable(IRQ_COM1, 0);
8644
8645     // Announce that we're here.
8646     for(p="xv6...\n"; *p; p++)
8647         uartputc(*p);
8648 }
8649

```

```

8650 void
8651 uartputc(int c)
8652 {
8653     int i;
8654
8655     if(!uart)
8656         return;
8657     for(i = 0; i < 128 && !(inb(COM1+5) & 0x20); i++)
8658         microdelay(10);
8659     outb(COM1+0, c);
8660 }
8661
8662 static int
8663 uartgetc(void)
8664 {
8665     if(!uart)
8666         return -1;
8667     if(!(inb(COM1+5) & 0x01))
8668         return -1;
8669     return inb(COM1+0);
8670 }
8671
8672 void
8673 uartintr(void)
8674 {
8675     consoleintr(uartgetc);
8676 }
8677
8678
8679
8680
8681
8682
8683
8684
8685
8686
8687
8688
8689
8690
8691
8692
8693
8694
8695
8696
8697
8698
8699

```

```

8700 # Initial process execs /init.
8701
8702 #include "syscall.h"
8703 #include "traps.h"
8704
8705
8706 # exec(init, argv)
8707 .globl start
8708 start:
8709     pushl $argv
8710     pushl $init
8711     pushl $0 // where caller pc would be
8712     movl $SYS_exec, %eax
8713     int $T_SYSCALL
8714
8715 # for(;;) exit();
8716 exit:
8717     movl $SYS_exit, %eax
8718     int $T_SYSCALL
8719     jmp exit
8720
8721 # char init[] = "/init\0";
8722 init:
8723     .string "/init\0"
8724
8725 # char *argv[] = { init, 0 };
8726 .p2align 2
8727 argv:
8728     .long init
8729     .long 0
8730
8731
8732
8733
8734
8735
8736
8737
8738
8739
8740
8741
8742
8743
8744
8745
8746
8747
8748
8749

```

```

8750 #include "syscall.h"
8751 #include "traps.h"
8752
8753 #define SYSCALL(name) \
8754     .globl name; \
8755     name: \
8756     movl $SYS_ ## name, %eax; \
8757     int $T_SYSCALL; \
8758     ret
8759
8760 SYSCALL(fork)
8761 SYSCALL(exit)
8762 SYSCALL(wait)
8763 SYSCALL(pipe)
8764 SYSCALL(read)
8765 SYSCALL(write)
8766 SYSCALL(close)
8767 SYSCALL(kill)
8768 SYSCALL(exec)
8769 SYSCALL(open)
8770 SYSCALL(mknod)
8771 SYSCALL(unlink)
8772 SYSCALL(fstat)
8773 SYSCALL(link)
8774 SYSCALL(mkdir)
8775 SYSCALL(chdir)
8776 SYSCALL(dup)
8777 SYSCALL(getpid)
8778 SYSCALL(sbrk)
8779 SYSCALL(sleep)
8780 SYSCALL(uptime)
8781 SYSCALL(halt)
8782 SYSCALL(date)
8783 SYSCALL(getuid)
8784 SYSCALL(getgid)
8785 SYSCALL(getppid)
8786 SYSCALL(setuid)
8787 SYSCALL(setgid)
8788 SYSCALL(getprocs)
8789
8790
8791
8792
8793
8794
8795
8796
8797
8798
8799

```

```

8800 // init: The initial user-level program
8801
8802 #include "types.h"
8803 #include "stat.h"
8804 #include "user.h"
8805 #include "fcntl.h"
8806
8807 char *argv[] = { "sh", 0 };
8808
8809 int
8810 main(void)
8811 {
8812     int pid, wpid;
8813
8814     if(open("console", O_RDWR) < 0){
8815         mknod("console", 1, 1);
8816         open("console", O_RDWR);
8817     }
8818     dup(0); // stdout
8819     dup(0); // stderr
8820
8821     for(;;){
8822         printf(1, "init: starting sh\n");
8823         pid = fork();
8824         if(pid < 0){
8825             printf(1, "init: fork failed\n");
8826             exit();
8827         }
8828         if(pid == 0){
8829             exec("sh", argv);
8830             printf(1, "init: exec sh failed\n");
8831             exit();
8832         }
8833         while((wpid=wait()) >= 0 && wpid != pid)
8834             printf(1, "zombie!\n");
8835     }
8836 }
8837
8838
8839
8840
8841
8842
8843
8844
8845
8846
8847
8848
8849

```

```

8850 // Shell.
8851 // 2015-12-21. Added very simple processing for builtin commands
8852
8853 #include "types.h"
8854 #include "user.h"
8855 #include "fcntl.h"
8856
8857 // Parsed command representation
8858 #define EXEC 1
8859 #define REDIR 2
8860 #define PIPE 3
8861 #define LIST 4
8862 #define BACK 5
8863
8864 #define MAXARGS 10
8865
8866 struct cmd {
8867     int type;
8868 };
8869
8870 struct execcmd {
8871     int type;
8872     char *argv[MAXARGS];
8873     char *eargv[MAXARGS];
8874 };
8875
8876 struct redircmd {
8877     int type;
8878     struct cmd *cmd;
8879     char *file;
8880     char *efile;
8881     int mode;
8882     int fd;
8883 };
8884
8885 struct pipecmd {
8886     int type;
8887     struct cmd *left;
8888     struct cmd *right;
8889 };
8890
8891 struct listcmd {
8892     int type;
8893     struct cmd *left;
8894     struct cmd *right;
8895 };
8896
8897
8898
8899

```



```

8900 struct backcmd {
8901     int type;
8902     struct cmd *cmd;
8903 };
8904
8905 int fork1(void); // Fork but panics on failure.
8906 void panic(char*);
8907 struct cmd *parsecmd(char*);
8908
8909 // Execute cmd. Never returns.
8910 void
8911 runcmd(struct cmd *cmd)
8912 {
8913     int p[2];
8914     struct backcmd *bcmd;
8915     struct execcmd *ecmd;
8916     struct listcmd *lcmd;
8917     struct pipecmd *pcmd;
8918     struct redircmd *rcmd;
8919
8920     if(cmd == 0)
8921         exit();
8922
8923     switch(cmd->type){
8924     default:
8925         panic("runcmd");
8926
8927     case EXEC:
8928         ecmd = (struct execcmd*)cmd;
8929         if(ecmd->argv[0] == 0)
8930             exit();
8931         exec(ecmd->argv[0], ecmd->argv);
8932         printf(2, "exec %s failed\n", ecmd->argv[0]);
8933         break;
8934
8935     case REDIR:
8936         rcmd = (struct redircmd*)cmd;
8937         close(rcmd->fd);
8938         if(open(rcmd->file, rcmd->mode) < 0){
8939             printf(2, "open %s failed\n", rcmd->file);
8940             exit();
8941         }
8942         runcmd(rcmd->cmd);
8943         break;
8944
8945     case LIST:
8946         lcmd = (struct listcmd*)cmd;
8947         if(fork1() == 0)
8948             runcmd(lcmd->left);
8949         wait();

```

```

8950     runcmd(lcmd->right);
8951     break;
8952
8953     case PIPE:
8954         pcmd = (struct pipecmd*)cmd;
8955         if(pipe(p) < 0)
8956             panic("pipe");
8957         if(fork1() == 0){
8958             close(1);
8959             dup(p[1]);
8960             close(p[0]);
8961             close(p[1]);
8962             runcmd(pcmd->left);
8963         }
8964         if(fork1() == 0){
8965             close(0);
8966             dup(p[0]);
8967             close(p[0]);
8968             close(p[1]);
8969             runcmd(pcmd->right);
8970         }
8971         close(p[0]);
8972         close(p[1]);
8973         wait();
8974         wait();
8975         break;
8976
8977     case BACK:
8978         bcmd = (struct backcmd*)cmd;
8979         if(fork1() == 0)
8980             runcmd(bcmd->cmd);
8981         break;
8982     }
8983     exit();
8984 }
8985
8986 int
8987 getcmd(char *buf, int nbuf)
8988 {
8989     printf(2, "$ ");
8990     memset(buf, 0, nbuf);
8991     gets(buf, nbuf);
8992     if(buf[0] == 0) // EOF
8993         return -1;
8994     return 0;
8995 }
8996
8997
8998
8999

```

```

9000 // ***** processing for shell builtins begins here *****
9001
9002 int
9003 strncmp(const char *p, const char *q, uint n)
9004 {
9005     while(n > 0 && *p && *p == *q)
9006         n--, p++, q++;
9007     if(n == 0)
9008         return 0;
9009     return (uchar)*p - (uchar)*q;
9010 }
9011
9012 int
9013 makeint(char *p)
9014 {
9015     int val = 0;
9016
9017     while ((*p >= '0') && (*p <= '9')) {
9018         val = 10*val + (*p-'0');
9019         ++p;
9020     }
9021     return val;
9022 }
9023
9024 int
9025 setbuiltin(char *p)
9026 {
9027     int i;
9028
9029     p += strlen("_set");
9030     while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
9031     if (strncmp("uid", p, 3) == 0) {
9032         p += strlen("uid");
9033         while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
9034         i = makeint(p); // ugly
9035         return (setuid(i));
9036     } else
9037     if (strncmp("gid", p, 3) == 0) {
9038         p += strlen("gid");
9039         while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
9040         i = makeint(p); // ugly
9041         return (setgid(i));
9042     }
9043     printf(2, "Invalid _set parameter\n");
9044     return -1;
9045 }
9046
9047
9048
9049

```

```

9050 int
9051 getbuiltin(char *p)
9052 {
9053     p += strlen("_get");
9054     while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
9055     if (strncmp("uid", p, 3) == 0) {
9056         printf(2, "%d\n", getuid());
9057         return 0;
9058     }
9059     if (strncmp("gid", p, 3) == 0) {
9060         printf(2, "%d\n", getgid());
9061         return 0;
9062     }
9063     printf(2, "Invalid _get parameter\n");
9064     return -1;
9065 }
9066
9067 typedef int funcPtr_t(char *);
9068 typedef struct {
9069     char *cmd;
9070     funcPtr_t *name;
9071 } dispatchTableEntry_t;
9072
9073 // Use a simple function dispatch table (FDT) to process builtin commands
9074 dispatchTableEntry_t fdt[] = {
9075     {"_set", setbuiltin},
9076     {"_get", getbuiltin}
9077 };
9078 int FDTcount = sizeof(fdt) / sizeof(fdt[0]); // # entris in FDT
9079
9080 void
9081 dobuiltin(char *cmd) {
9082     int i;
9083
9084     for (i=0; i<FDTcount; i++)
9085         if (strncmp(cmd, fdt[i].cmd, strlen(fdt[i].cmd)) == 0)
9086             (*fdt[i].name)(cmd);
9087 }
9088
9089
9090
9091
9092
9093
9094
9095
9096
9097
9098
9099

```

```

9100 // ***** processing for shell builtins ends here *****
9101
9102 int
9103 main(void)
9104 {
9105     static char buf[100];
9106     int fd;
9107
9108     // Assumes three file descriptors open.
9109     while((fd = open("console", O_RDWR)) >= 0){
9110         if(fd >= 3){
9111             close(fd);
9112             break;
9113         }
9114     }
9115
9116     // Read and run input commands.
9117     while(getcmd(buf, sizeof(buf)) >= 0){
9118 // add support for built-ins here. cd is a built-in
9119         if(buf[0] == 'c' && buf[1] == 'd' && buf[2] == ' '){
9120             // Clumsy but will have to do for now.
9121             // Chdir has no effect on the parent if run in the child.
9122             buf[strlen(buf)-1] = 0; // chop \n
9123             if(chdir(buf+3) < 0)
9124                 printf(2, "cannot cd %s\n", buf+3);
9125             continue;
9126         }
9127         if (buf[0] == '_') { // assume it is a builtin command
9128             dobuiltin(buf);
9129             continue;
9130         }
9131         if(fork1() == 0)
9132             runcmd(parsecmd(buf));
9133         wait();
9134     }
9135     exit();
9136 }
9137
9138 void
9139 panic(char *s)
9140 {
9141     printf(2, "%s\n", s);
9142     exit();
9143 }
9144
9145
9146
9147
9148
9149

```

```

9150 int
9151 fork1(void)
9152 {
9153     int pid;
9154
9155     pid = fork();
9156     if(pid == -1)
9157         panic("fork");
9158     return pid;
9159 }
9160
9161
9162
9163
9164
9165
9166
9167
9168
9169
9170
9171
9172
9173
9174
9175
9176
9177
9178
9179
9180
9181
9182
9183
9184
9185
9186
9187
9188
9189
9190
9191
9192
9193
9194
9195
9196
9197
9198
9199

```

```

9200 // Constructors
9201
9202 struct cmd*
9203 execcmd(void)
9204 {
9205     struct execcmd *cmd;
9206
9207     cmd = malloc(sizeof(*cmd));
9208     memset(cmd, 0, sizeof(*cmd));
9209     cmd->type = EXEC;
9210     return (struct cmd*)cmd;
9211 }
9212
9213 struct cmd*
9214 redircmd(struct cmd *subcmd, char *file, char *efile, int mode, int fd)
9215 {
9216     struct redircmd *cmd;
9217
9218     cmd = malloc(sizeof(*cmd));
9219     memset(cmd, 0, sizeof(*cmd));
9220     cmd->type = REDIR;
9221     cmd->cmd = subcmd;
9222     cmd->file = file;
9223     cmd->efile = efile;
9224     cmd->mode = mode;
9225     cmd->fd = fd;
9226     return (struct cmd*)cmd;
9227 }
9228
9229 struct cmd*
9230 pipecmd(struct cmd *left, struct cmd *right)
9231 {
9232     struct pipecmd *cmd;
9233
9234     cmd = malloc(sizeof(*cmd));
9235     memset(cmd, 0, sizeof(*cmd));
9236     cmd->type = PIPE;
9237     cmd->left = left;
9238     cmd->right = right;
9239     return (struct cmd*)cmd;
9240 }
9241
9242
9243
9244
9245
9246
9247
9248
9249

```

```

9250 struct cmd*
9251 listcmd(struct cmd *left, struct cmd *right)
9252 {
9253     struct listcmd *cmd;
9254
9255     cmd = malloc(sizeof(*cmd));
9256     memset(cmd, 0, sizeof(*cmd));
9257     cmd->type = LIST;
9258     cmd->left = left;
9259     cmd->right = right;
9260     return (struct cmd*)cmd;
9261 }
9262
9263 struct cmd*
9264 backcmd(struct cmd *subcmd)
9265 {
9266     struct backcmd *cmd;
9267
9268     cmd = malloc(sizeof(*cmd));
9269     memset(cmd, 0, sizeof(*cmd));
9270     cmd->type = BACK;
9271     cmd->cmd = subcmd;
9272     return (struct cmd*)cmd;
9273 }
9274
9275
9276
9277
9278
9279
9280
9281
9282
9283
9284
9285
9286
9287
9288
9289
9290
9291
9292
9293
9294
9295
9296
9297
9298
9299

```

```

9300 // Parsing
9301
9302 char whitespace[] = " \t\r\n\v";
9303 char symbols[] = "<|>&()";
9304
9305 int
9306 gettoken(char **ps, char *es, char **q, char **eq)
9307 {
9308     char *s;
9309     int ret;
9310
9311     s = *ps;
9312     while(s < es && strchr(whitespace, *s))
9313         s++;
9314     if(*q)
9315         *q = s;
9316     ret = *s;
9317     switch(*s){
9318     case 0:
9319         break;
9320     case '|':
9321     case '(':
9322     case ')':
9323     case ';':
9324     case '&':
9325     case '<':
9326         s++;
9327         break;
9328     case '>':
9329         s++;
9330         if(*s == '>'){
9331             ret = '+';
9332             s++;
9333         }
9334         break;
9335     default:
9336         ret = 'a';
9337         while(s < es && !strchr(whitespace, *s) && !strchr(symbols, *s))
9338             s++;
9339         break;
9340     }
9341     if(eq)
9342         *eq = s;
9343
9344     while(s < es && strchr(whitespace, *s))
9345         s++;
9346     *ps = s;
9347     return ret;
9348 }
9349

```

```

9350 int
9351 peek(char **ps, char *es, char *toks)
9352 {
9353     char *s;
9354
9355     s = *ps;
9356     while(s < es && strchr(whitespace, *s))
9357         s++;
9358     *ps = s;
9359     return *s && strchr(toks, *s);
9360 }
9361
9362 struct cmd *parseline(char**, char*);
9363 struct cmd *parsepipe(char**, char*);
9364 struct cmd *parseexec(char**, char*);
9365 struct cmd *nulterminate(struct cmd*);
9366
9367 struct cmd*
9368 parsecmd(char *s)
9369 {
9370     char *es;
9371     struct cmd *cmd;
9372
9373     es = s + strlen(s);
9374     cmd = parseline(&s, es);
9375     peek(&s, es, "");
9376     if(s != es){
9377         printf(2, "leftovers: %s\n", s);
9378         panic("syntax");
9379     }
9380     nulterminate(cmd);
9381     return cmd;
9382 }
9383
9384 struct cmd*
9385 parseline(char **ps, char *es)
9386 {
9387     struct cmd *cmd;
9388
9389     cmd = parsepipe(ps, es);
9390     while(peek(ps, es, "&")){
9391         gettoken(ps, es, 0, 0);
9392         cmd = backcmd(cmd);
9393     }
9394     if(peek(ps, es, ";")){
9395         gettoken(ps, es, 0, 0);
9396         cmd = listcmd(cmd, parseline(ps, es));
9397     }
9398     return cmd;
9399 }

```

```

9400 struct cmd*
9401 parsepipe(char **ps, char *es)
9402 {
9403     struct cmd *cmd;
9404
9405     cmd = parseexec(ps, es);
9406     if(peek(ps, es, "|")){
9407         gettoken(ps, es, 0, 0);
9408         cmd = pipecmd(cmd, parsepipe(ps, es));
9409     }
9410     return cmd;
9411 }
9412
9413 struct cmd*
9414 parseredirs(struct cmd *cmd, char **ps, char *es)
9415 {
9416     int tok;
9417     char *q, *eq;
9418
9419     while(peek(ps, es, "<>")){
9420         tok = gettoken(ps, es, 0, 0);
9421         if(gettoken(ps, es, &q, &eq) != 'a')
9422             panic("missing file for redirection");
9423         switch(tok){
9424             case '<':
9425                 cmd = redircmd(cmd, q, eq, O_RDONLY, 0);
9426                 break;
9427             case '>':
9428                 cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
9429                 break;
9430             case '+': // >>
9431                 cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
9432                 break;
9433         }
9434     }
9435     return cmd;
9436 }
9437
9438
9439
9440
9441
9442
9443
9444
9445
9446
9447
9448
9449

```

```

9450 struct cmd*
9451 parseblock(char **ps, char *es)
9452 {
9453     struct cmd *cmd;
9454
9455     if(!peek(ps, es, "("))
9456         panic("parseblock");
9457     gettoken(ps, es, 0, 0);
9458     cmd = parseline(ps, es);
9459     if(!peek(ps, es, ")"))
9460         panic("syntax - missing )");
9461     gettoken(ps, es, 0, 0);
9462     cmd = parseredirs(cmd, ps, es);
9463     return cmd;
9464 }
9465
9466 struct cmd*
9467 parseexec(char **ps, char *es)
9468 {
9469     char *q, *eq;
9470     int tok, argc;
9471     struct execcmd *cmd;
9472     struct cmd *ret;
9473
9474     if(peek(ps, es, "("))
9475         return parseblock(ps, es);
9476
9477     ret = execcmd();
9478     cmd = (struct execcmd*)ret;
9479
9480     argc = 0;
9481     ret = parseredirs(ret, ps, es);
9482     while(!peek(ps, es, "|)&")){
9483         if((tok=gettoken(ps, es, &q, &eq)) == 0)
9484             break;
9485         if(tok != 'a')
9486             panic("syntax");
9487         cmd->argv[argc] = q;
9488         cmd->eargv[argc] = eq;
9489         argc++;
9490         if(argc >= MAXARGS)
9491             panic("too many args");
9492         ret = parseredirs(ret, ps, es);
9493     }
9494     cmd->argv[argc] = 0;
9495     cmd->eargv[argc] = 0;
9496     return ret;
9497 }
9498
9499

```

```

9500 // NUL-terminate all the counted strings.
9501 struct cmd*
9502 nulterminate(struct cmd *cmd)
9503 {
9504     int i;
9505     struct backcmd *bcmd;
9506     struct execcmd *ecmd;
9507     struct listcmd *lcmd;
9508     struct pipecmd *pcmd;
9509     struct redircmd *rcmd;
9510
9511     if(cmd == 0)
9512         return 0;
9513
9514     switch(cmd->type){
9515     case EXEC:
9516         ecmd = (struct execcmd*)cmd;
9517         for(i=0; ecmd->argv[i]; i++)
9518             *ecmd->eargv[i] = 0;
9519         break;
9520
9521     case REDIR:
9522         rcmd = (struct redircmd*)cmd;
9523         nulterminate(rcmd->cmd);
9524         *rcmd->efile = 0;
9525         break;
9526
9527     case PIPE:
9528         pcmd = (struct pipecmd*)cmd;
9529         nulterminate(pcmd->left);
9530         nulterminate(pcmd->right);
9531         break;
9532
9533     case LIST:
9534         lcmd = (struct listcmd*)cmd;
9535         nulterminate(lcmd->left);
9536         nulterminate(lcmd->right);
9537         break;
9538
9539     case BACK:
9540         bcmd = (struct backcmd*)cmd;
9541         nulterminate(bcmd->cmd);
9542         break;
9543     }
9544     return cmd;
9545 }
9546
9547
9548
9549

```

```

9550 #include "asm.h"
9551 #include "memlayout.h"
9552 #include "mmu.h"
9553
9554 # Start the first CPU: switch to 32-bit protected mode, jump into C.
9555 # The BIOS loads this code from the first sector of the hard disk into
9556 # memory at physical address 0x7c00 and starts executing in real mode
9557 # with %cs=0 %ip=7c00.
9558
9559 .code16                                # Assemble for 16-bit mode
9560 .globl start
9561 start:
9562     cli                                # BIOS enabled interrupts; disable
9563
9564     # Zero data segment registers DS, ES, and SS.
9565     xorw    %ax,%ax                    # Set %ax to zero
9566     movw    %ax,%ds                    # -> Data Segment
9567     movw    %ax,%es                    # -> Extra Segment
9568     movw    %ax,%ss                    # -> Stack Segment
9569
9570     # Physical address line A20 is tied to zero so that the first PCs
9571     # with 2 MB would run software that assumed 1 MB. Undo that.
9572 seta20.1:
9573     inb     $0x64,%al                  # Wait for not busy
9574     testb   $0x2,%al
9575     jnz     seta20.1
9576
9577     movb     $0xd1,%al                 # 0xd1 -> port 0x64
9578     outb     %al,$0x64
9579
9580 seta20.2:
9581     inb     $0x64,%al                  # Wait for not busy
9582     testb   $0x2,%al
9583     jnz     seta20.2
9584
9585     movb     $0xdf,%al                 # 0xdf -> port 0x60
9586     outb     %al,$0x60
9587
9588     # Switch from real to protected mode. Use a bootstrap GDT that makes
9589     # virtual addresses map directly to physical addresses so that the
9590     # effective memory map doesn't change during the transition.
9591     lgdt     gdtdesc
9592     movl     %cr0,%eax
9593     orl      $CR0_PE,%eax
9594     movl     %eax,%cr0
9595
9596
9597
9598
9599

```

```

9600 # Complete transition to 32-bit protected mode by using long jmp
9601 # to reload %cs and %eip. The segment descriptors are set up with no
9602 # translation, so that the mapping is still the identity mapping.
9603 ljmp $(SEG_KCODE<<3), $start32
9604
9605 .code32 # Tell assembler to generate 32-bit code now.
9606 start32:
9607 # Set up the protected-mode data segment registers
9608 movw $(SEG_KDATA<<3), %ax # Our data segment selector
9609 movw %ax, %ds # -> DS: Data Segment
9610 movw %ax, %es # -> ES: Extra Segment
9611 movw %ax, %ss # -> SS: Stack Segment
9612 movw $0, %ax # Zero segments not ready for use
9613 movw %ax, %fs # -> FS
9614 movw %ax, %gs # -> GS
9615
9616 # Set up the stack pointer and call into C.
9617 movl $start, %esp
9618 call bootmain
9619
9620 # If bootmain returns (it shouldn't), trigger a Bochs
9621 # breakpoint if running under Bochs, then loop.
9622 movw $0x8a00, %ax # 0x8a00 -> port 0x8a00
9623 movw %ax, %dx
9624 outw %ax, %dx
9625 movw $0x8ae0, %ax # 0x8ae0 -> port 0x8a00
9626 outw %ax, %dx
9627 spin:
9628 jmp spin
9629
9630 # Bootstrap GDT
9631 .p2align 2 # force 4 byte alignment
9632 gdt:
9633 SEG_NULLASM # null seg
9634 SEG_ASM(STA_X|STA_R, 0x0, 0xffffffff) # code seg
9635 SEG_ASM(STA_W, 0x0, 0xffffffff) # data seg
9636
9637 gdtdesc:
9638 .word (gdtdesc - gdt - 1) # sizeof(gdt) - 1
9639 .long gdt # address gdt
9640
9641
9642
9643
9644
9645
9646
9647
9648
9649

```

```

9650 // Boot loader.
9651 //
9652 // Part of the boot block, along with bootasm.S, which calls bootmain().
9653 // bootasm.S has put the processor into protected 32-bit mode.
9654 // bootmain() loads an ELF kernel image from the disk starting at
9655 // sector 1 and then jumps to the kernel entry routine.
9656
9657 #include "types.h"
9658 #include "elf.h"
9659 #include "x86.h"
9660 #include "memlayout.h"
9661
9662 #define SECTSIZE 512
9663
9664 void readseg(uchar*, uint, uint);
9665
9666 void
9667 bootmain(void)
9668 {
9669     struct elfhdr *elf;
9670     struct proghdr *ph, *eph;
9671     void (*entry)(void);
9672     uchar* pa;
9673
9674     elf = (struct elfhdr*)0x10000; // scratch space
9675
9676     // Read 1st page off disk
9677     readseg((uchar*)elf, 4096, 0);
9678
9679     // Is this an ELF executable?
9680     if(elf->magic != ELF_MAGIC)
9681         return; // let bootasm.S handle error
9682
9683     // Load each program segment (ignores ph flags).
9684     ph = (struct proghdr*)((uchar*)elf + elf->phoff);
9685     eph = ph + elf->phnum;
9686     for(; ph < eph; ph++){
9687         pa = (uchar*)ph->paddr;
9688         readseg(pa, ph->filesz, ph->off);
9689         if(ph->memsz > ph->filesz)
9690             stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
9691     }
9692
9693     // Call the entry point from the ELF header.
9694     // Does not return!
9695     entry = (void(*) (void))(elf->entry);
9696     entry();
9697 }
9698
9699

```



```

9700 void
9701 waitdisk(void)
9702 {
9703     // Wait for disk ready.
9704     while((inb(0x1F7) & 0xC0) != 0x40)
9705         ;
9706 }
9707
9708 // Read a single sector at offset into dst.
9709 void
9710 readsect(void *dst, uint offset)
9711 {
9712     // Issue command.
9713     waitdisk();
9714     outb(0x1F2, 1);    // count = 1
9715     outb(0x1F3, offset);
9716     outb(0x1F4, offset >> 8);
9717     outb(0x1F5, offset >> 16);
9718     outb(0x1F6, (offset >> 24) | 0xE0);
9719     outb(0x1F7, 0x20); // cmd 0x20 - read sectors
9720
9721     // Read data.
9722     waitdisk();
9723     insl(0x1F0, dst, SECTSIZE/4);
9724 }
9725
9726 // Read 'count' bytes at 'offset' from kernel into physical address 'pa'.
9727 // Might copy more than asked.
9728 void
9729 readseg(uchar* pa, uint count, uint offset)
9730 {
9731     uchar* epa;
9732
9733     epa = pa + count;
9734
9735     // Round down to sector boundary.
9736     pa -= offset % SECTSIZE;
9737
9738     // Translate from bytes to sectors; kernel starts at sector 1.
9739     offset = (offset / SECTSIZE) + 1;
9740
9741     // If this is too slow, we could read lots of sectors at a time.
9742     // We'd write more to memory than asked, but it doesn't matter --
9743     // we load in increasing order.
9744     for(; pa < epa; pa += SECTSIZE, offset++)
9745         readsect(pa, offset);
9746 }
9747
9748
9749

```

```

9750 #include "types.h"
9751 #include "user.h"
9752 #include "date.h"
9753
9754 int main (int argc, char *argv[])
9755 {
9756     struct rtcdate r;
9757
9758     if(date(&r)){
9759         printf(2, "date_failed\n");
9760         exit();
9761     }
9762
9763     printf(1, "%d/%d/%d %d: %d: %d\n", r.month, r.day, r.year, r.hour, r.minut
9764         exit();
9765 }
9766
9767
9768
9769
9770
9771
9772
9773
9774
9775
9776
9777
9778
9779
9780
9781
9782
9783
9784
9785
9786
9787
9788
9789
9790
9791
9792
9793
9794
9795
9796
9797
9798
9799

```

```

9800 #include "types.h"
9801 #include "user.h"
9802 #include "date.h"
9803
9804 int main(int argc, char *argv[]){
9805     int min;
9806     int sec;
9807
9808     struct rtcdate startTime;
9809     struct rtcdate endTime;
9810
9811
9812
9813     if(date(&startTime)){
9814         printf(2, "date_failed");
9815         exit();
9816     }
9817
9818
9819
9820
9821     int pid = fork();
9822
9823     if(pid > 0){
9824         wait();
9825
9826         if(date(&endTime)){
9827             printf(2, "date2_failed");
9828             exit();
9829         }
9830
9831         min = endTime.minute - startTime.minute;
9832         sec = endTime.second - startTime.second;
9833
9834         if(sec < 0){
9835             sec = sec + 60;
9836             min = min-1;
9837         }
9838
9839         printf(1, "%s %s %d %s %d %s\n", argv[1], "runs in", min, "minute(s)", sec
9840
9841         exit();
9842     }
9843     else if(pid == 0){
9844         exec(argv[1], argv+1);
9845         printf(1, "Error: exec returned");
9846
9847         exit();
9848     }
9849     else{

```

```

9850         printf(2, "fork error\n");
9851         exit();
9852     }
9853
9854
9855
9856
9857
9858     exit();
9859
9860 }
9861
9862
9863
9864
9865
9866
9867
9868
9869
9870
9871
9872
9873
9874
9875
9876
9877
9878
9879
9880
9881
9882
9883
9884
9885
9886
9887
9888
9889
9890
9891
9892
9893
9894
9895
9896
9897
9898
9899

```

```
9900 struct stat;
9901 struct rtcdate;
9902 struct uproc;
9903
9904 // system calls
9905 int fork(void);
9906 int exit(void) __attribute__((noreturn));
9907 int wait(void);
9908 int pipe(int*);
9909 int write(int, void*, int);
9910 int read(int, void*, int);
9911 int close(int);
9912 int kill(int);
9913 int exec(char*, char**);
9914 int open(char*, int);
9915 int mknod(char*, short, short);
9916 int unlink(char*);
9917 int fstat(int fd, struct stat*);
9918 int link(char*, char*);
9919 int mkdir(char*);
9920 int chdir(char*);
9921 int dup(int);
9922 int getpid(void);
9923 char* sbrk(int);
9924 int sleep(int);
9925 int uptime(void);
9926 int halt(void);
9927 int date(struct rtcdate*);
9928 int getuid(void);
9929 int getgid(void);
9930 int getppid(void);
9931 int setuid(int);
9932 int setgid(int);
9933 int getprocs(int, struct uproc*);
9934 // ulib.c
9935 int stat(char*, struct stat*);
9936 char* strcpy(char*, char*);
9937 void *memmove(void*, void*, int);
9938 char* strchr(const char*, char c);
9939 int strcmp(const char*, const char*);
9940 void printf(int, char*, ...);
9941 char* gets(char*, int max);
9942 uint strlen(char*);
9943 void* memset(void*, int, uint);
9944 void* malloc(uint);
9945 void free(void*);
9946 int atoi(const char*);
9947
9948
9949
```

```
9950 // halt the system.
9951 #include "types.h"
9952 #include "user.h"
9953
9954 int
9955 main(void) {
9956     halt();
9957     return 0;
9958 }
9959
9960
9961
9962
9963
9964
9965
9966
9967
9968
9969
9970
9971
9972
9973
9974
9975
9976
9977
9978
9979
9980
9981
9982
9983
9984
9985
9986
9987
9988
9989
9990
9991
9992
9993
9994
9995
9996
9997
9998
9999
```

```

10000 struct uproc {
10001     int pid;
10002     uint uid;
10003     uint gid;
10004     int ppid;
10005     uint size;
10006     char name[16];
10007     char state[20];
10008 };
10009
10010
10011
10012
10013
10014
10015
10016
10017
10018
10019
10020
10021
10022
10023
10024
10025
10026
10027
10028
10029
10030
10031
10032
10033
10034
10035
10036
10037
10038
10039
10040
10041
10042
10043
10044
10045
10046
10047
10048
10049

```

```

10050 #include "types.h"
10051 #include "user.h"
10052 #include "ps.h"
10053
10054 #define MAX_PROC 7
10055 int
10056 main(int argc, char *argv[])
10057 {
10058     int processes;
10059     int i;
10060     struct uproc prc[MAX_PROC];
10061
10062     processes = getprocs(MAX_PROC, prc);
10063     // printf(1, "%d\n", processes);
10064
10065     if(processes == -1){
10066         printf(2, "getprocs failed\n");
10067         exit();
10068     }
10069
10070     printf(1, "\nPID PPID UID GID STATE SIZE NAME");
10071
10072     for(i=0; i<processes;i++){
10073         printf(1, "\n%d %d %d %d %s %d %s\n ", prc[i].pid, prc[i].ppid,
10074             }
10075     }
10076     exit();
10077 }
10078
10079
10080
10081
10082
10083
10084
10085
10086
10087
10088
10089
10090
10091
10092
10093
10094
10095
10096
10097
10098
10099

```

```
10100 #include "param.h"
10101 #include "types.h"
10102 #include "stat.h"
10103 #include "user.h"
10104 #include "fs.h"
10105 #include "fcntl.h"
10106 #include "syscall.h"
10107 #include "traps.h"
10108 #include "memlayout.h"
10109
10110
10111 int
10112 main(int argc, char *argv[])
10113 {
10114     int uid, gid, ppid;
10115     uid = getuid();
10116     printf(1, "Current UID is: %d\n", uid);
10117     printf(1, "Setting UID to 4\n");
10118
10119     setuid(4);
10120     uid = getuid();
10121     printf(1, "Current UID is: %d\n", uid);
10122
10123     gid = getgid();
10124     printf(1, "Current GID is: %d\n", gid);
10125     printf(1, "Setting GID to 100\n");
10126
10127     setgid(100);
10128     gid = getgid();
10129     printf(1, "Current GID is: %d\n", gid);
10130
10131     ppid = getppid();
10132     printf(1, "Current PPID is: %d\n", ppid);
10133
10134
10135
10136     exit();
10137 }
10138
10139
10140
10141
10142
10143
10144
10145
10146
10147
10148
10149
```